



Vis  
VAST • INFOVIS • SCIVIS  
BIOVIS • LDAV  
2013

# Nanocubes for Real-Time Exploration of Spatiotemporal Datasets

Lauro Lins, James T. Klosowski, Carlos Scheidegger  
AT&T Labs Research

<http://nanocubes.net>

We are interested in phenomena  
with a spatiotemporal component

tweets

<u>Latitude</u>	<u>Longitude</u>	<u>device</u>	<u>time</u>
23.4	-40.3	iPhone	Aug 3, 2011 10:00
31.2	-41.3	Android	Aug 3, 2011 10:05
27.8	-39.3	iPhone	Aug 3, 2011 10:07
...	...	...	...
26.1	-38.1	Android	Jun 8, 2012 21:03
27.2	-44.3	Android	Jun 8, 2012 21:04
24.2	-39.7	iPhone	Jun 8, 2012 21:10

We are interested in phenomena  
with a spatiotemporal component

## tweets

<u>Latitude</u>	<u>Longitude</u>	<u>device</u>	<u>time</u>
23.4	-40.3	iPhone	Aug 3, 2011 10:00
31.2	-41.3	Android	Aug 3, 2011 10:05
27.8	-39.3	iPhone	Aug 3, 2011 10:07
...	...	...	...
26.1	-38.1	Android	Jun 8, 2012 21:03
27.2	-44.3	Android	Jun 8, 2012 21:04
24.2	-39.7	iPhone	Jun 8, 2012 21:

How do we  
make sense of  
this...

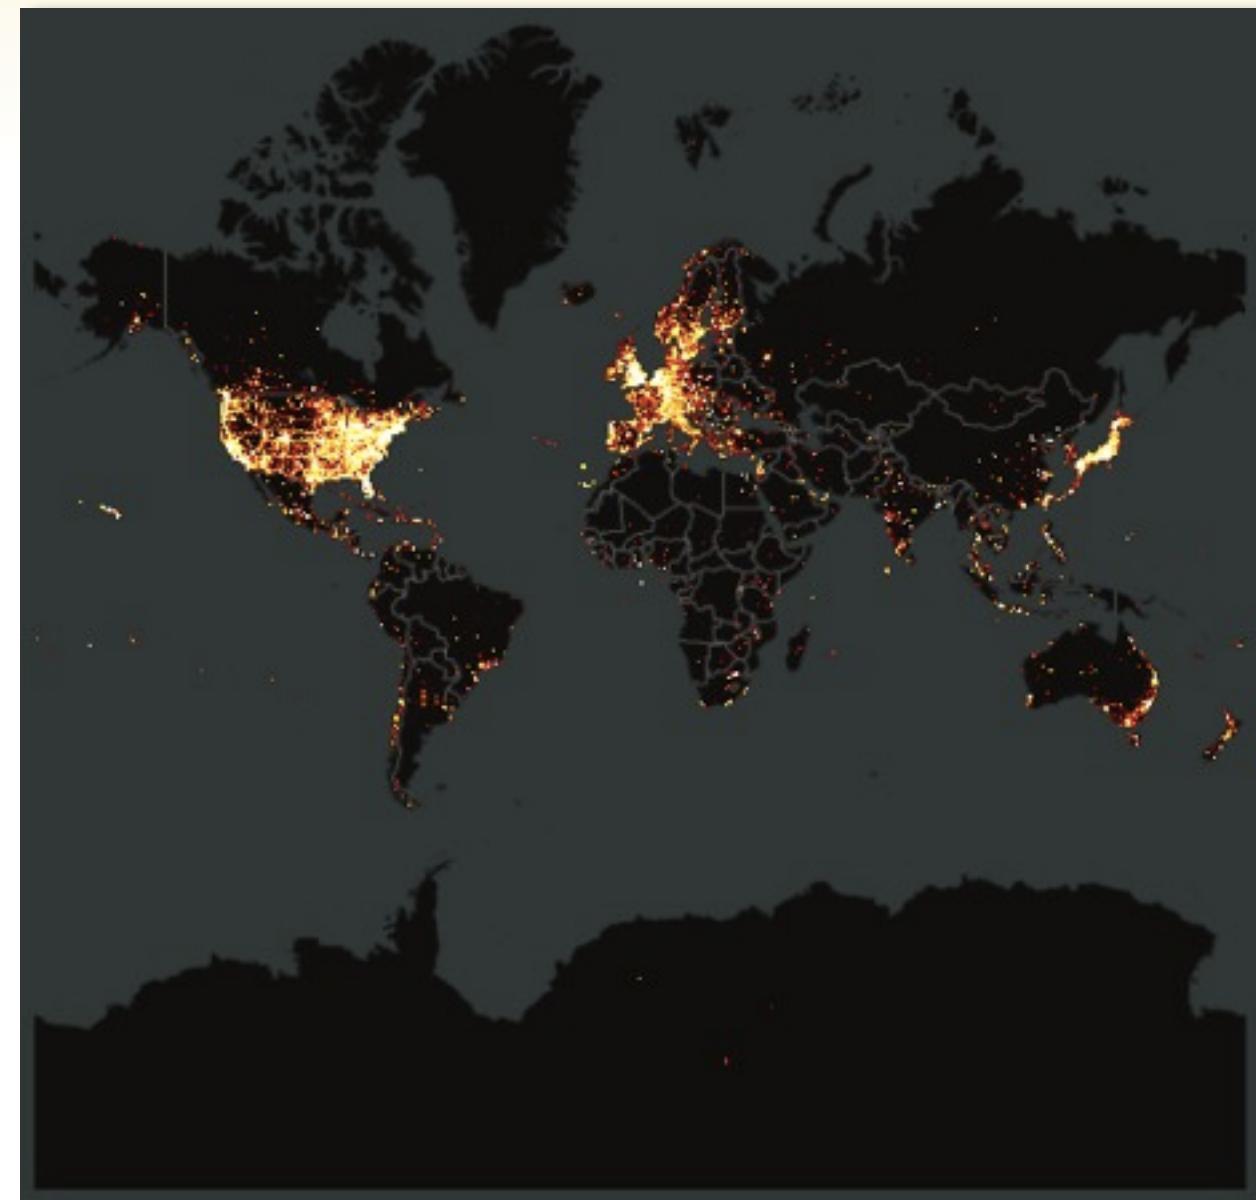
when there's a  
billion of  
these?

Demo

# Problem Statement

How can we build interactive visualization interfaces when spatiotemporal datasets are large?

Query: produce a count heatmap of the world  
for all points in my dataset



Result

Query: produce a count heatmap of the world for all points in my database



Result

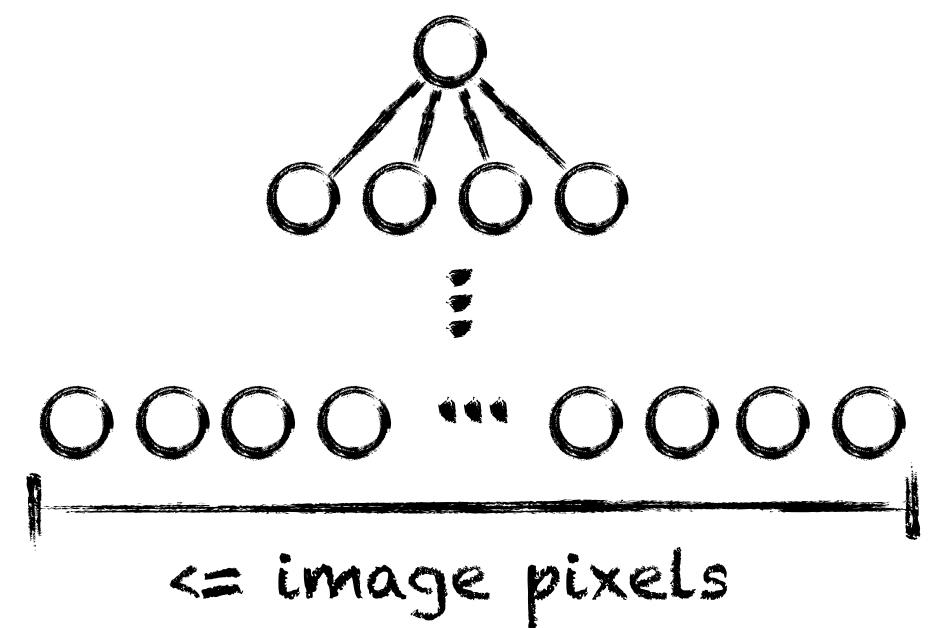
Latitude	longitude	device	time
23.4	-40.3	iPhone	Aug 3, 2011 10:00
31.2	-41.3	Android	Aug 3, 2011 10:05
27.8	-39.3	iPhone	Aug 3, 2011 10:07
-	..	..	..
26.1	-38.1	Android	Jun 8, 2012 21:03
27.2	-44.3	Android	Jun 8, 2012 21:04
24.2	-39.7	iPhone	Jun 8, 2012 21:10

If no aggregation was pre-computed then this query is proportional to "n"

Query: produce a count heatmap of the world for all points in my database



Result



Latitude	Longitude	device	time
23.4	-40.3	iPhone	Aug 3, 2011 10:00
31.2	-41.3	Android	Aug 3, 2011 10:05
27.8	-39.3	iPhone	Aug 3, 2011 10:07
-	...	...	...
26.1	-38.1	Android	Jun 8, 2012 21:03
27.2	-44.3	Android	Jun 8, 2012 21:04
24.2	-39.7	iPhone	Jun 8, 2012 21:10

If no aggregation was pre-computed then this query is proportional to "n"

If we pre-aggregate counts (e.g. quadtree) the query time becomes proportional to the number of reported pixels

Query: produce a count heatmap of the world  
for all points in my database

latitude	longitude	device	time
37.7749	-122.4194	iPhone 6S	2016-02-14T18:00:00Z
37.7749	-122.4194	iPhone 6S	2016-02-14T18:00:00Z

What if data has time dimension? and  
other categorical variables?



Result

What to pre-compute?



<= image pixels

pixels

if we pre-aggregate counts  
(e.g. quadtree) the query  
time becomes proportional  
to the number of reported

# Data Cubes!

[1987, J.Gray et al. Data cube: A relational aggregation operator generalizing group-by, cross-tab, and sub-totals]

... well known in Vis!

[2002, C. Stolle et al. Polaris]

[2003, C. Stolle et al. Multiscale visualization using data cubes]

[2007, J. Mackinlay et al. Show me]

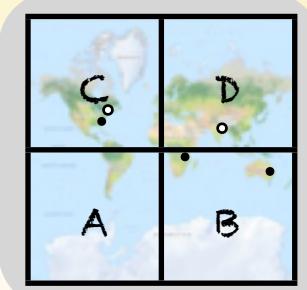
[2012, S. Kandol et al. Profiler]

[2013, Z. Liu et al. imMens]

# What is a Nanocube?

Nanocube is a **Sparse** **in-Memory** **Data Cube** of **Spatiotemporal** **Events**

# Building a Nanocube

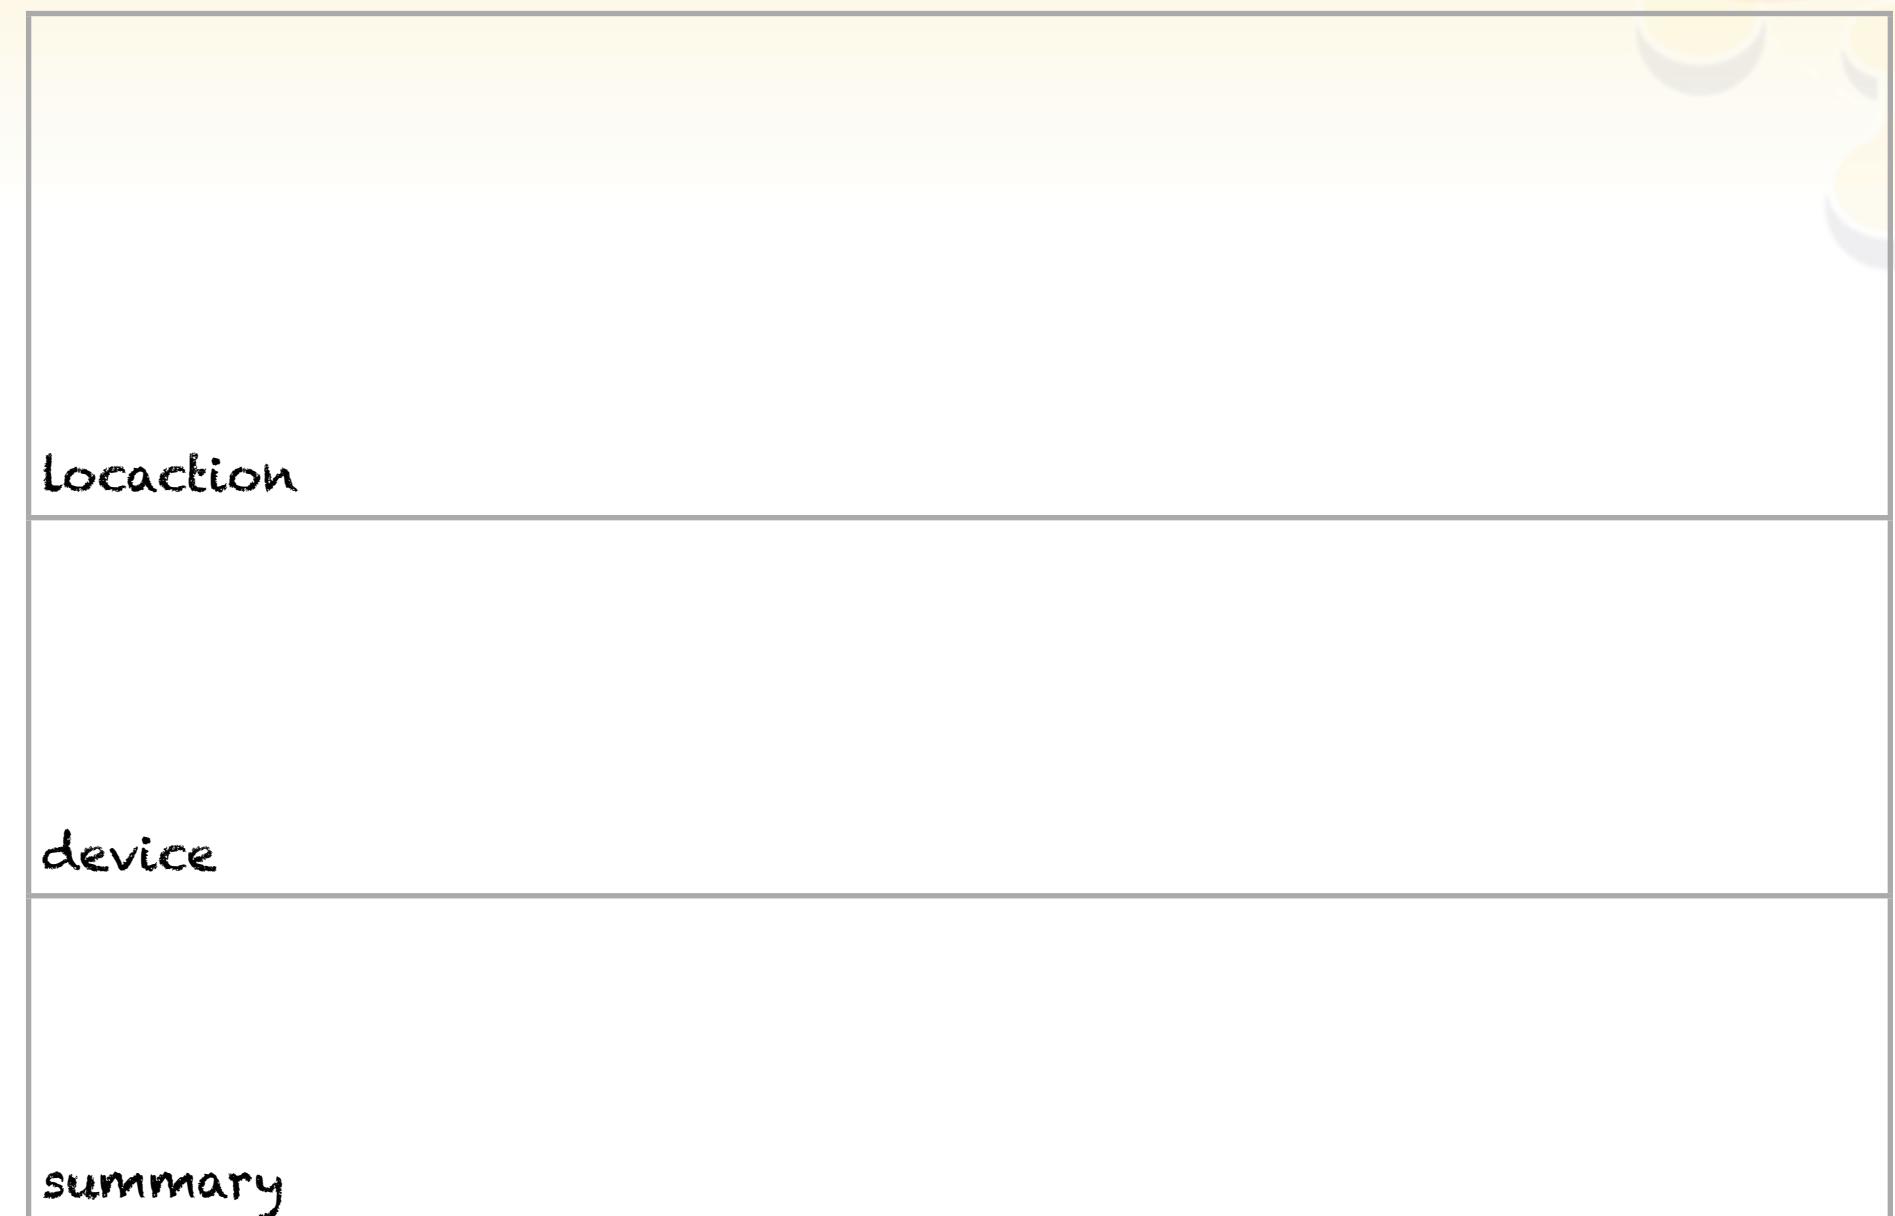


C3	C4	D3	D4
C1	C2	D1	D2
A3	A4	B3	B4
A1	A2	B1	B2

Android  
iPhone

<loc1, loc2>	<device>	time
C, C2	Android	1
C, C2	iPhone	2
B, B3	iPhone	2
D, D1	Android	3
B, B4	iPhone	5

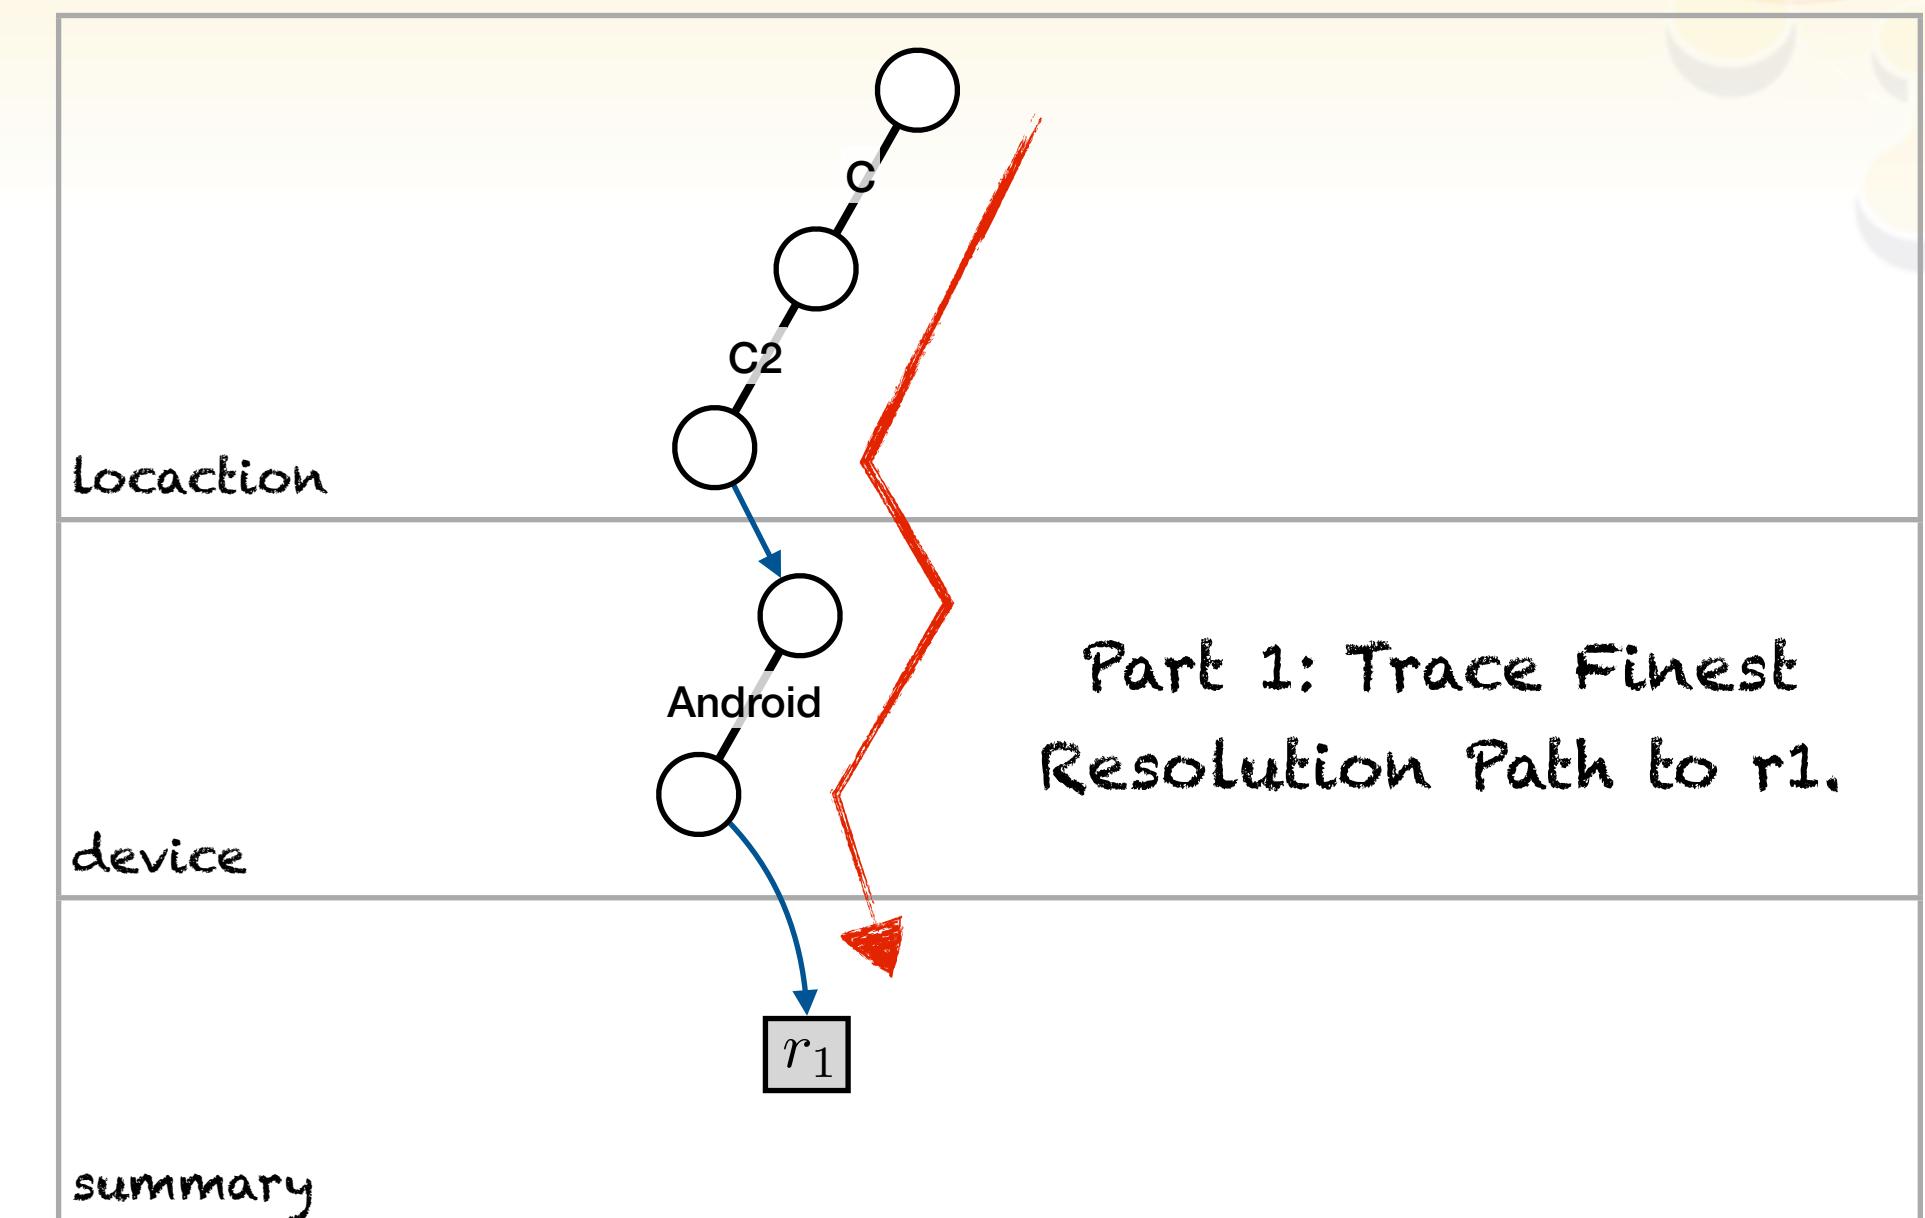
Insert ( $r_1$ )



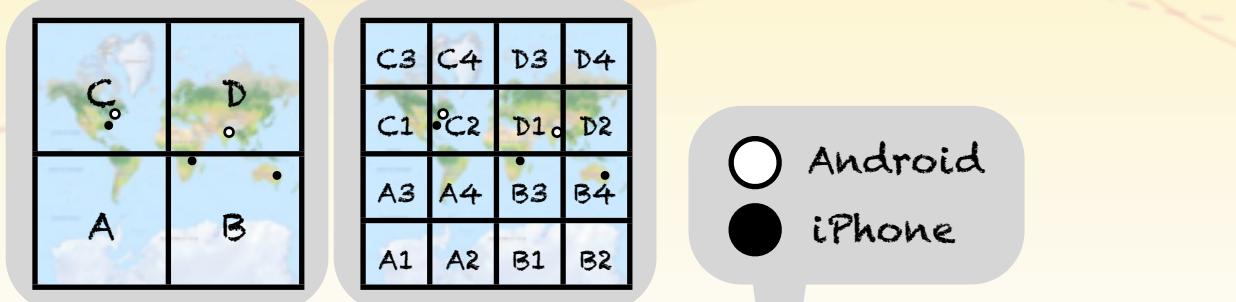
# Building a Nanocube

**Insert ( $r_1$ )**

$\langle loc1, loc2 \rangle$	$\langle device \rangle$	time
C, C2	Android	1
C, C2	iPhone	2
B, B3	iPhone	2
D, D1	Android	3
B, B4	iPhone	5

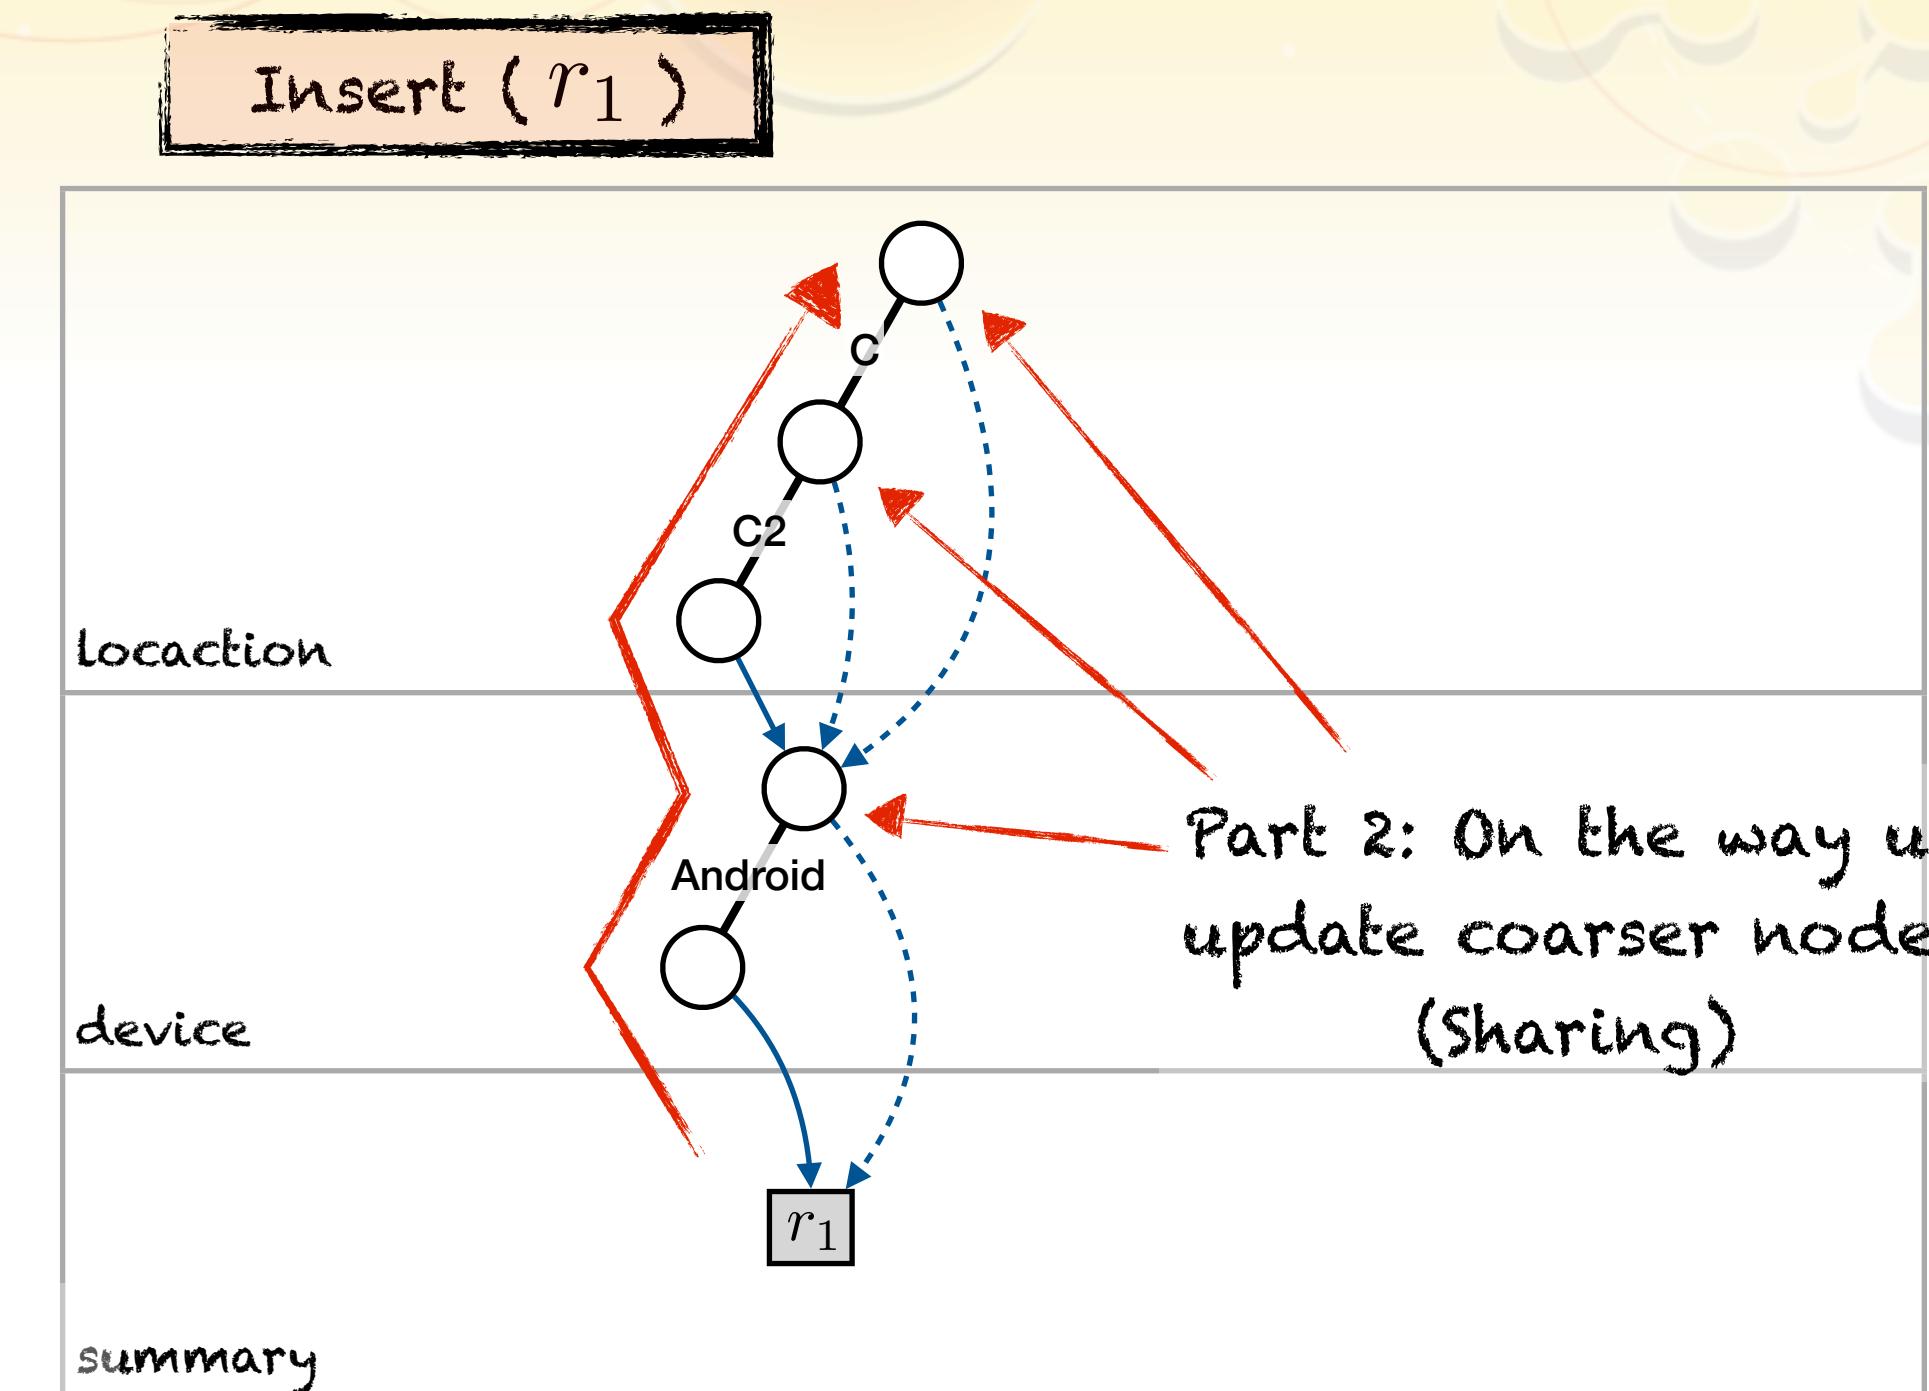


# Building a Nanocube



The diagram illustrates the hierarchical construction of a Nanocube. At the top left is a 2x2 grid labeled A and B, with a point C in the top-left cell. To its right is a 4x4 matrix labeled A1 through A4, B1 through B4, C1 through C4, and D1 through D4. Below these are two legends: one for 'Android' represented by an open circle, and one for 'iPhone' represented by a solid black circle.

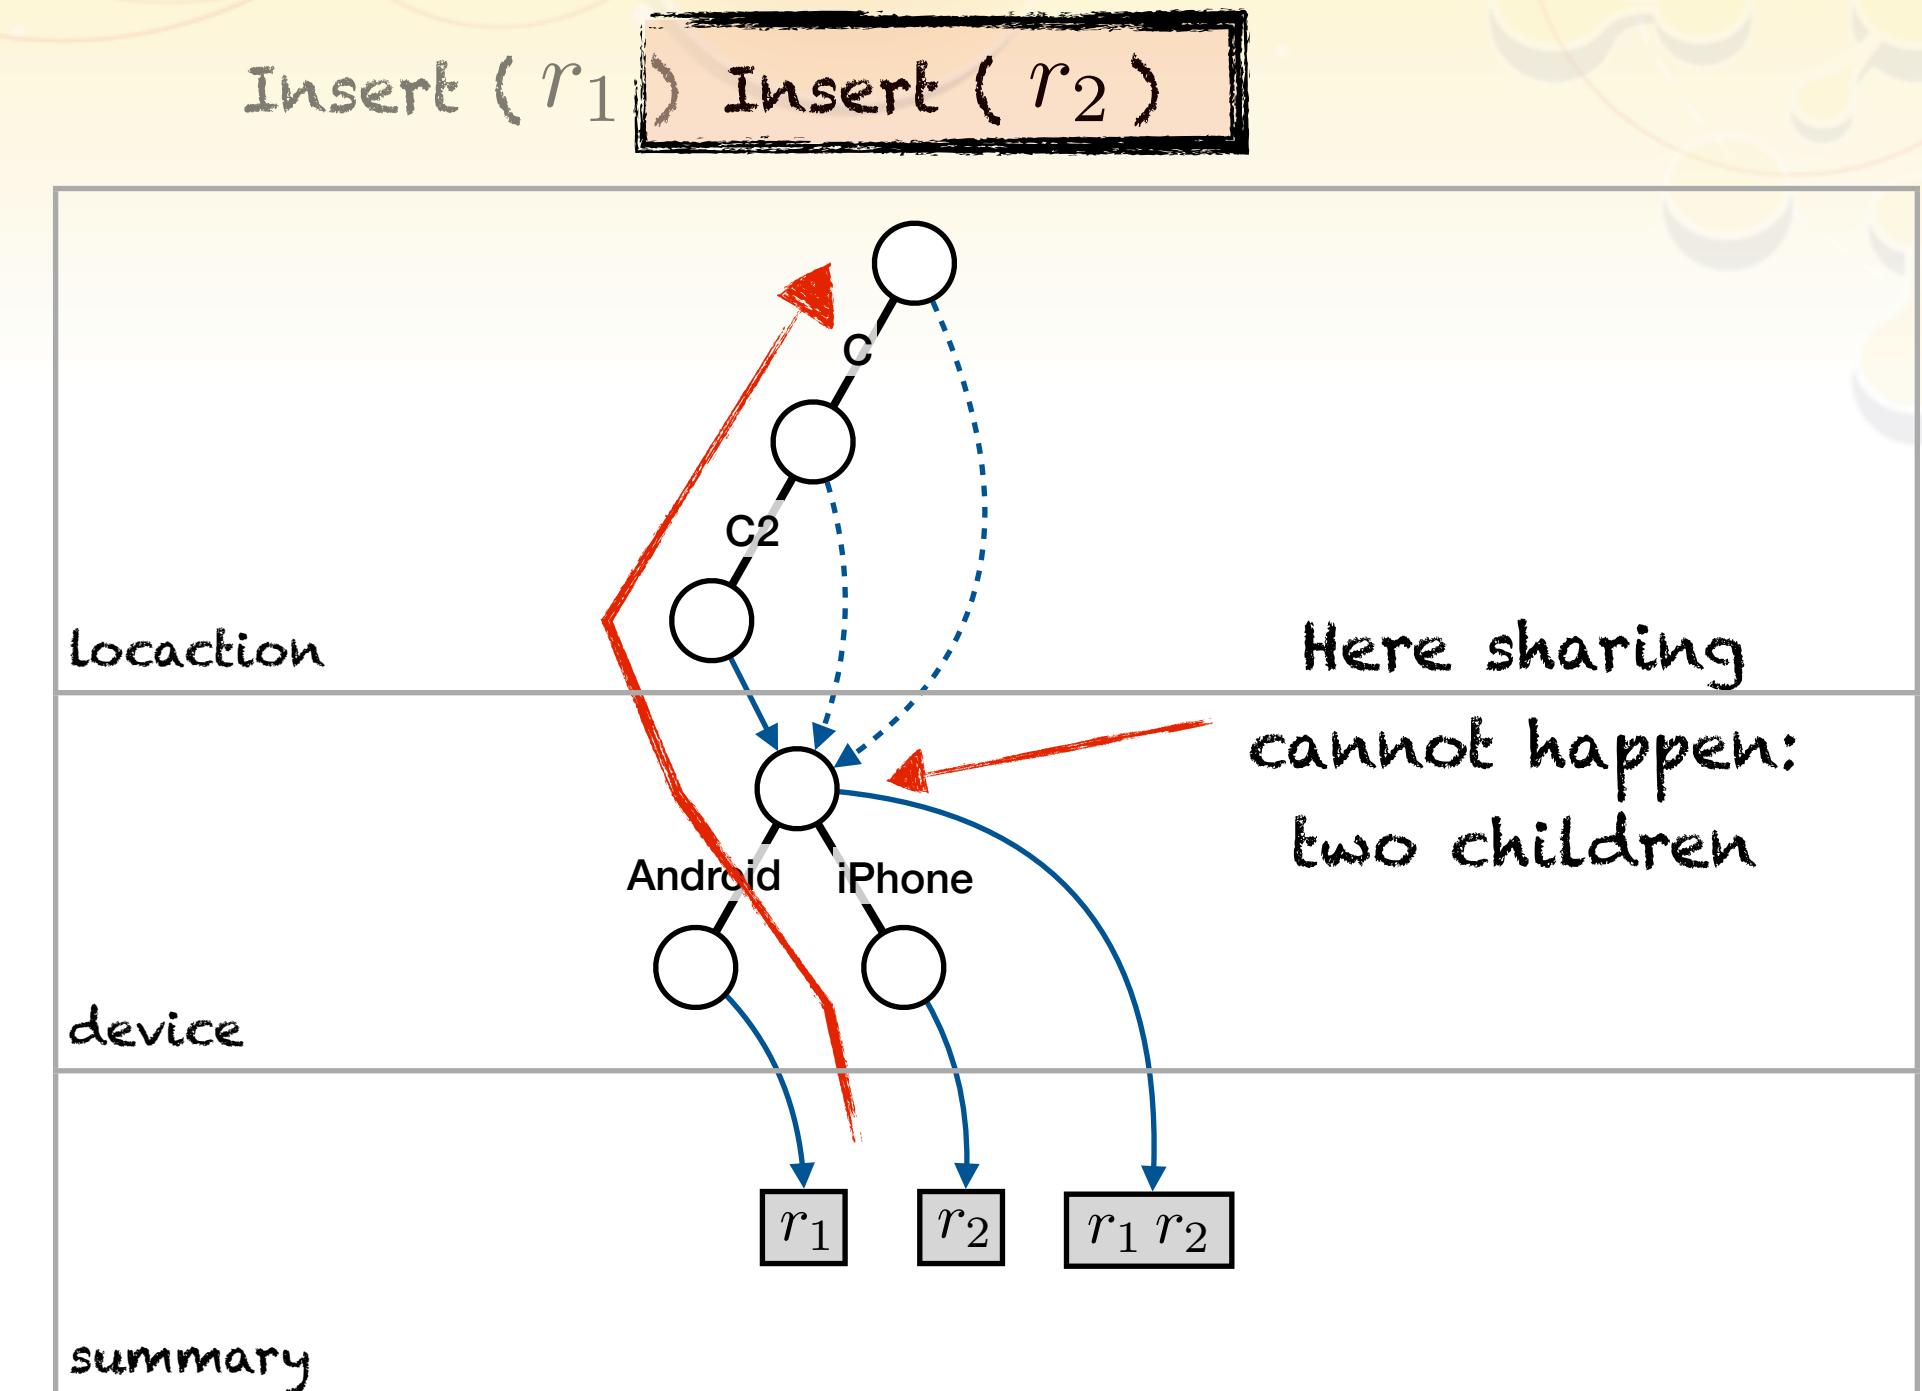
<loc1, loc2>		<device>	time
C, C2	Android	1	
C, C2	iPhone	2	
B, B3	iPhone	2	
D, D1	Android	3	
B, B4	iPhone	5	



# Building a Nanocube

`<loc1, loc2>` `<device>` time

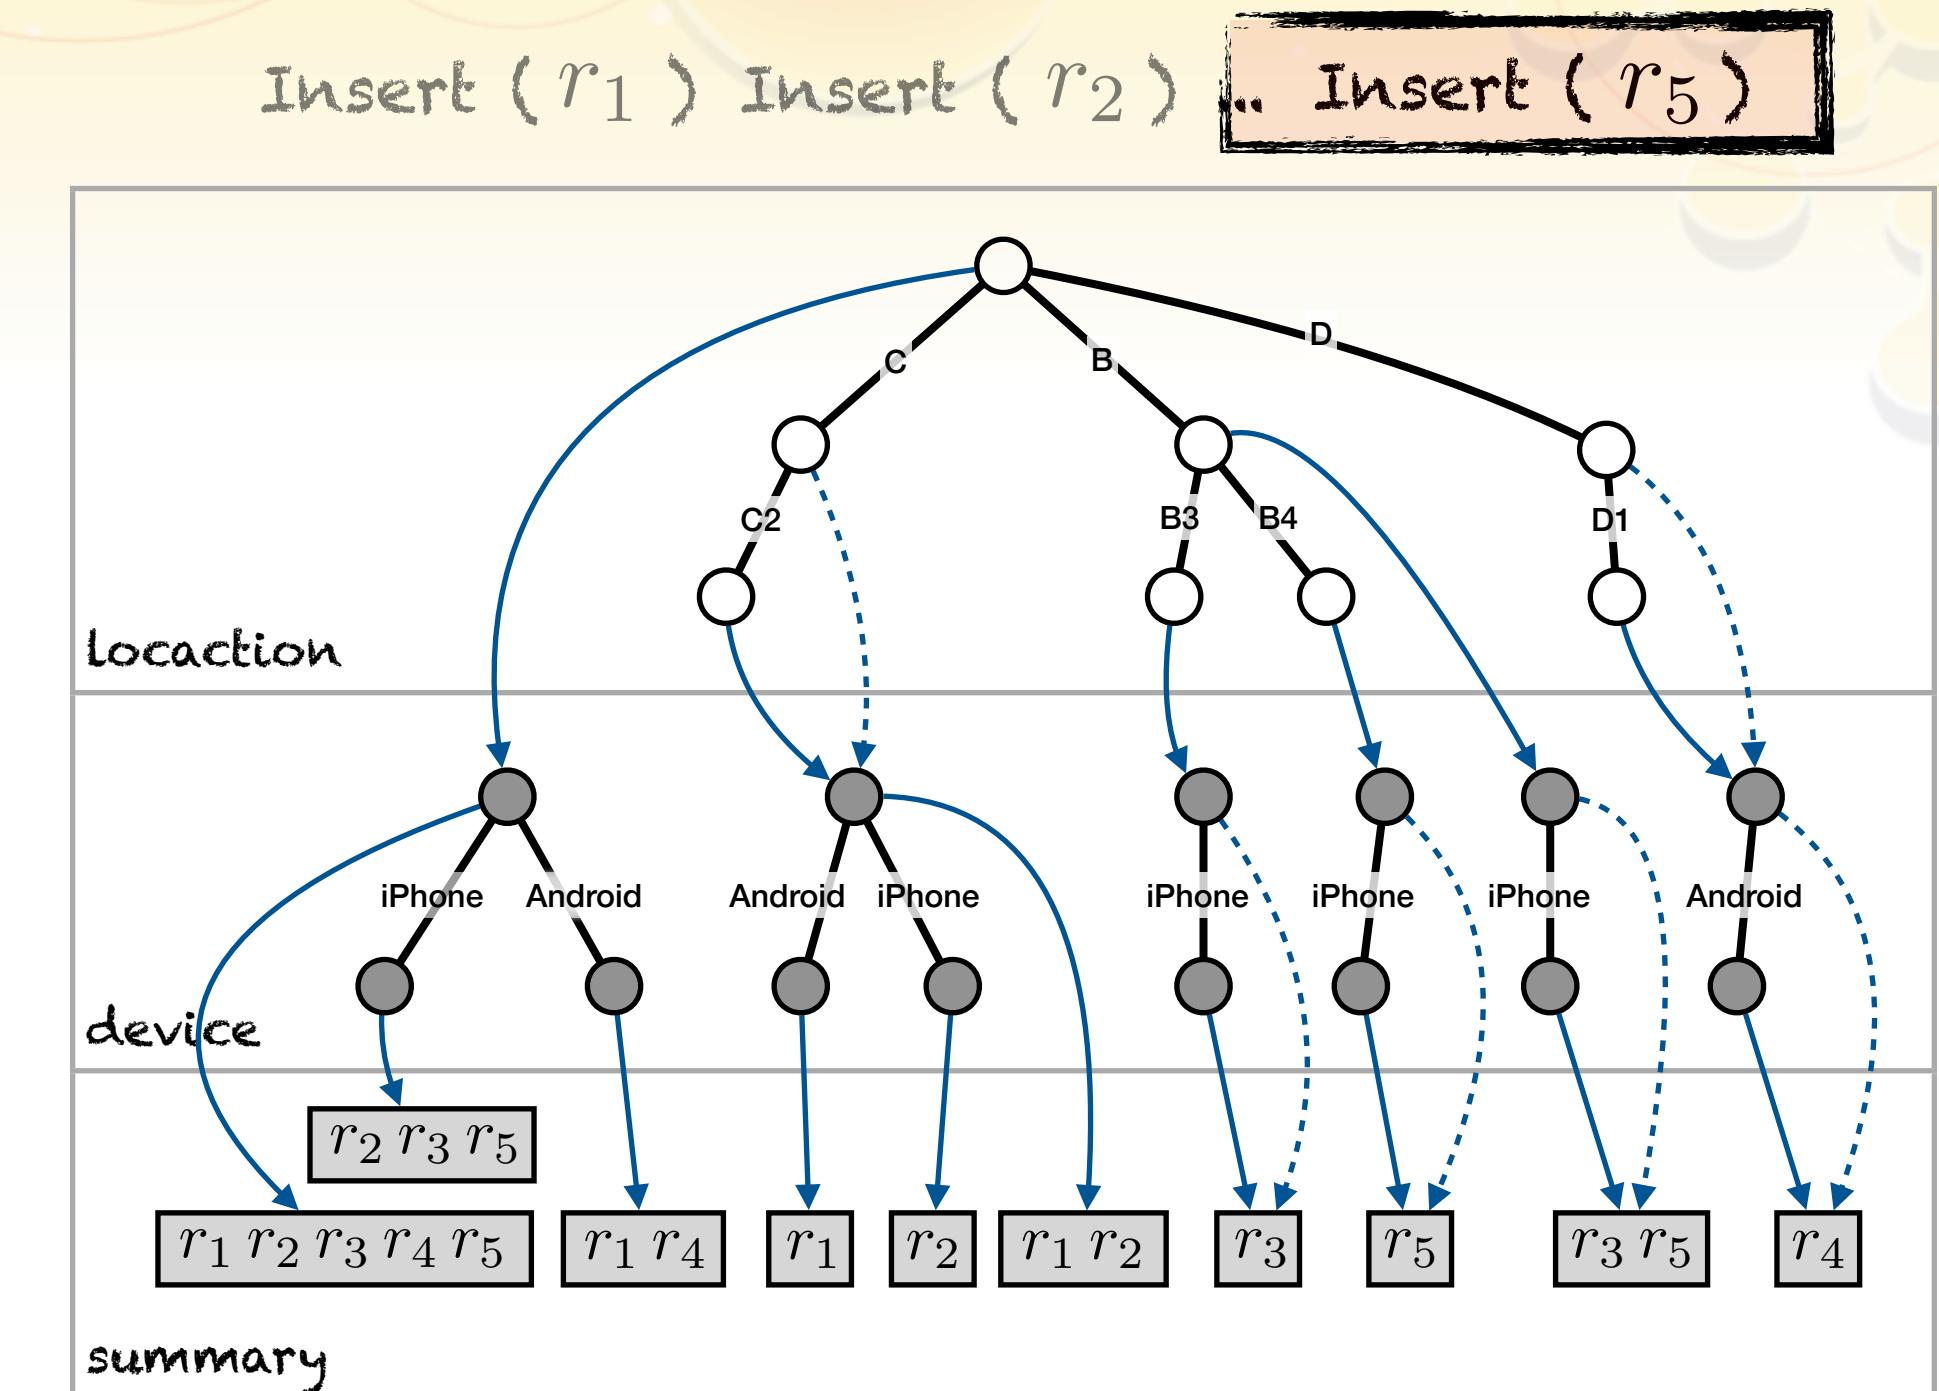
<code>&lt;loc1, loc2&gt;</code>	<code>&lt;device&gt;</code>	time
C, C2	Android	1
C, C2	iPhone	2
B, B3	iPhone	2
D, D1	Android	3
B, B4	iPhone	5



# Building a Nanocube

**Table:**

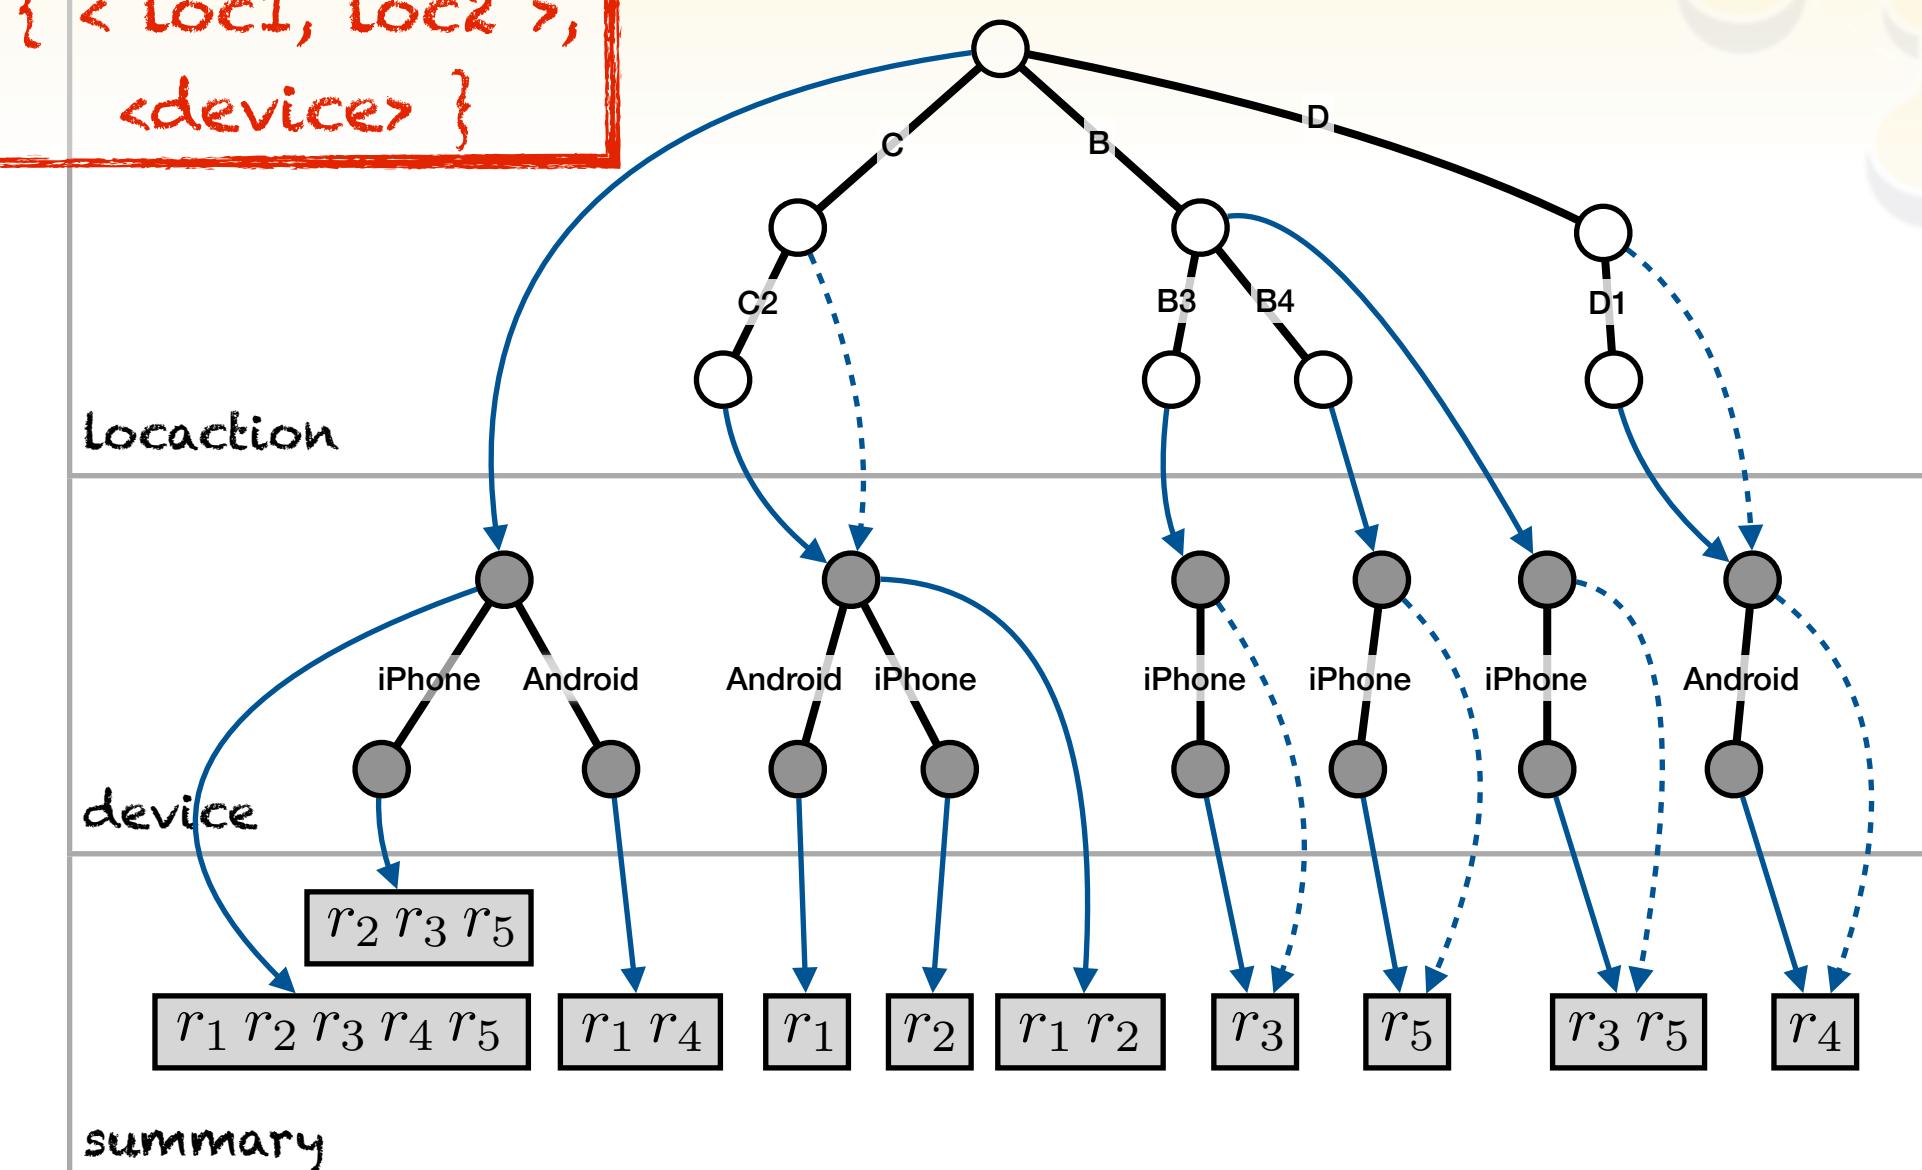
<code>&lt;loc1, loc2&gt;</code>	<code>&lt;device&gt;</code>	time
C, C2	Android	1
C, C2	iPhone	2
B, B3	iPhone	2
D, D1	Android	3
B, B4	iPhone	5



# Building a Nanocube

Loc1	Loc2	device	orig. records
*	*	*	{r1,r2,r3,r4,r5}
B	*	*	{r3,r5}
C	*	*	{r1,r2}
D	*	*	{r4}
*	Android		{r1,r4}
*	iPhone		{r2,r3,r5}
C2	*	*	{r1,r2}
B3	*	*	{r3}
B4	*	*	{r5}
D1	*	*	{r4}
Android			{r1}
iPhone			{r2}
C	*	*	{r3,r5}
C	*	*	{r4}
B	*	*	{r1}
D	*	*	{r2}
C2	*	*	{r1}
C2	*	*	{r2}
B3	*	*	{r3}
B4	*	*	{r5}
D	D1	Android	{r4}

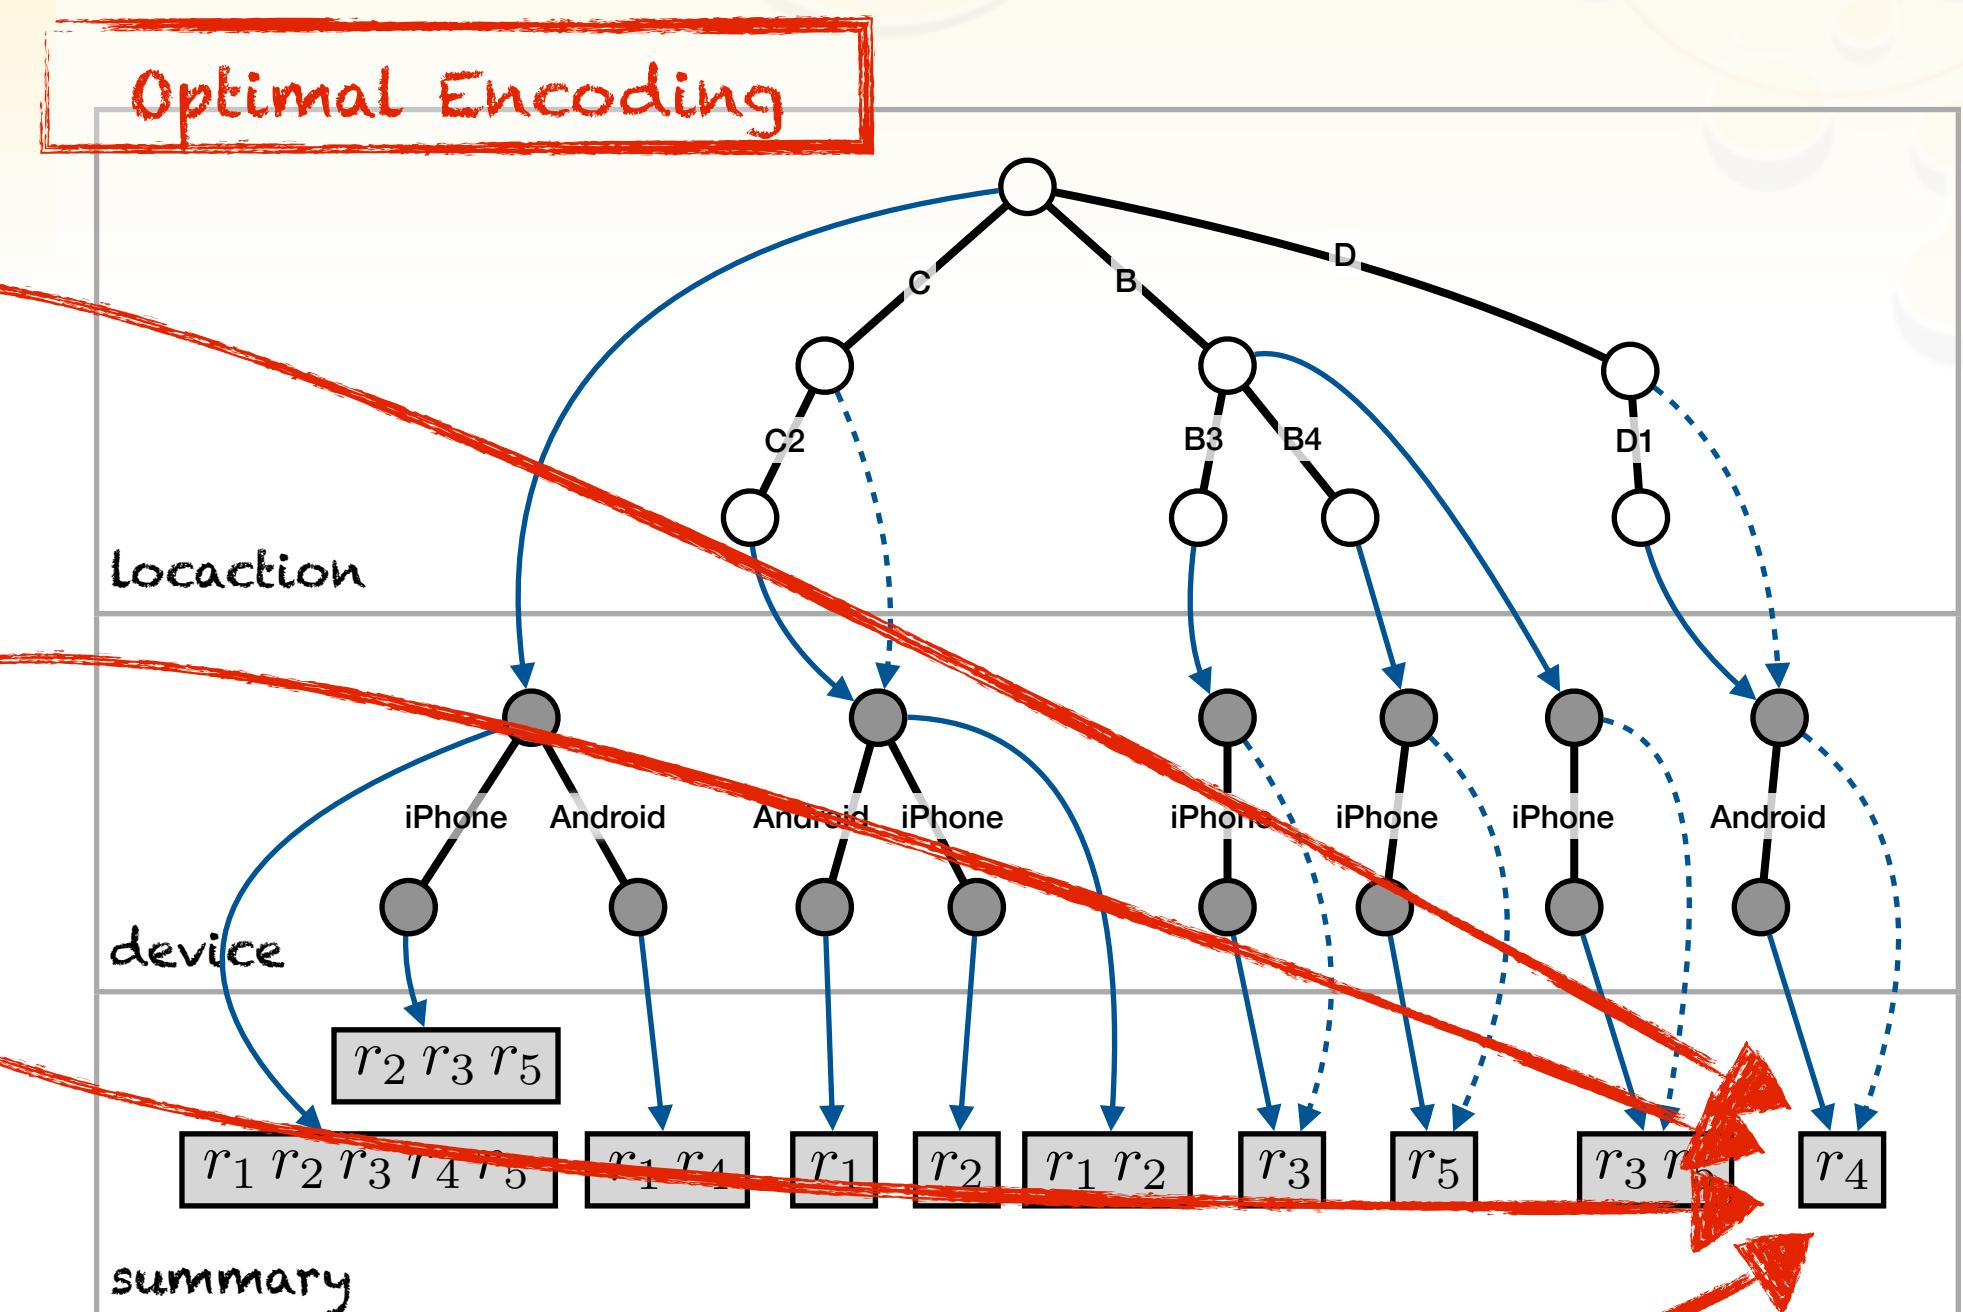
Roll Up Cube  
{ < Loc1, Loc2 >,  
< device > }



Roll Up Cube  
 { < loc1, loc2 >,  
 <device> }

# Building a Nanocube

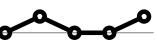
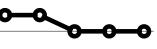
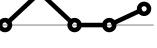
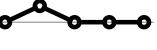
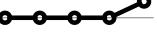
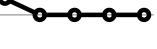
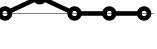
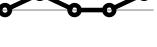
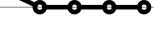
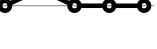
Loc1	Loc2	device	orig. records
*	*	*	{r1,r2,r3,r4,r5}
B	*	*	{r3,r5}
C	*	*	{r1,r2}
D	*	*	{r4}
*	*	Android	{r1,r4}
*	*	iPhone	{r2,r3,r5}
C2	*	*	{r1,r2}
B3	*	*	{r3}
B4	*	*	{r5}
D1	*	*	{r4}
Android	*	Android	{r1}
iPhone	*	iPhone	{r2}
iPhone	*	iPhone	{r3}
Android	*	Android	{r5}
Android	*	Android	{r4}
C2	*	Android	{r1}
C2	*	iPhone	{r2}
C2	*	iPhone	{r3}
B3	*	iPhone	{r4}
B4	*	iPhone	{r5}
D1	*	Android	{r4}



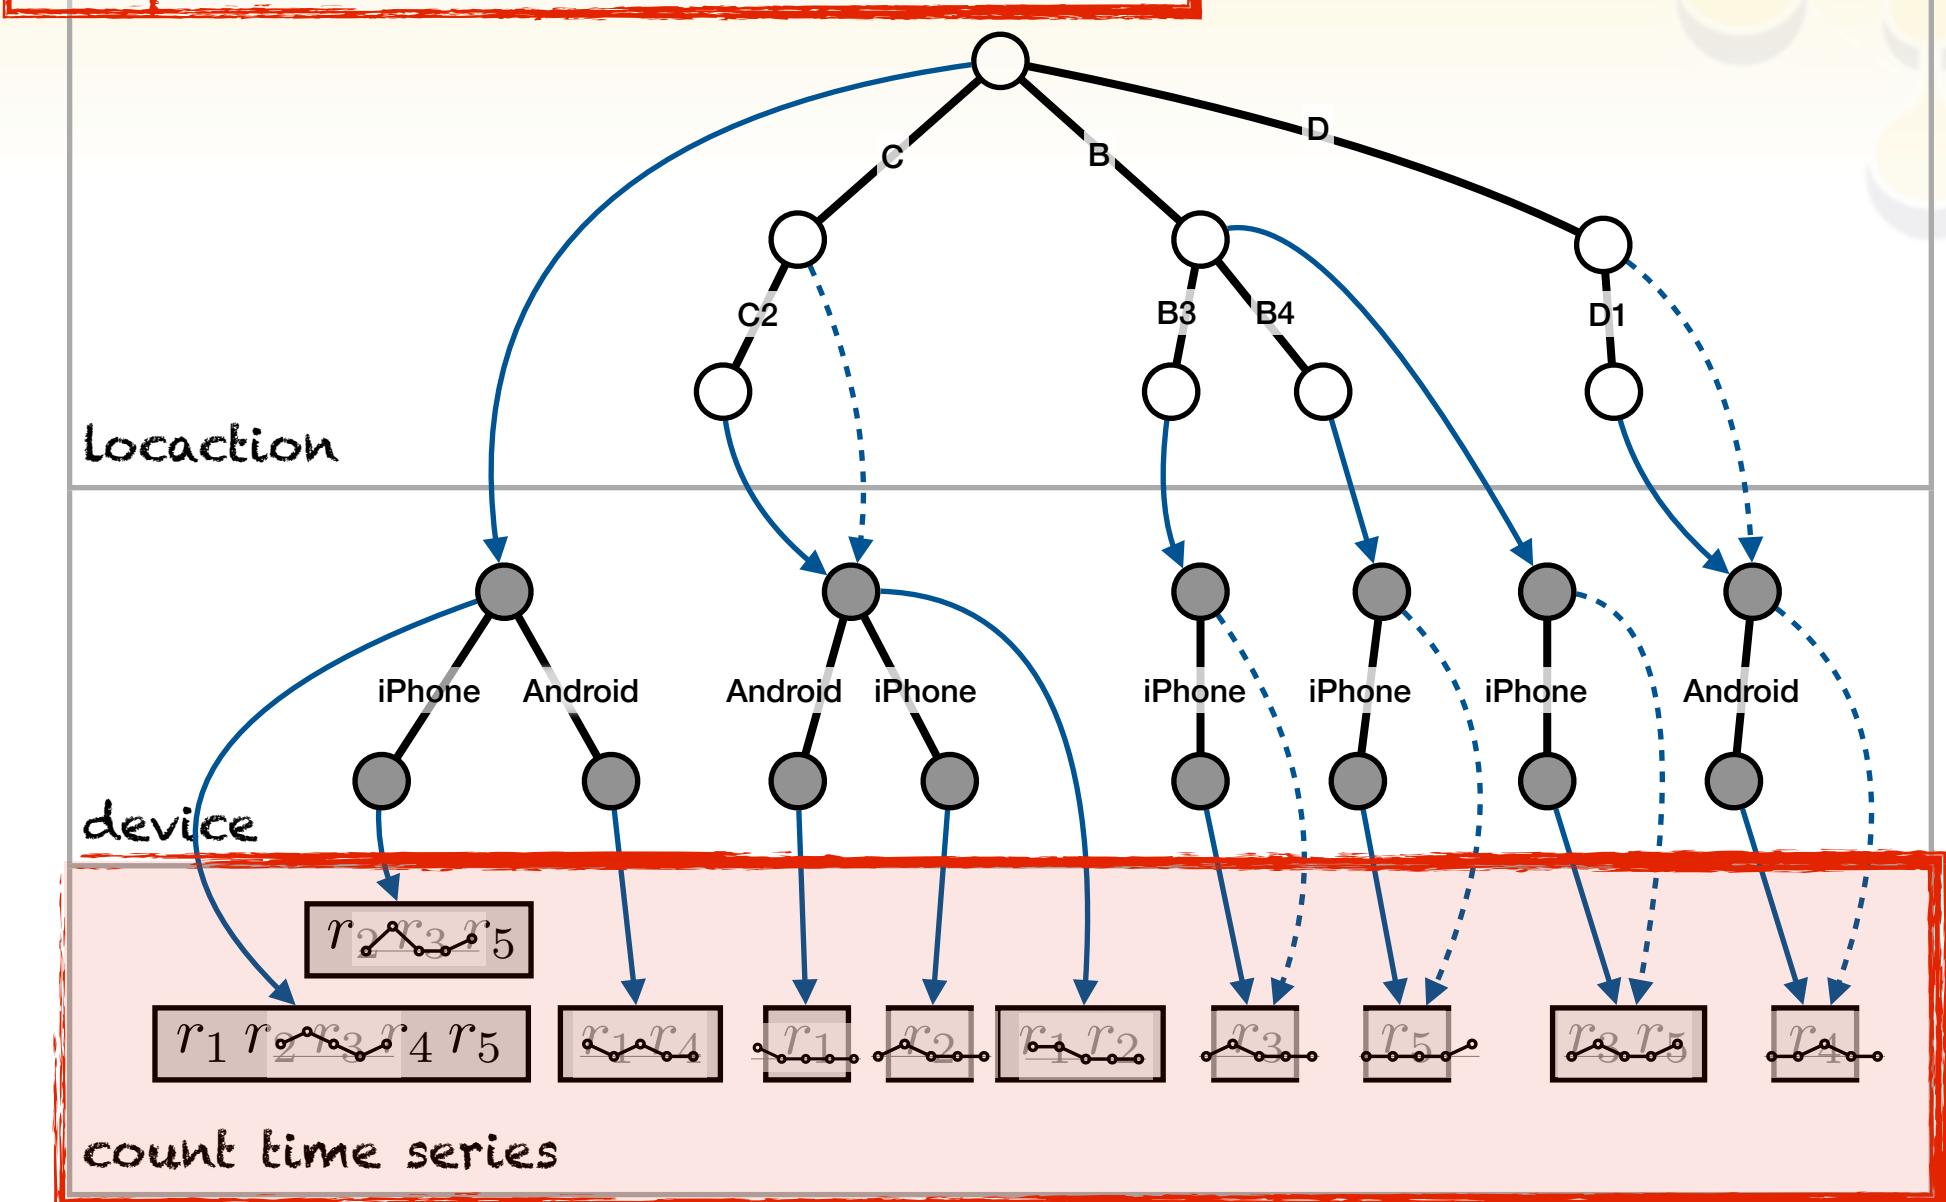
# Roll Up Cube

```
{ <loc1, loc2>,  
  <device> }
```

# Building a Nanocube

loc1	loc2	device	orig. records
*	*	*	
B	*	*	
C	*	*	
D	*	*	
*	*	Android	
*	*	iPhone	
C	C2	*	
B	B3	*	
B	B4	*	
D	D1	*	
C	*	Android	
C	*	iPhone	
B	*	iPhone	
D	*	Android	
C	C2	Android	
C	C2	iPhone	
B	B3	iPhone	
B	B4	iPhone	
D	D1	Android	

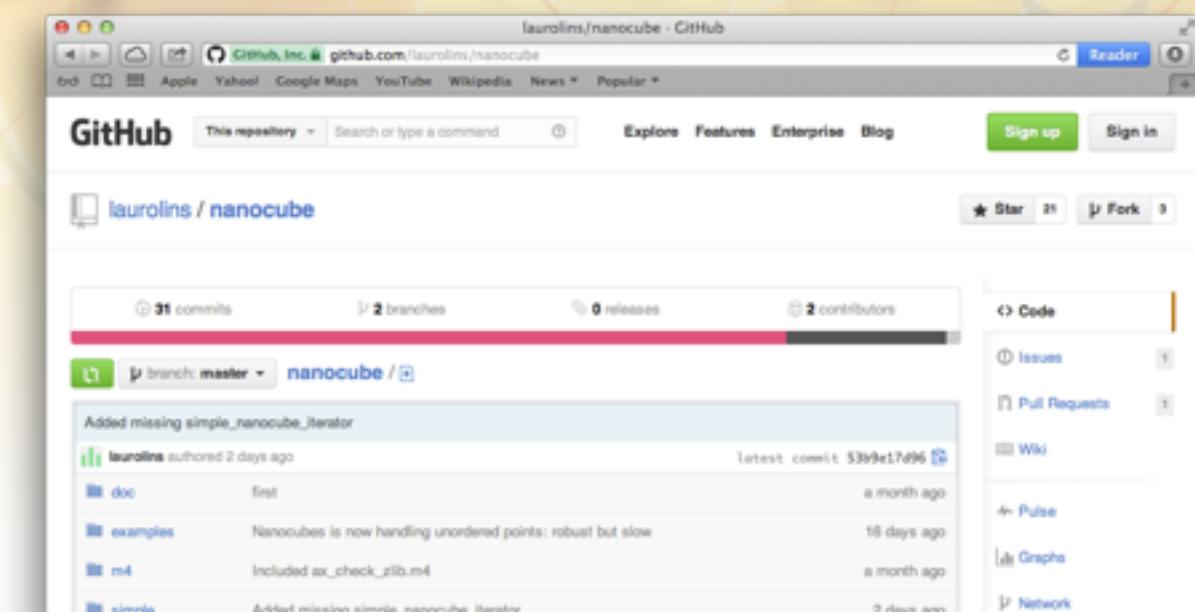
Time Series are encoded using  
sparse summed area tables



# Implementation

- Server - Back End
- HTTP Server written in C++11
- mongoose (http)
- templates and boost::mpl
- tagged pointers
- **open source**

<http://github.com/laurolins/nanocube>



<http://nanocubes.net>



# Results

dataset	n	memory	time	keys	cardinality	schema
brightkite	4.5 M	1.6 GB	3.50 m	3.5 M	$2^{74}$	spatial(50), time(16), weekday(3), hour(5)
customer tix	7.8 M	2.5 GB	8.47 m	7.8 M	$2^{69}$	spatial(50), time(16), type(3)
flights	121 M	2.3 GB	31.13 m	43.3 M	$2^{75}$	spatial(50), time(16), carrier(5), delay(4)
twitter-small	210 M	10.2 GB	1.23 h	116.0 M	$2^{53}$	spatial(34), time(16), device(3)
twitter	210 M	46.4 GB	5.87 h	136.0 M	$2^{60}$	spatial(34), time(16), lang(5), device(3), app(2)
splom-10	1 B	4.3 MB	4.13 h	7.4 K	$2^{20}$	d1(4), d2(4), d3(4), d4(4), d5(4)
splom-50	1 B	166 MB	4.72 h	1.9 M	$2^{30}$	d1(6), d2(6), d3(6), d4(6), d5(6)
cdrs	1 B	3.6 GB	3.08 h	96.3 M	$2^{69}$	spatial(50), time(16), duration(3)

every nanocube  
size  $\leq 11\text{GB RAM}$   
except for one

## Results

dataset	n	memory	time	keys	cardinality	schema
brightkite	4.5 M	1.6 GB	3.50 m	3.5 M	$2^{74}$	spatial(50), time(16), weekday(3), hour(5)
customer tix	7.8 M	2.5 GB	8.47 m	7.8 M	$2^{69}$	spatial(50), time(16), type(3)
flights	121 M	2.3 GB	31.13 m	43.3 M	$2^{75}$	spatial(50), time(16), carrier(5), delay(4)
twitter-small	210 M	10.2 GB	1.23 h	116.0 M	$2^{53}$	spatial(34), time(16), device(3)
twitter	210 M	46.4 GB	5.87 h	136.0 M	$2^{60}$	spatial(34), time(16), lang(5), device(3), app(2)
splom-10	1 B	4.3 MB	4.13 h	7.4 K	$2^{20}$	d1(4), d2(4), d3(4), d4(4), d5(4)
splom-50	1 B	166 MB	4.72 h	1.9 M	$2^{30}$	d1(6), d2(6), d3(6), d4(6), d5(6)
cdrs	1 B	3.6 GB	3.08 h	96.3 M	$2^{69}$	spatial(50), time(16), duration(3)

## Limitations

- Summed-area table trick can only be used on the last dimension
- In the worst case, size of a nanocube grows exponentially with the number of new dimensions

## Future and ongoing work

- Multiple spatial dimensions
- Nanocubes are insert-only
- Horizontal Scaling (Parallelism)
- Approximate Nanocubes

## Conclusions

- Nanocubes provide spatiotemporal in-memory data cubes for large datasets
- they enable Interactive Visual Interfaces for datasets much larger than previously possible
- You can explore big data with your laptop, tablet and phone!

## Acknowledgements:

Stephen North

Divesh Srivastava

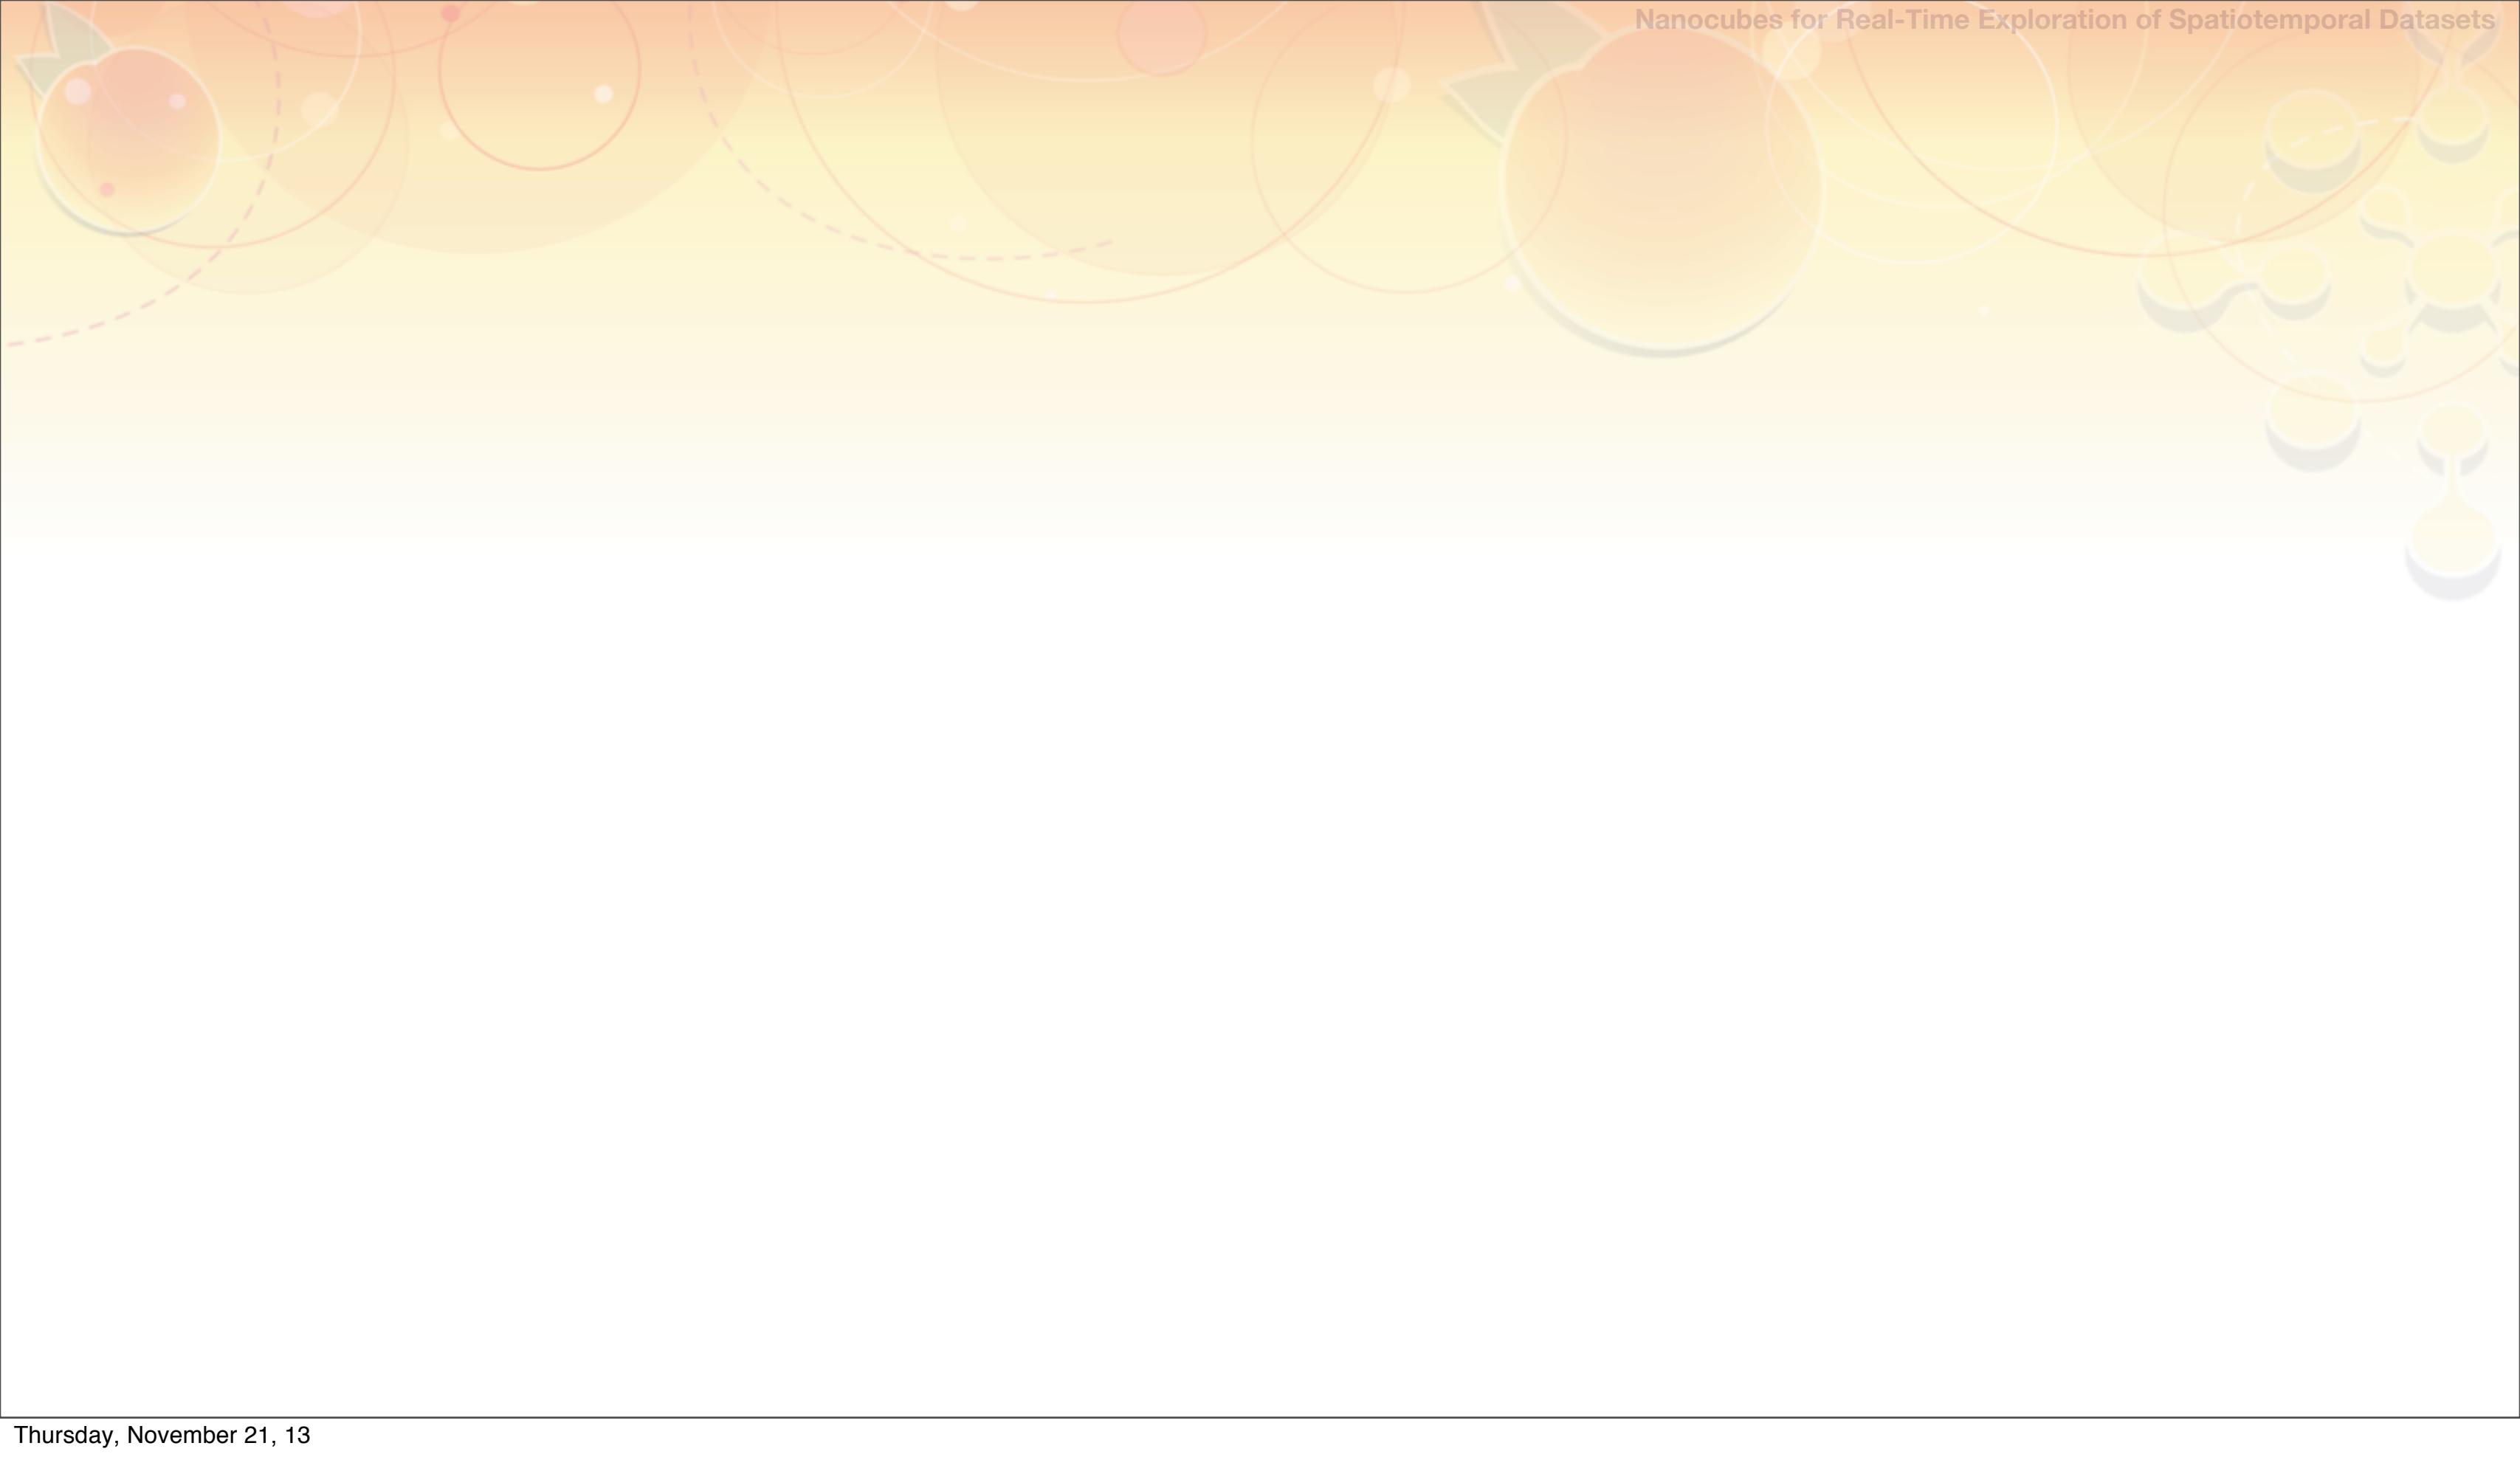
Luciano Barbosa

Thank you - questions?

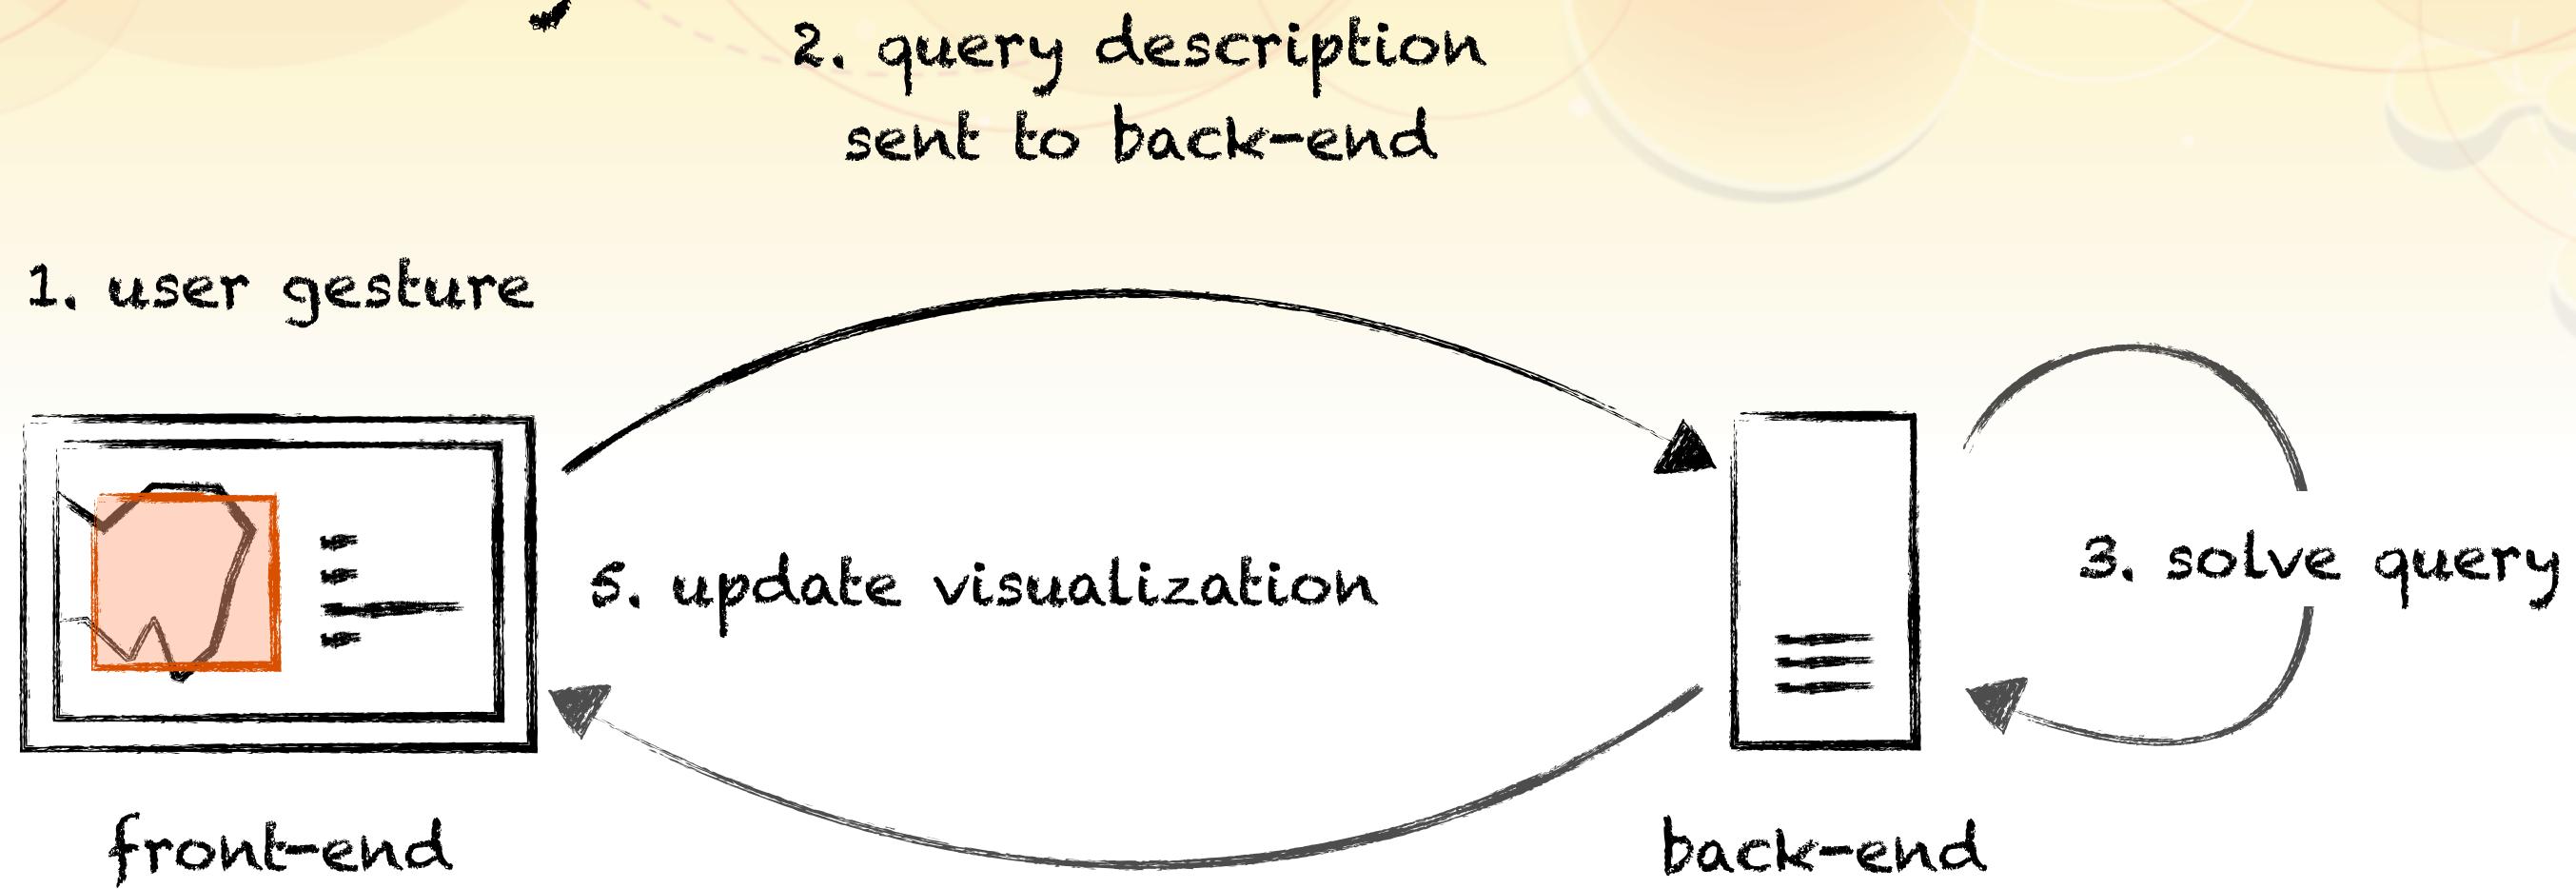
<http://nanocubes.net>

llins@research.att.com

<http://github.com/Laurolins/nanocube>

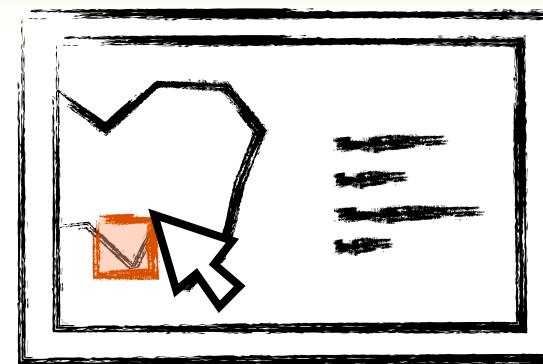


# Interaction Cycle

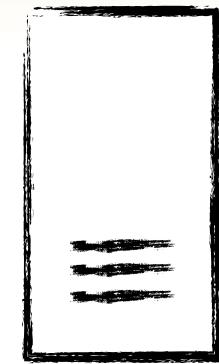


$$t_2 + t_3 + t_4 + t_5 \leq \begin{cases} 33.3 \text{ ms } (30 \text{ fps}) \\ 66.6 \text{ ms } (15 \text{ fps}) \\ 100 \text{ ms } (10 \text{ fps}) \end{cases}$$

# Interaction Cycle



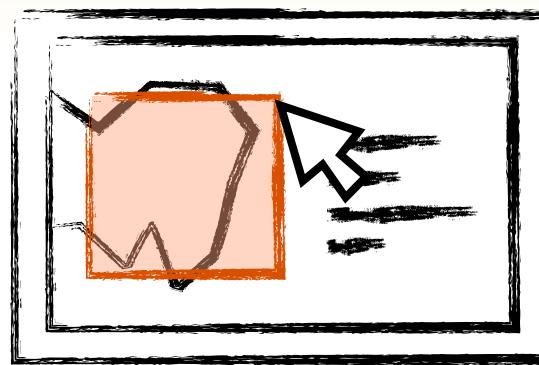
front-end



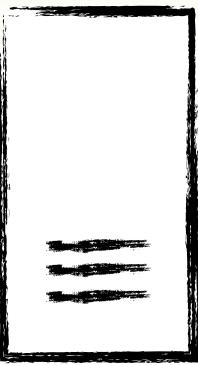
back-end

# Interaction Cycle

1. user gesture

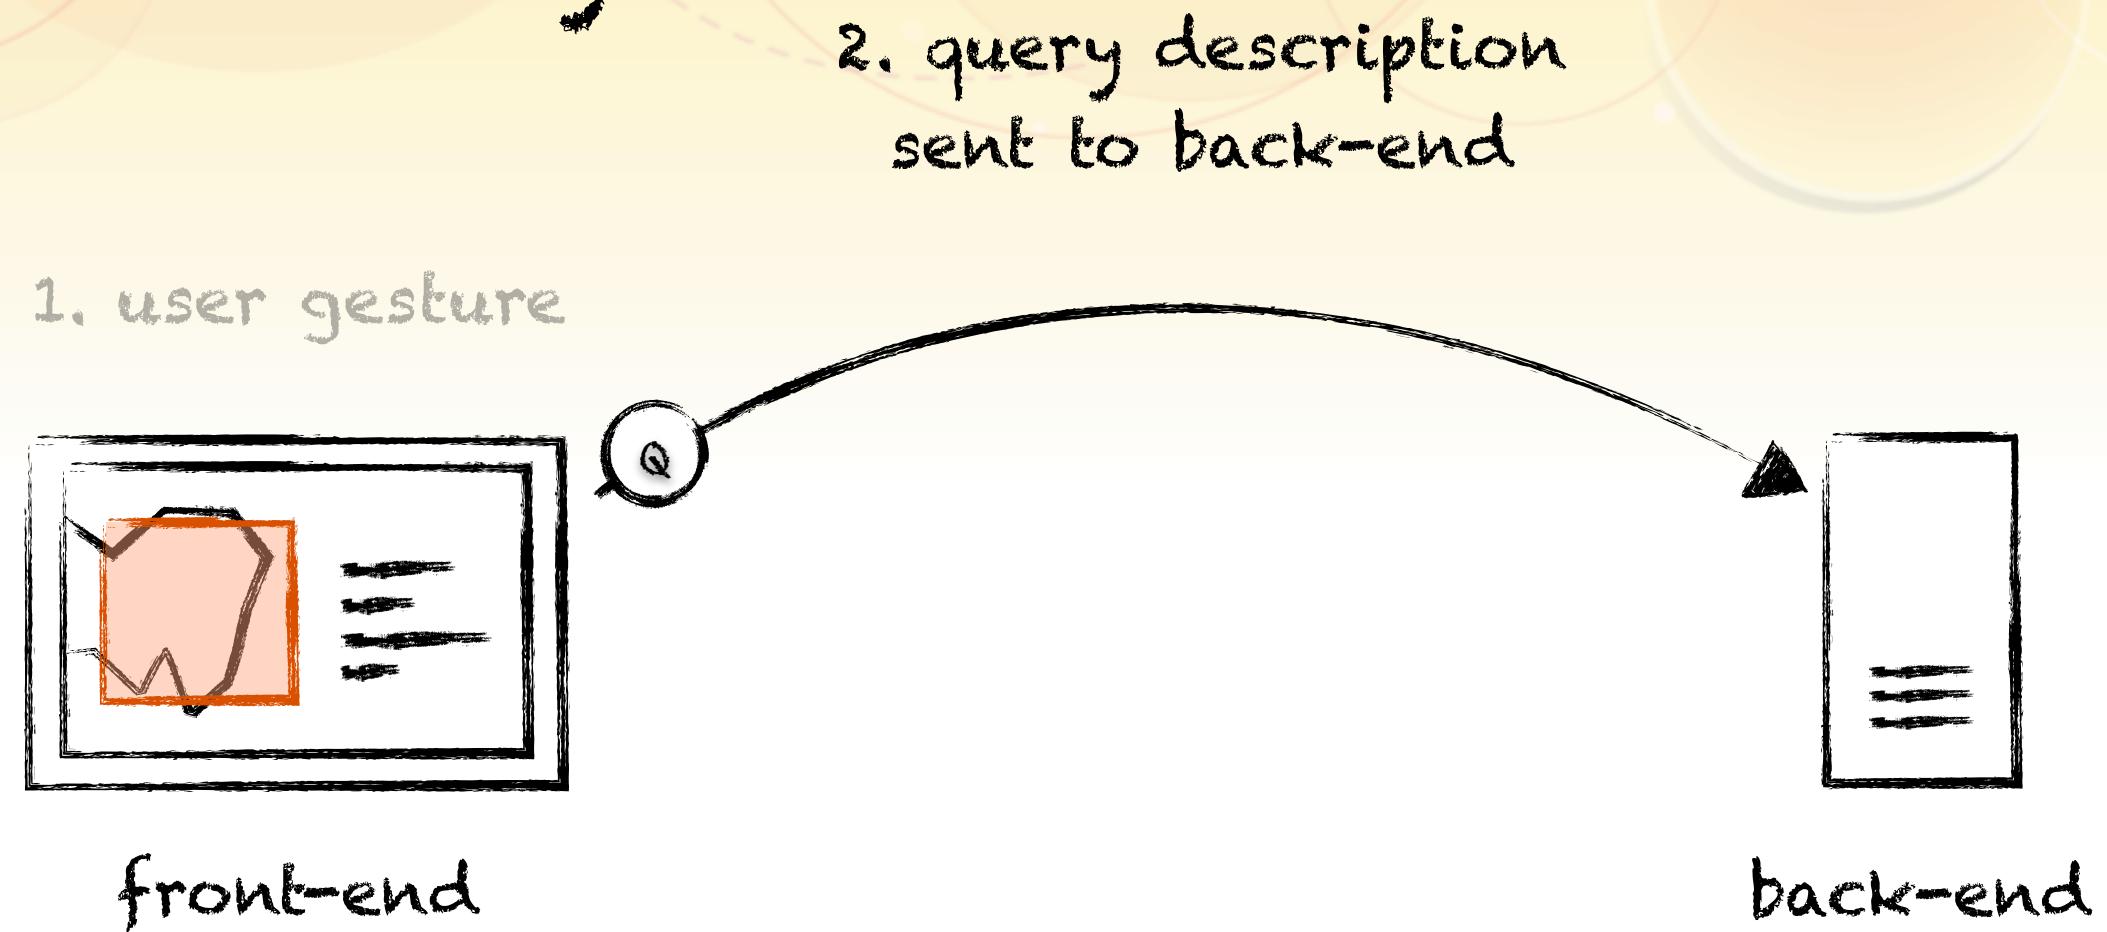


front-end



back-end

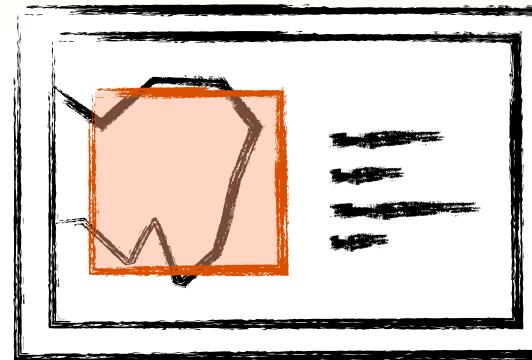
# Interaction Cycle



# Interaction Cycle

2. query description sent to back-end

1. user gesture

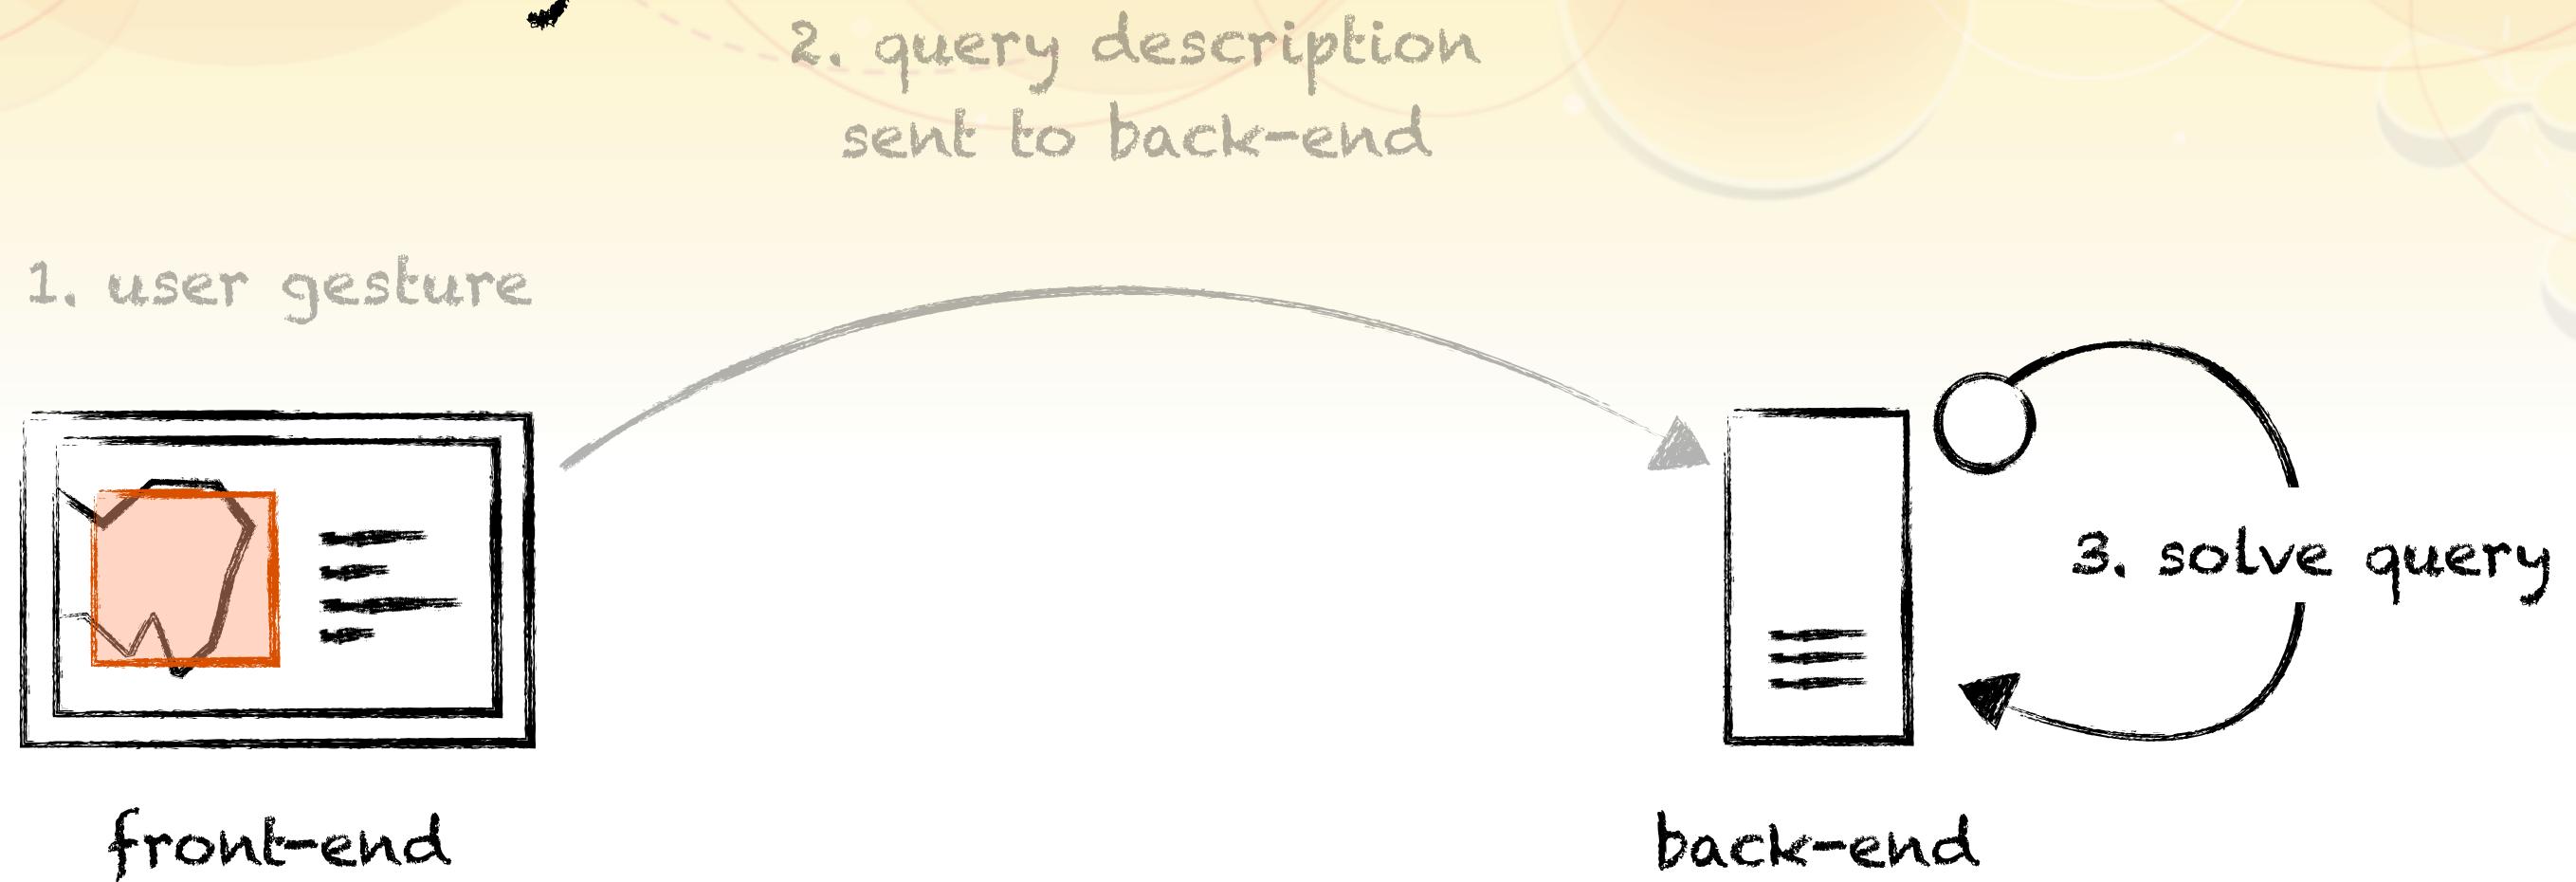


front-end

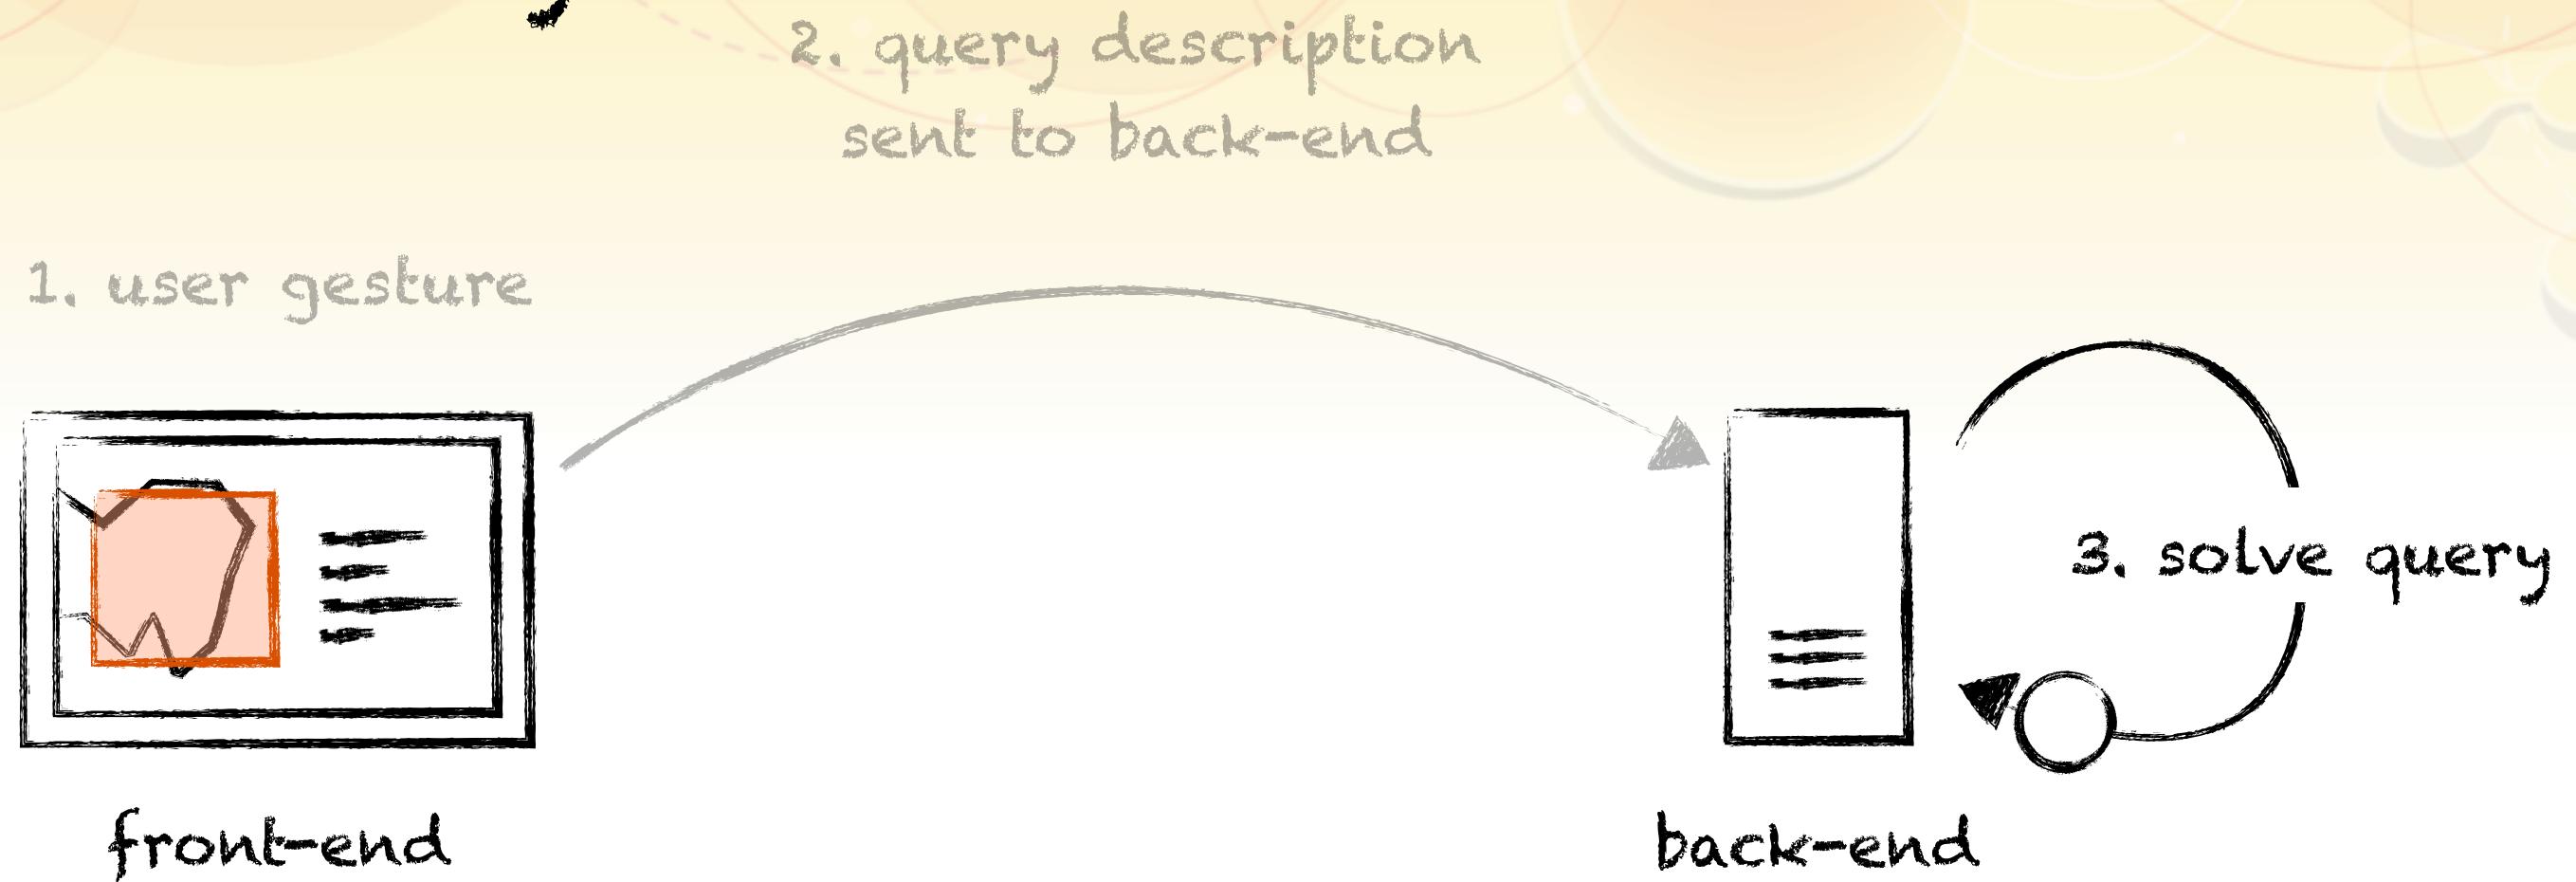


back-end

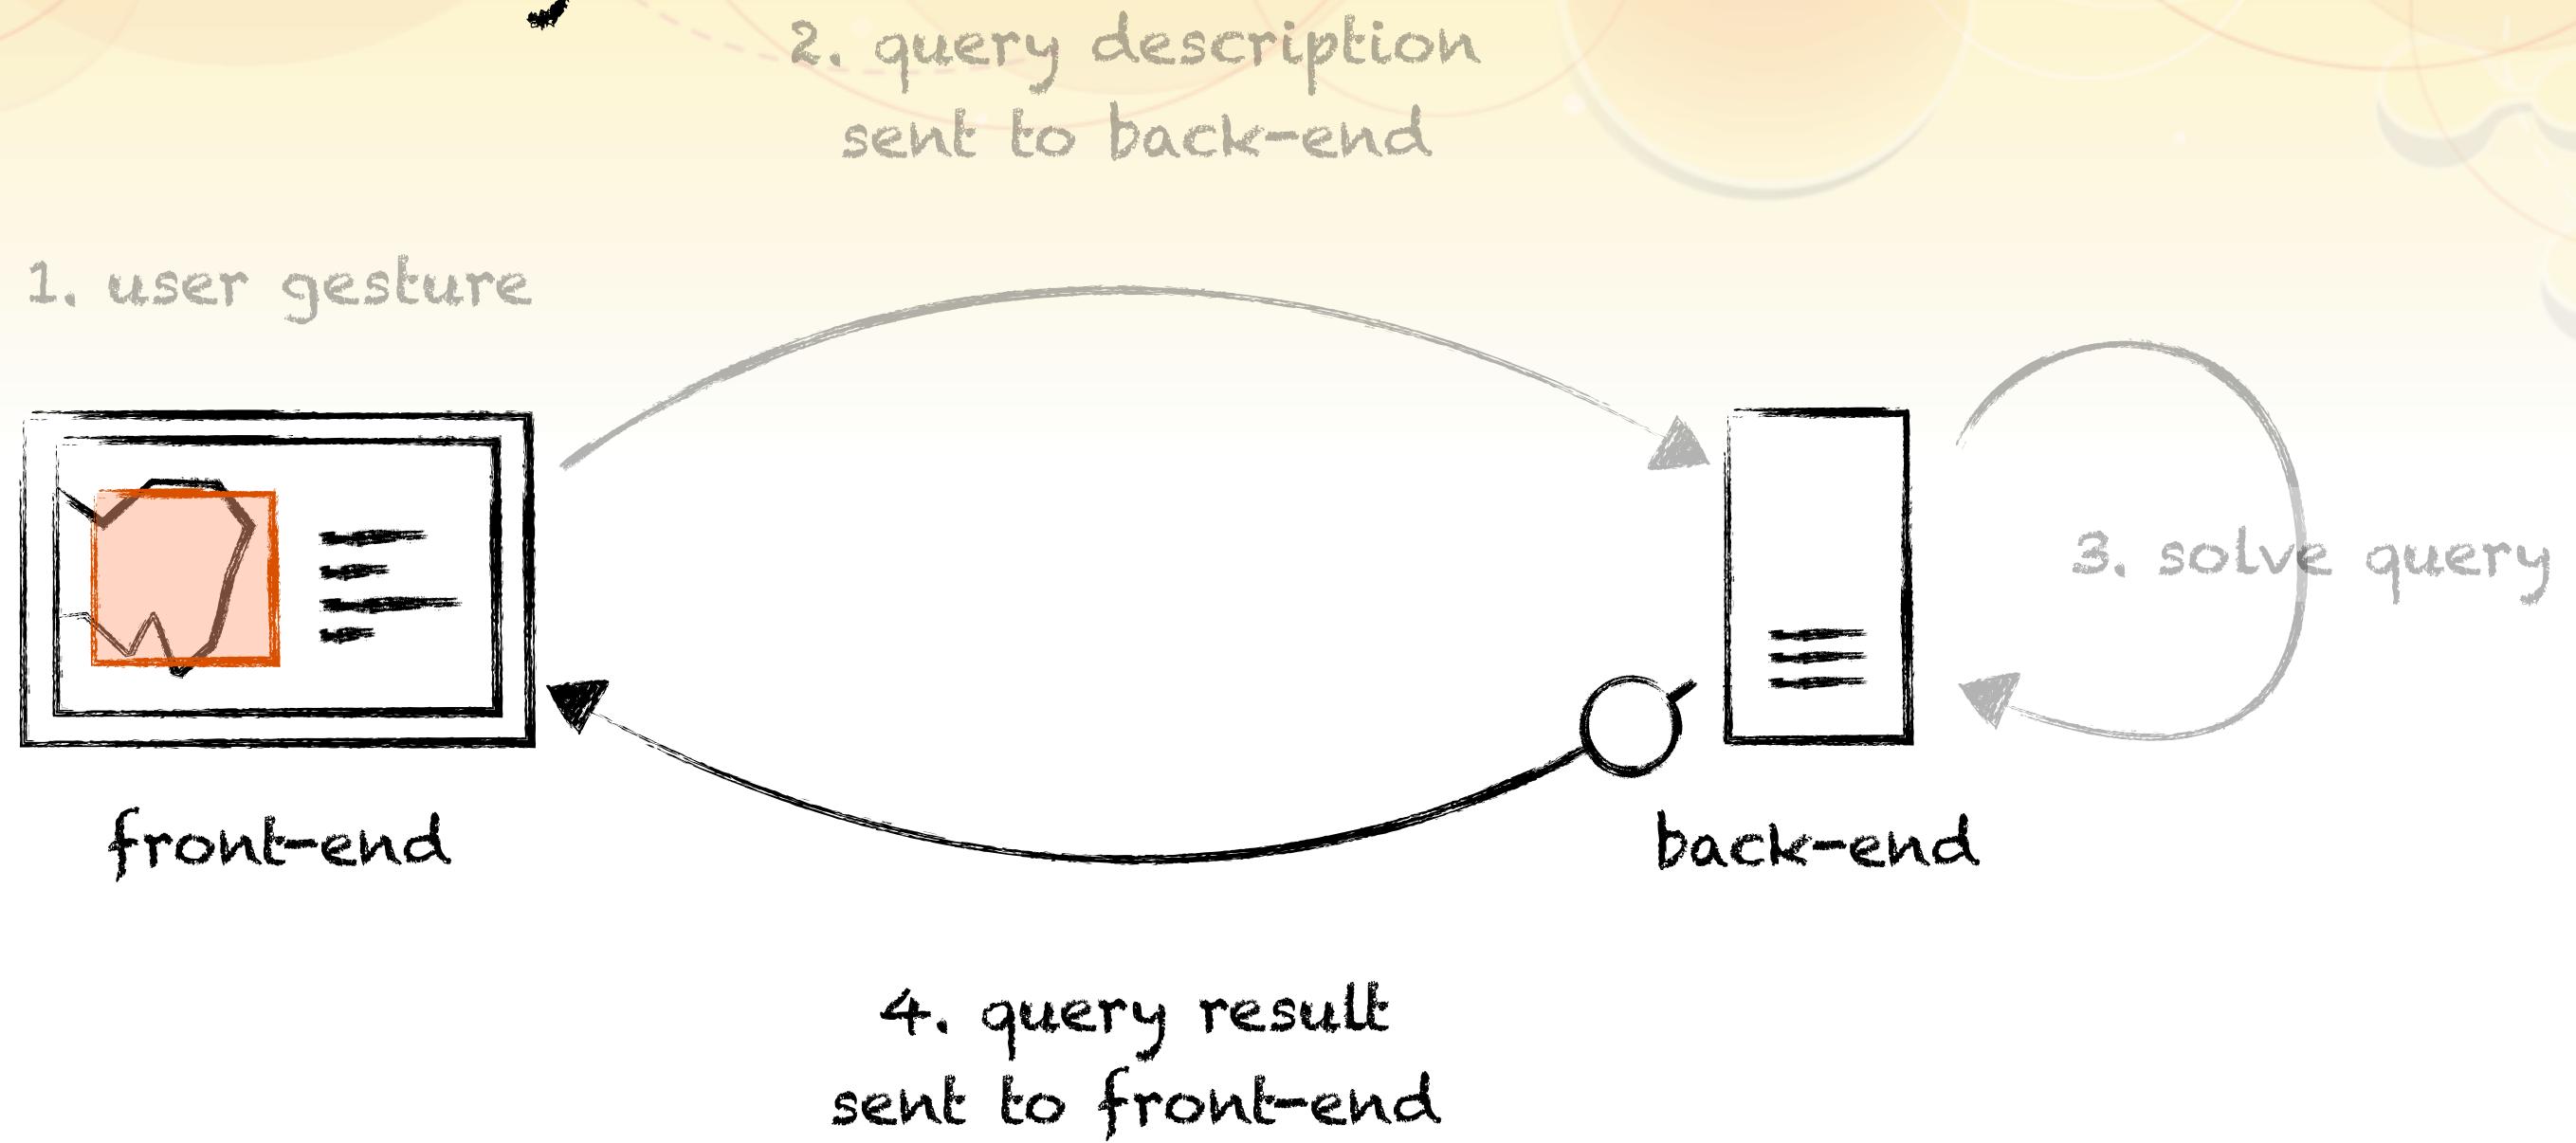
# Interaction Cycle



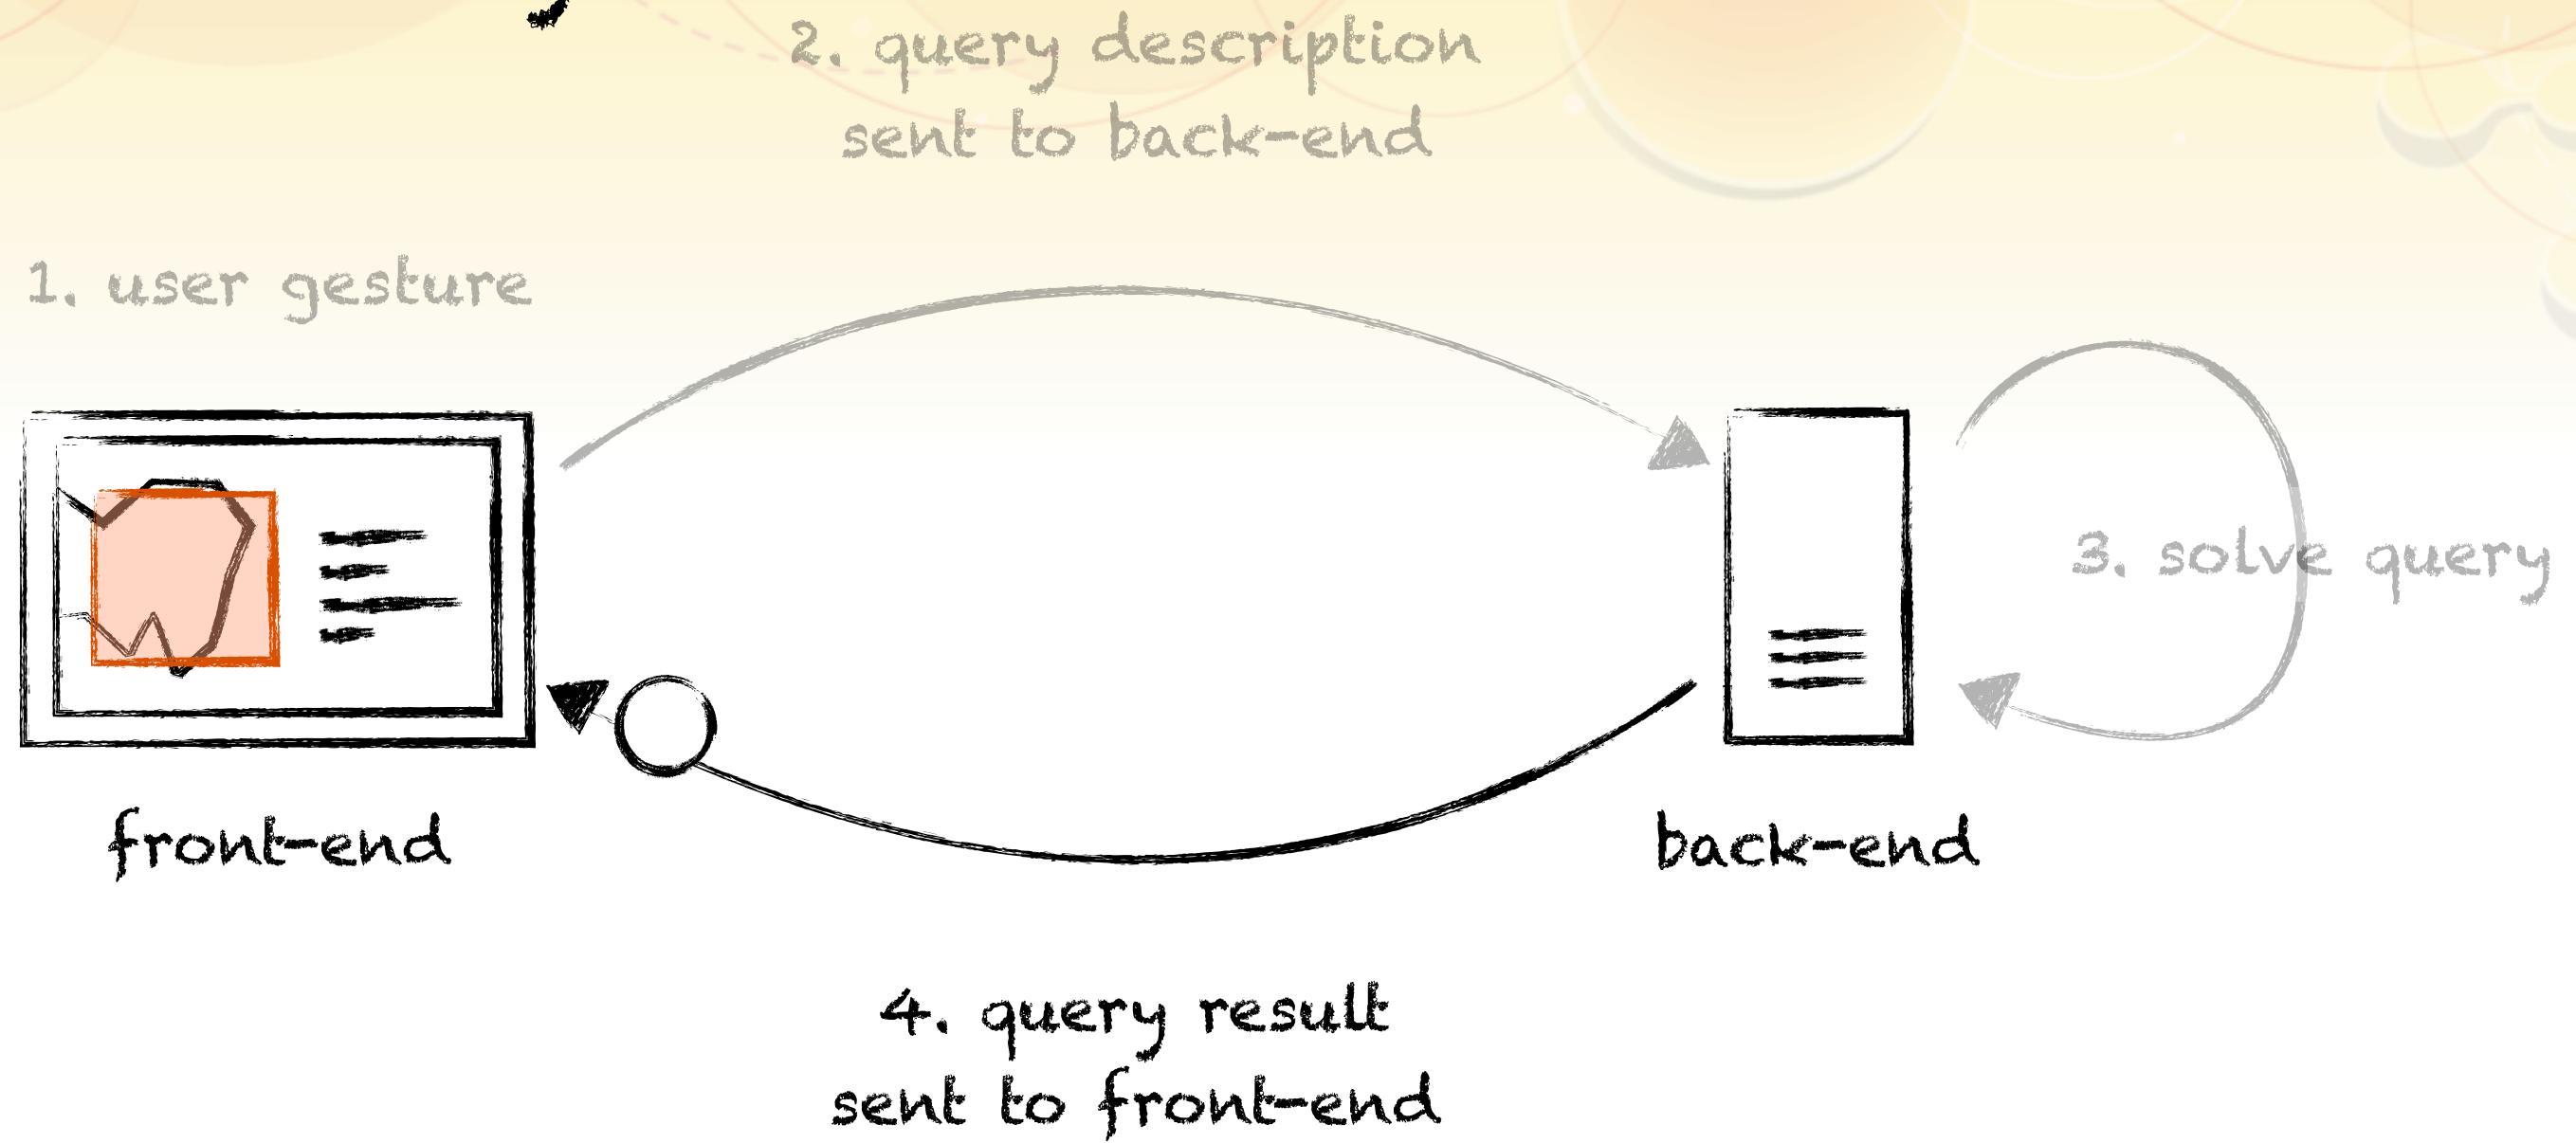
# Interaction Cycle



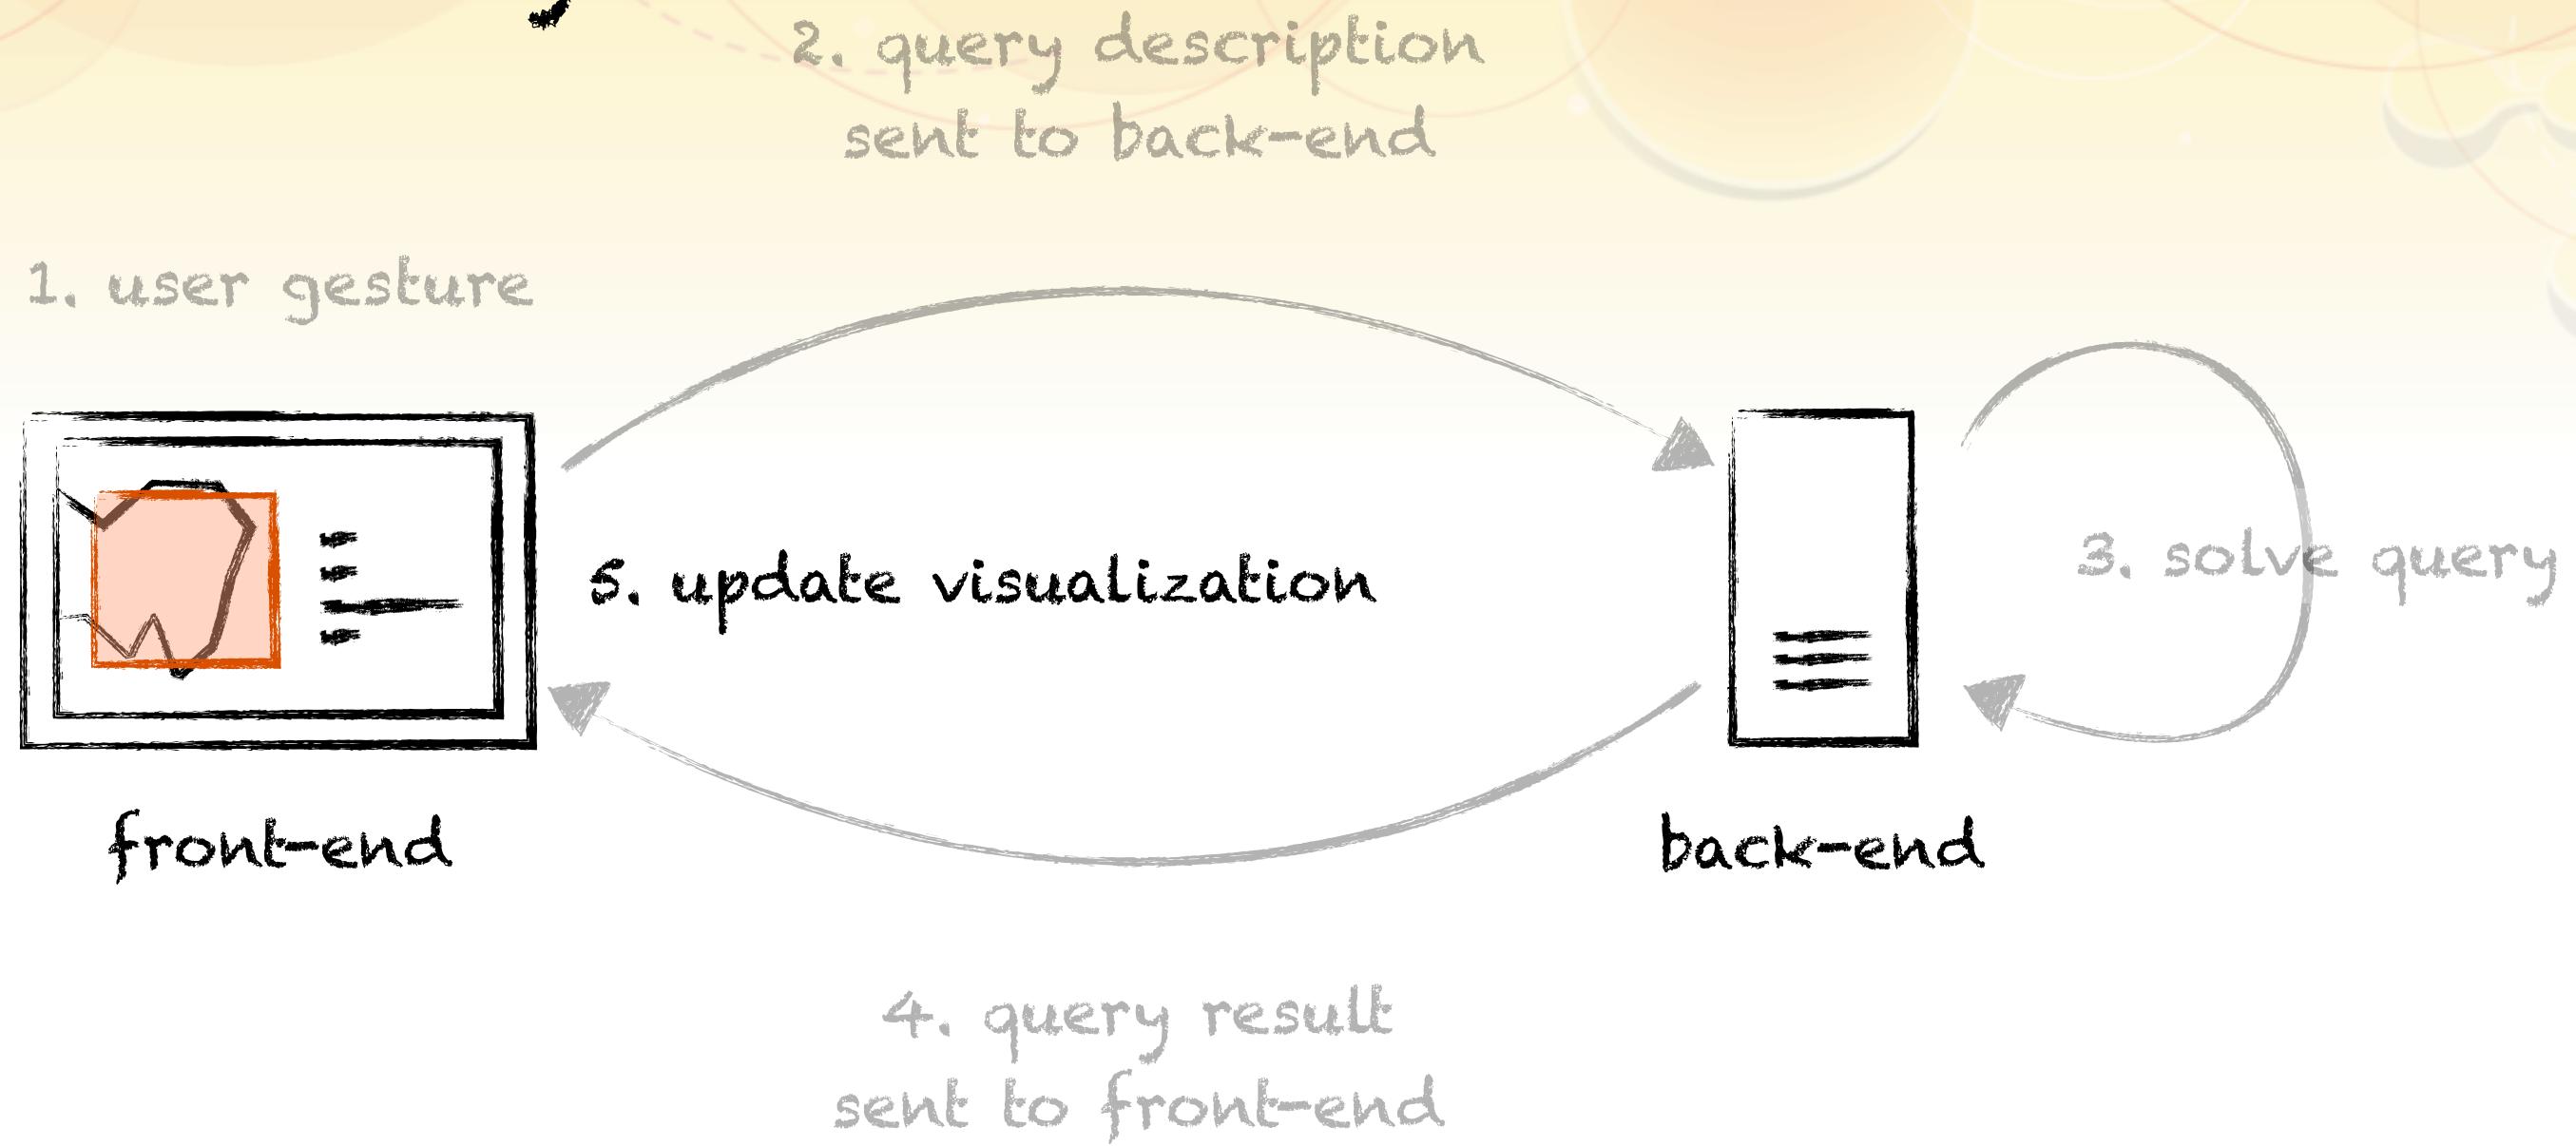
# Interaction Cycle



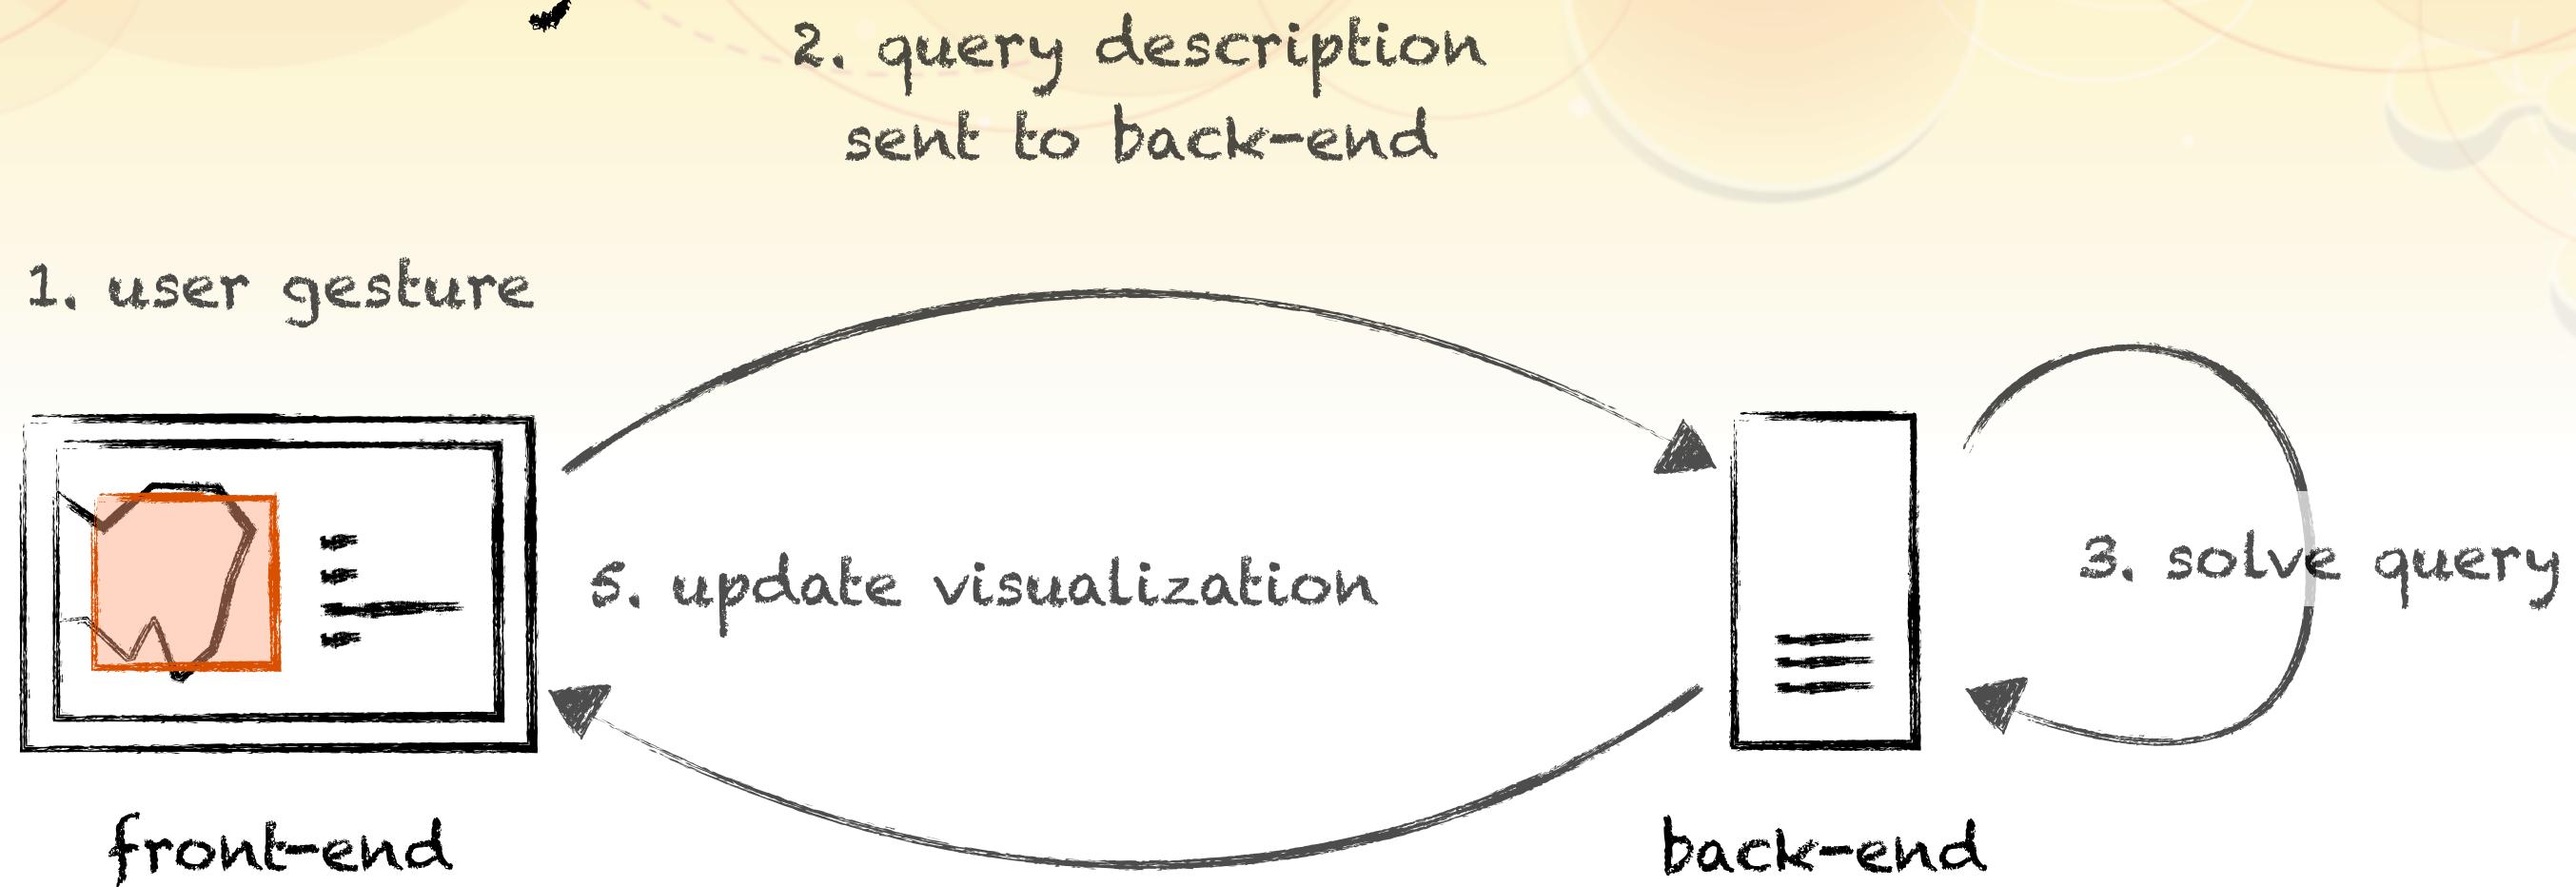
# Interaction Cycle



# Interaction Cycle

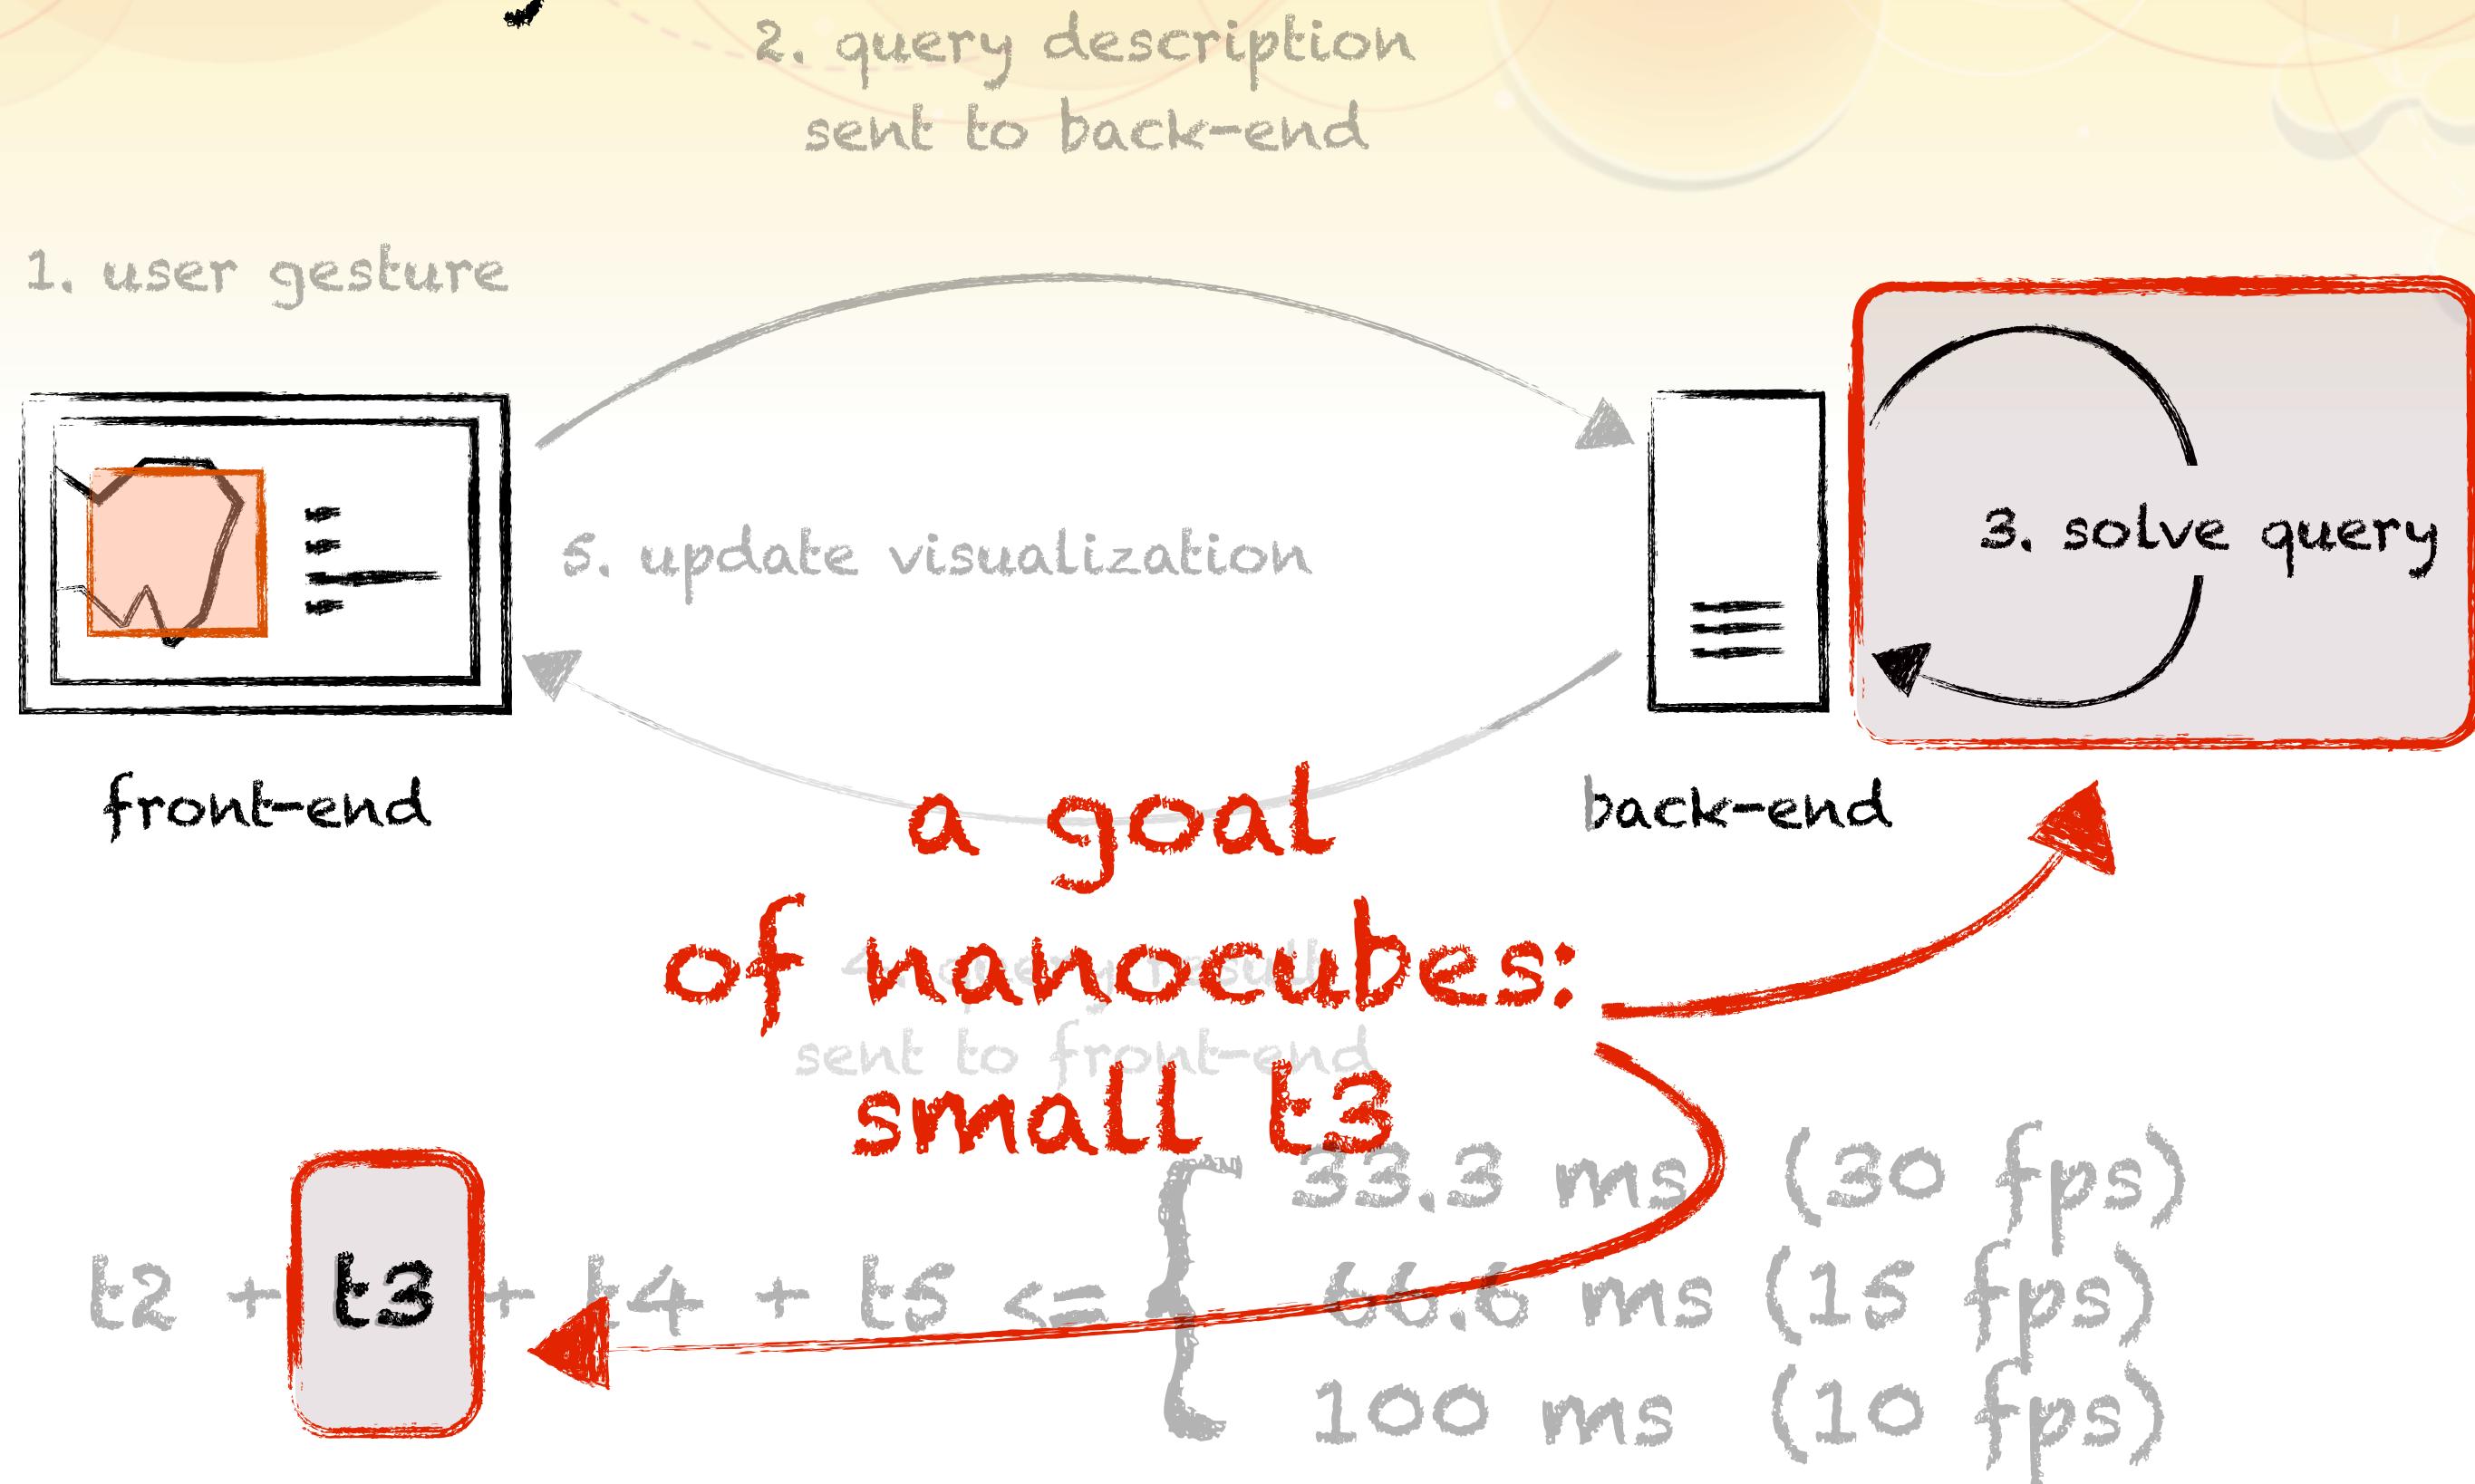


# Interaction Cycle



$$t_2 + t_3 + t_4 + t_5 \leq \begin{cases} 33.3 \text{ ms } (30 \text{ fps}) \\ 66.6 \text{ ms } (15 \text{ fps}) \\ 100 \text{ ms } (10 \text{ fps}) \end{cases}$$

# Interaction Cycle

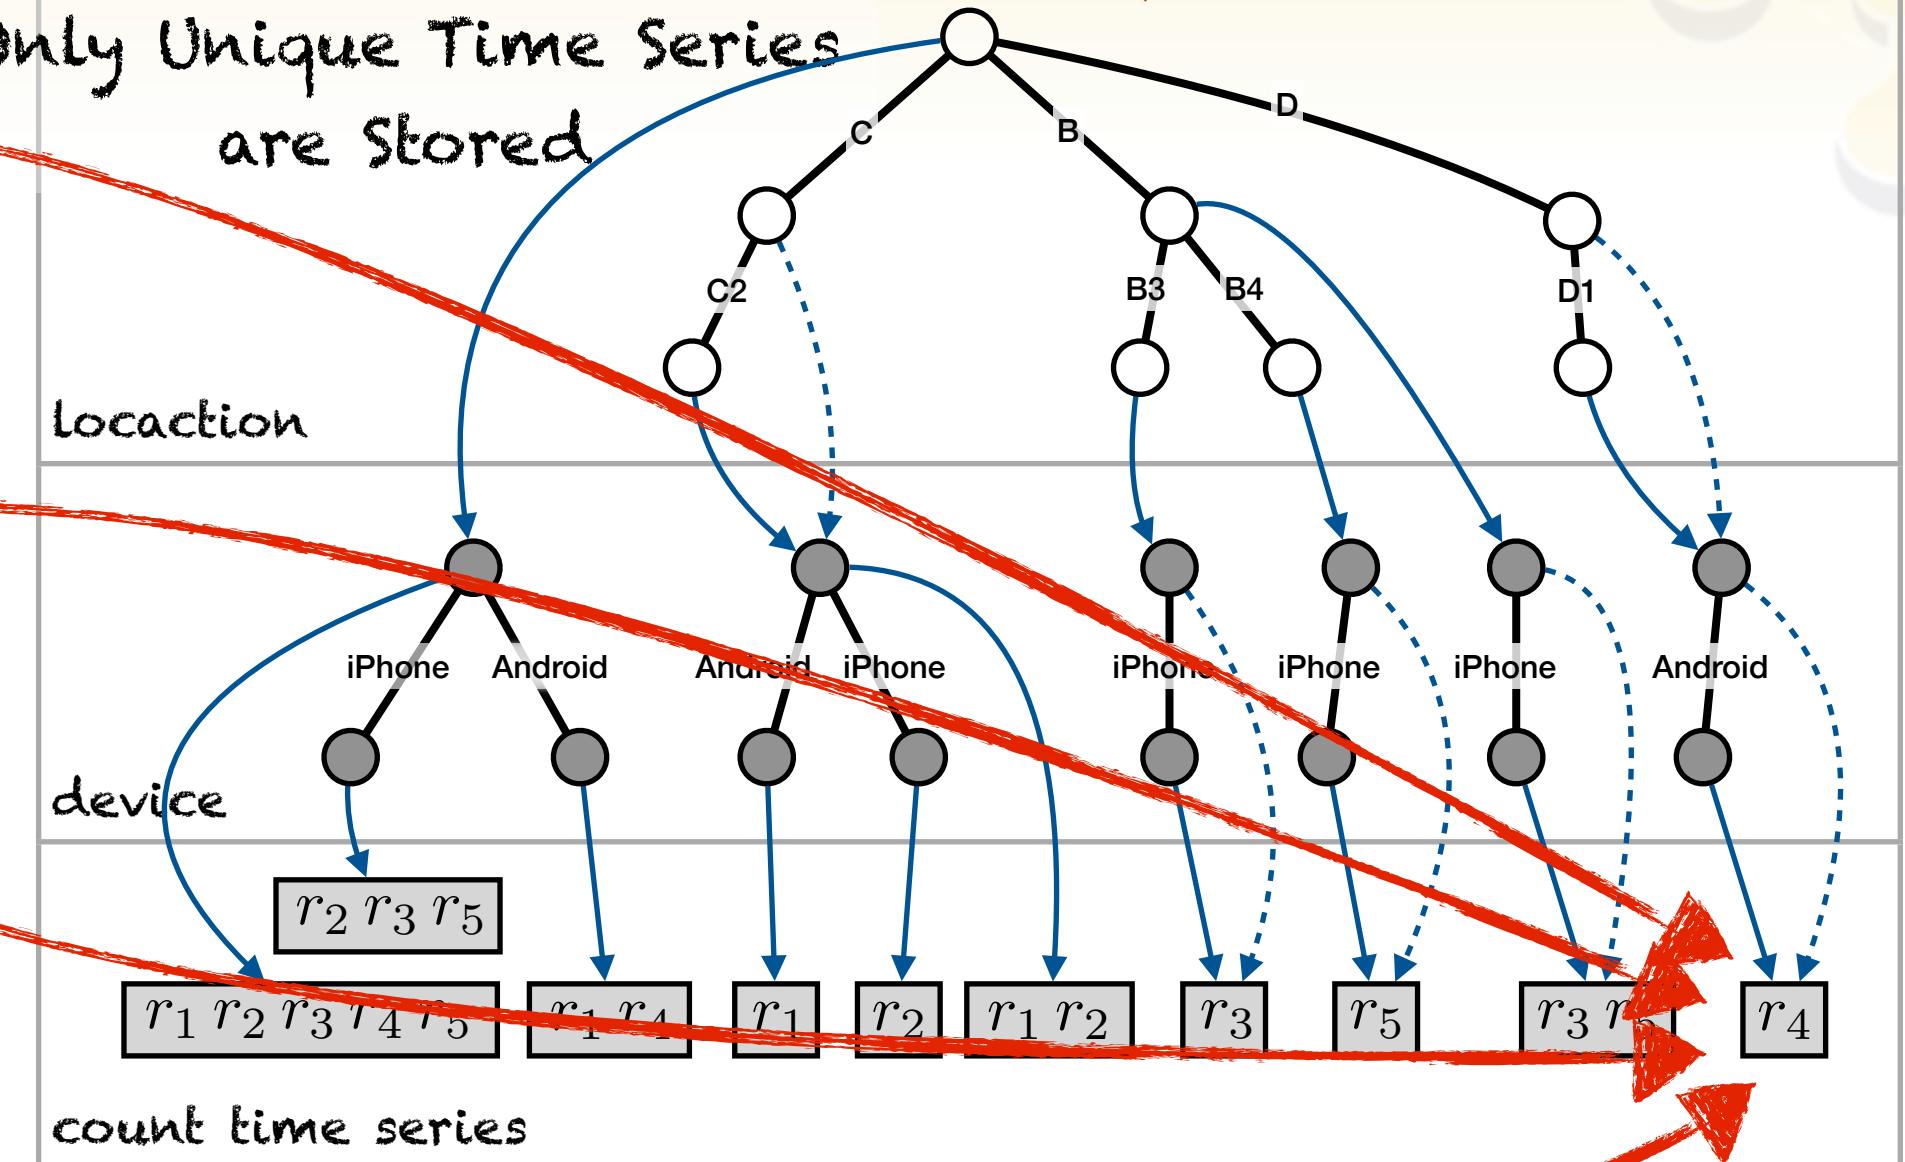


# Building a Nanocube

[2002, 2003 Sismanis et al.]

Loc1	Loc2	device	orig. records
*	*	*	{r1,r2,r3,r4,r5}
B	*	*	{r3,r5}
C	*	*	{r1,r2}
D	*	*	{r4}
*	*	Android	{r1,r4}
*	*	iPhone	{r2,r3,r5}
C2	*	*	{r1,r2}
B3	*	*	{r3}
B4	*	*	{r5}
D1	*	*	{r4}
Android	*	*	{r1}
iPhone	*	*	{r2}
C	*	Android	{r3,r5}
C	*	iPhone	{r4}
B	*	iPhone	{r3,r5}
D	*	Android	{r4}
C2	*	Android	{r1}
C2	*	iPhone	{r2}
C	*	iPhone	{r3}
B	*	Android	{r4}
B	*	Android	{r1}
B	*	iPhone	{r2}
B	*	iPhone	{r3}
B	*	iPhone	{r4}
D	*	Android	{r5}
D	*	Android	{r4}

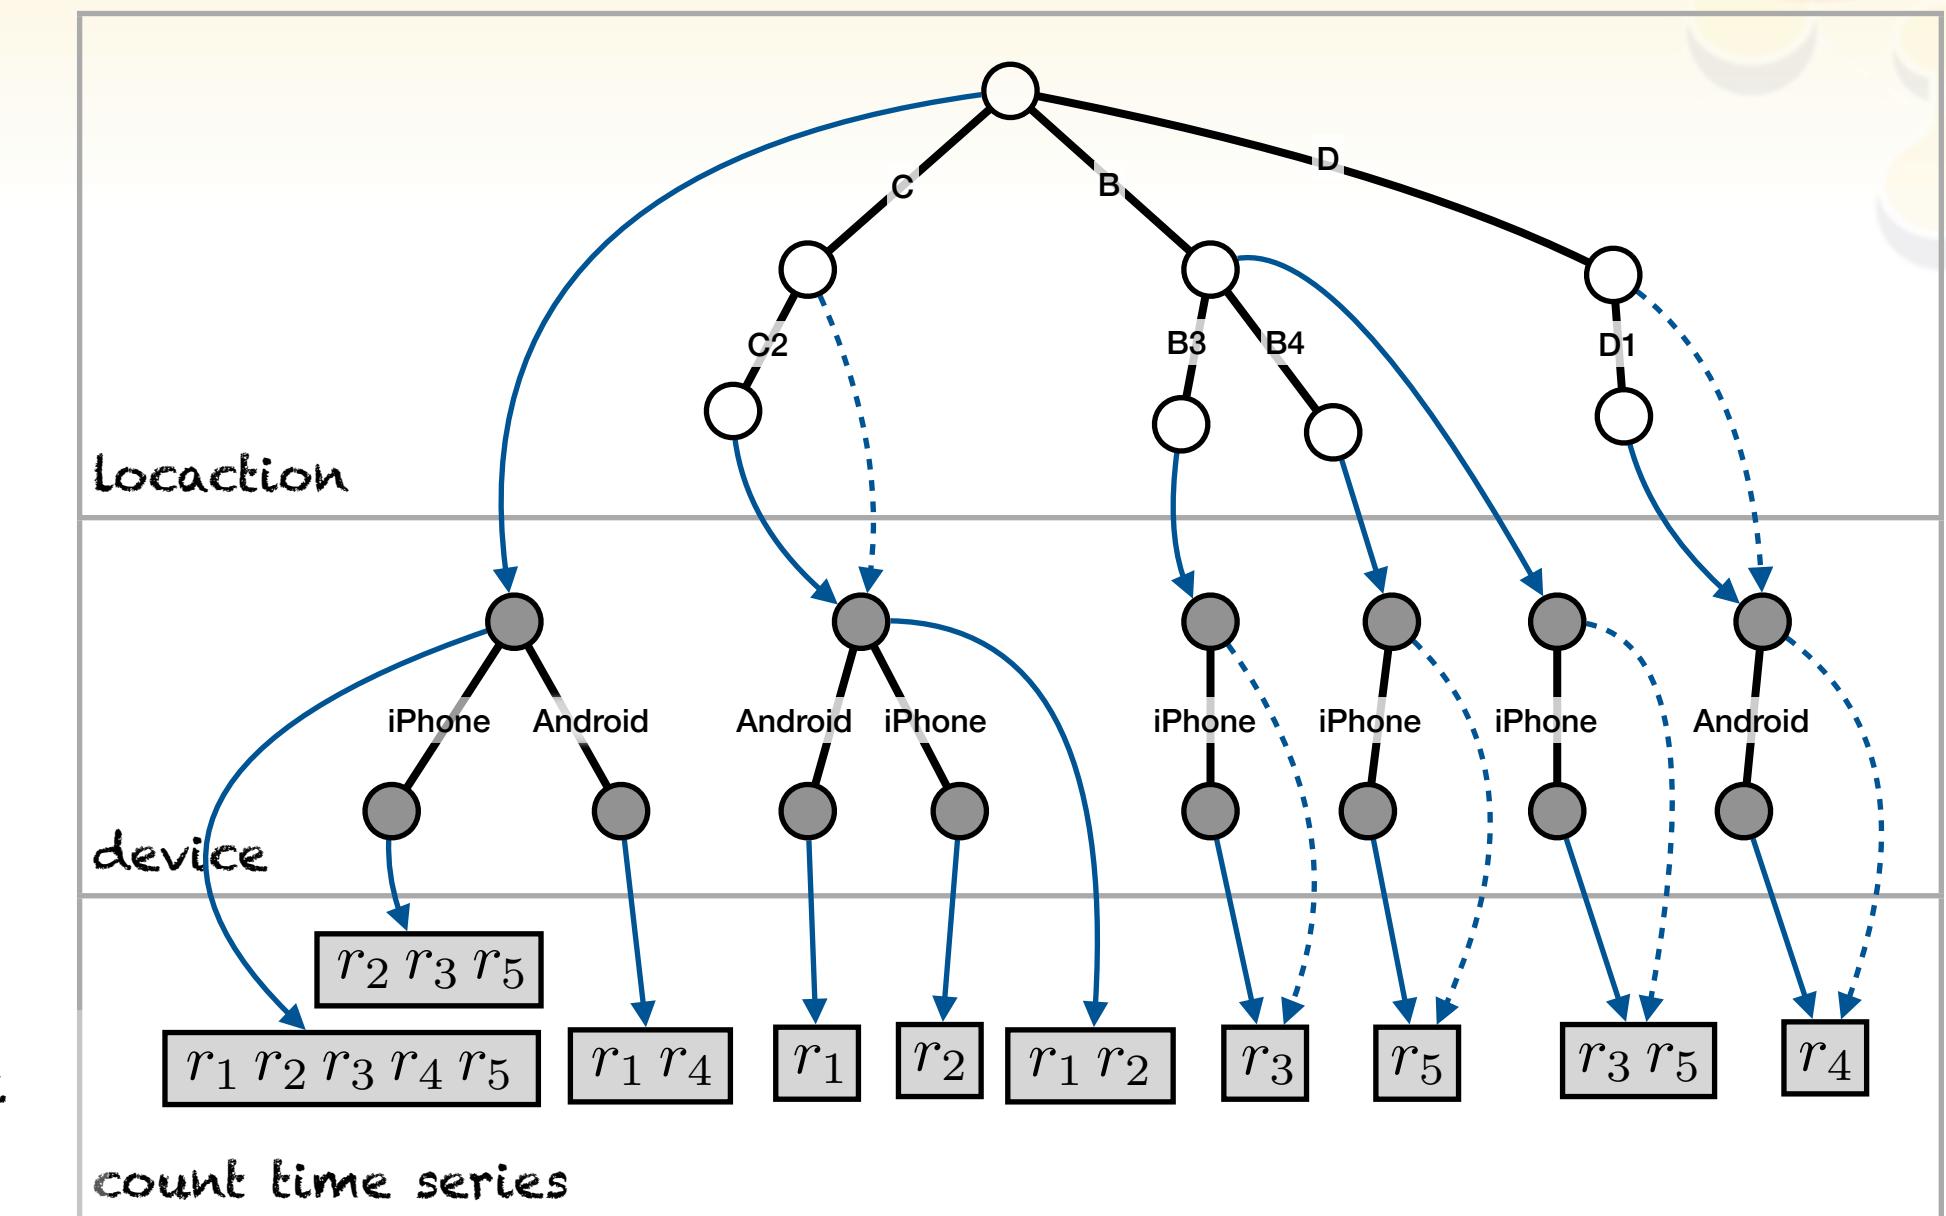
Efficient Encoding:  
Only Unique Time Series  
are Stored



# Building a Nanocube

loc1	loc2	device	time
C	C2	Android	1
C	C2	iPhone	2
B	B3	iPhone	3
D	D1	Android	4
B	B4	iPhone	5

Nanocube for Roll Up Cube  
{ < loc1, loc2 >, <device> }  
w/ Count Time Series



# Nanocube Trick #1 for Memory Efficiency

Loc1	Loc2	device	orig. records
*	*	*	{r1,r2,r3,r4,r5}
B	*	*	{r3,r5}
C	*	*	{r1,r2}
D	*	*	{r4}
*	*	Android	{r1,r4}
*	*	iPhone	{r2,r3,r5}
C	C2	*	{r1,r2}
B	B3	*	{r3}
B	B4	*	{r5}
D	D1	*	{r4}
C	*	Android	{r1}
C	*	iPhone	{r2}
B	*	iPhone	{r3,r5}
D	*	Android	{r4}
C	C2	Android	{r1}
C	C2	iPhone	{r2}
B	B3	iPhone	{r3}
B	B4	iPhone	{r5}
D	D1	Android	{r4}

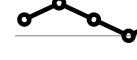
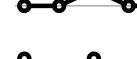
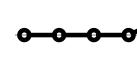
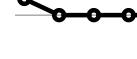
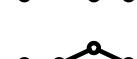
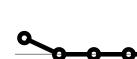
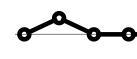
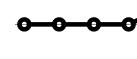
# Nanocube Trick #1 for Memory Efficiency

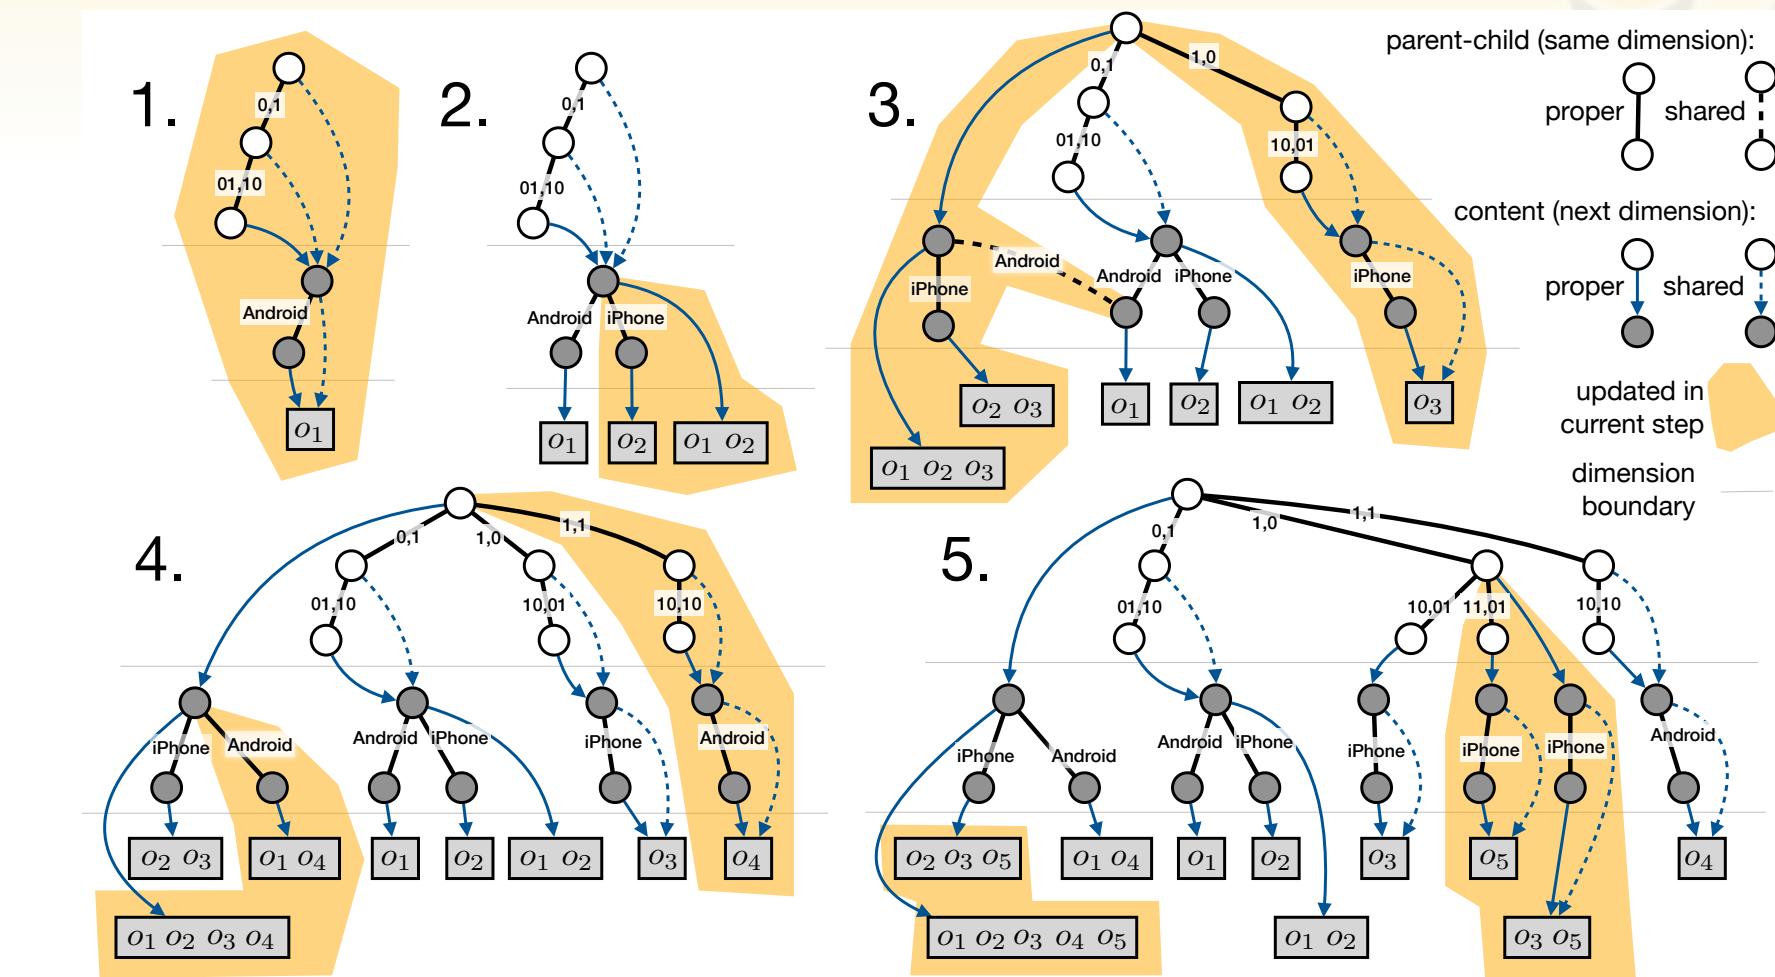
Loc1	Loc2	device	orig. records
*	*	*	{r1,r2,r3,r4,r5}
B	*	*	{r3,r5}
C	*	*	{r1,r2}
D	*	*	{r4}
*	*	Android	{r1,r4}
*	*	iPhone	{r2,r3,r5}
C	C2	*	{r1,r2}
B	B3	*	{r3}
B	B4	*	{r5}
D	D1	*	{r4}
C	*	Android	{r1}
C	*	iPhone	{r2}
B	*	iPhone	{r3,r5}
D	*	Android	{r4}
C	C2	Android	{r1}
C	C2	iPhone	{r2}
B	B3	iPhone	{r3}
B	B4	iPhone	{r5}
D	D1	Android	{r4}

# Nanocube Trick #1 for Memory Efficiency

Loc1	Loc2	device	orig. records
*	*	*	{r1,r2,r3,r4,r5}
B	*	*	{r3,r5}
C	*	*	{r1,r2}
D	*	*	{r4}
*	*	Android	{r1,r4}
*	*	iPhone	{r2,r3,r5}
C	C2	*	{r1,r2}
B	B3	*	{r3}
B	B4	*	{r5}
D	D1	*	{r4}
C	*	Android	{r1}
C	*	iPhone	{r2}
B	*	iPhone	{r3,r5}
D	*	Android	{r4}
C	C2	Android	{r1}
C	C2	iPhone	{r2}
B	B3	iPhone	{r3}
B	B4	iPhone	{r5}
D	D1	Android	{r4}

# Nanocube's Encoding

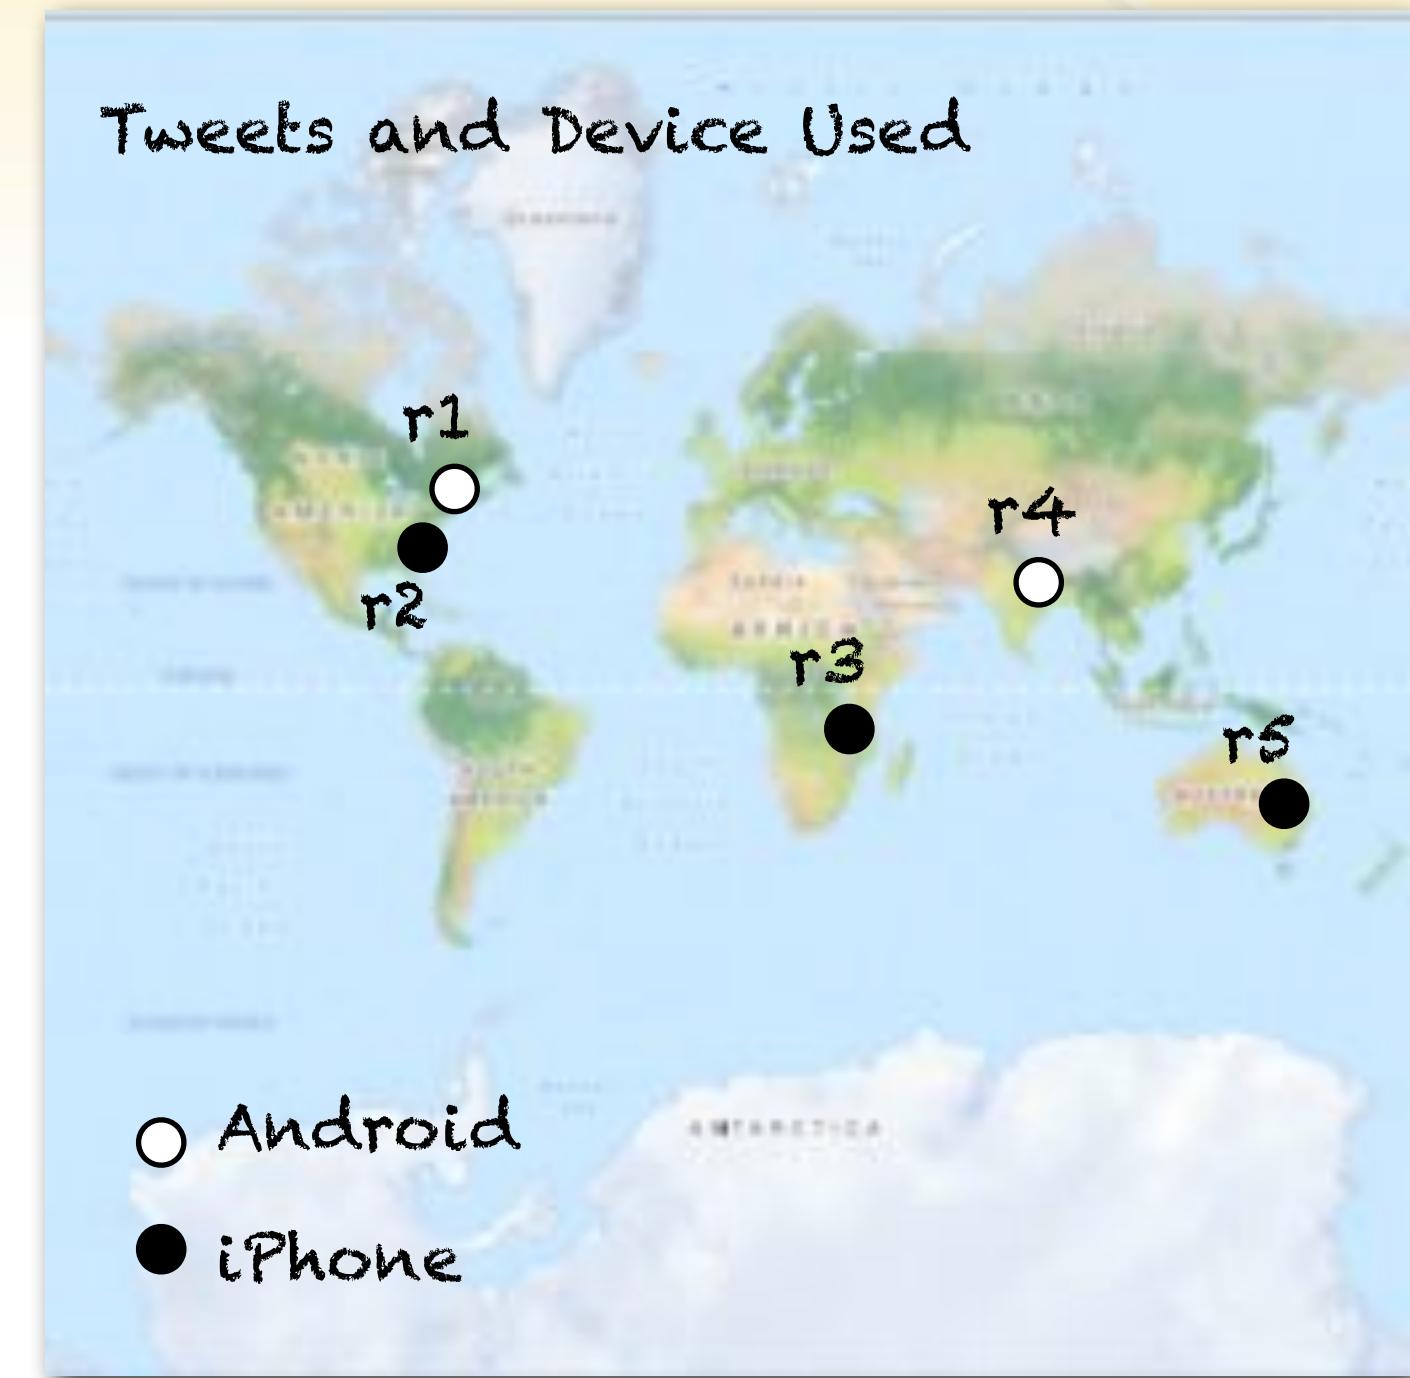
loc1	loc2	device	time series
*	*	*	
B	*	*	
C	*	*	
D	*	*	
*	*	Android	
*	*	iPhone	
C2	*	*	
B3	*	*	
B4	*	*	
D1	*	Android	
C	*	iPhone	
C	*	iPhone	
B	*	Android	
D	*	Android	
C	C2	Android	
C	C2	iPhone	
B	B3	iPhone	
B	B4	iPhone	
D	D1	Android	



# Nanocube Trick #1 for Memory Efficiency

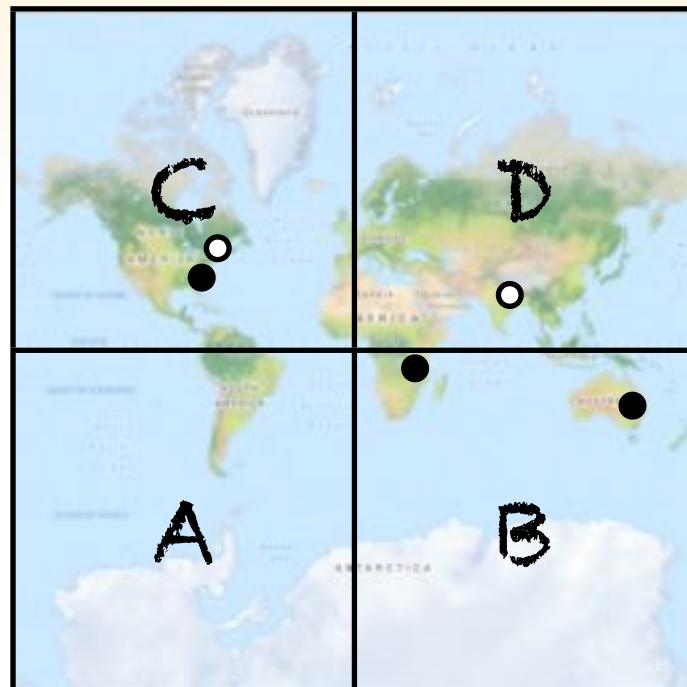
Loc1	Loc2	device	orig. records
*	*	*	{r1,r2,r3,r4,r5}
B	*	*	{r3,r5}
C	*	*	{r1,r2}
D	*	*	{r4}
*	*	Android	{r1,r4}
*	*	iPhone	{r2,r3,r5}
C	C2	*	{r1,r2}
B	B3	*	{r3}
B	B4	*	{r5}
D	D1	*	{r4}
C	*	Android	{r1}
C	*	iPhone	{r2}
B	*	iPhone	{r3,r5}
D	*	Android	{r4}
C	C2	Android	{r1}
C	C2	iPhone	{r2}
B	B3	iPhone	{r3}
B	B4	iPhone	{r5}
D	D1	Android	{r4}

# What to pre-compute? Data Cubes!



# What to pre-compute? Data Cubes!

Loc1



Loc2

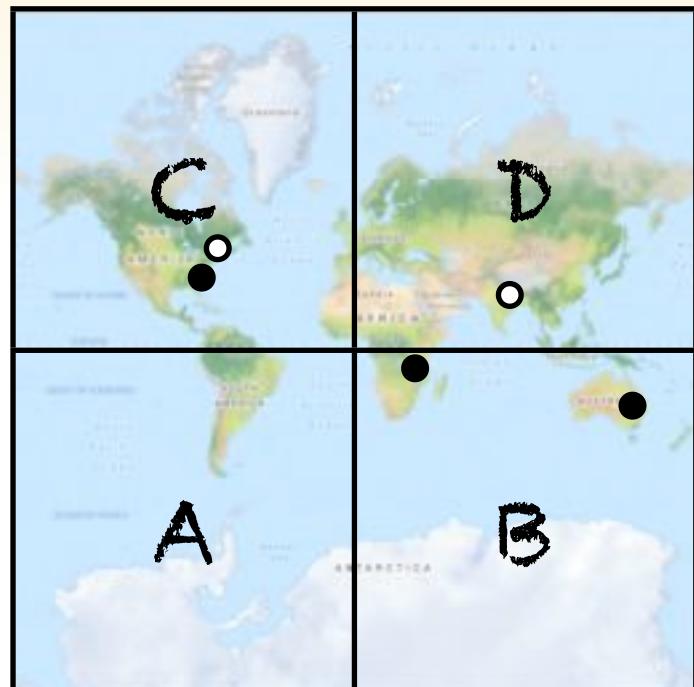
C3	C4	D3	D4
C1	C2	D1	D2
A3	A4	B3	B4
A1	A2	B1	B2

device:

- Android
- iPhone

# What to pre-compute? Data Cubes!

loc1



loc2

C3	C4	D3	D4
C1	C2	D1	D2
A3	A4	B3	B4
A1	A2	B1	B2

device:

- Android
- iPhone

loc1	loc2	device
C	C2	Android
C	C2	iPhone
B	B3	iPhone
D	D1	Android
B	B4	iPhone

relation

# What to pre-compute? Data Cubes!

loc1	loc2	device
C	C2	Android
C	C2	iPhone
B	B3	iPhone
D	D1	Android
B	B4	iPhone

group by { device }



Loc1	Loc2	device	orig. rows
*	*	Android	{ r1, r4 }
*	*	iPhone	{ r2, r3, r5 }

relation

# What to pre-compute? Data Cubes!

loc1	loc2	device
C	C2	Android
C	C2	iPhone
B	B3	iPhone
D	D1	Android
B	B4	iPhone

relation

group by { loc1, device }



Loc1	Loc2	device	orig. rows
C	*	Android	{ r1 }
C	*	iPhone	{ r2 }
B	*	iPhone	{ r3, r5 }
D	*	Android	{ r4 }

# What to pre-compute? Data Cubes!

loc1	loc2	device
C	C2	Android
C	C2	iPhone
B	B3	iPhone
D	D1	Android
B	B4	iPhone

relation

**cube { loc1, loc2, device }**



= group by {} +  
group by { loc1 } +  
group by { loc2 } +  
group by { device } +  
group by { loc1, loc2 } +  
group by { loc1, device } +  
group by { loc2, device } +  
group by { loc1, loc2, device }

# What to pre-compute? Data Cubes!

loc1	loc2	device
c	c2	Android
c	c2	iPhone
		iPhone
		Android
b	b4	iPhone

By construction  
these two cases

are redundant

relation

$\text{cube } \{ \text{loc1, loc2, device} \}$   
 $= \text{group by } \{ \} +$   
 $\text{group by } \{ \text{loc1} \} +$   
 $\boxed{\text{group by } \{ \text{loc2} \} +}$   
 $\text{group by } \{ \text{device} \} +$   
 $\boxed{\text{group by } \{ \text{loc1, loc2} \} +}$   
 $\text{group by } \{ \text{loc1, device} \} +$   
 $\text{group by } \{ \text{loc2, device} \} +$   
 $\text{group by } \{ \text{loc1, loc2, device} \}$

# What to pre-compute? Data Cubes!

loc1	loc2	device
C	C2	Android
C	C2	iPhone
B	B3	iPhone
D	D3	Android
B	B4	iPhone

By construction  
these two cases  
are redundant

cube { loc1, loc2, device }

= group by {} +

group by { loc1 } +

group by { loc2 } +

group by { device } +

group by { loc1, loc2 } +

Roll Up on < loc1, loc2 > = group by {} +  
group by { loc1 } + group by { loc1, loc2 }

# What to pre-compute? Data Cubes!

ROLL UP Cube { < loc1, loc2 >, <device> }

loc1	loc2	device
C	C2	Android
C	C2	iPhone
B	B3	iPhone
D	D1	Android
B	B4	iPhone

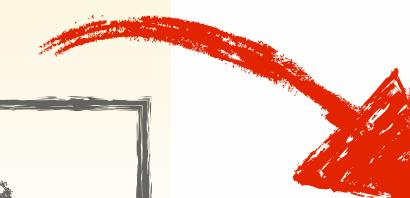
relation



= group by {} +  
group by { loc1 } +  
~~group by { loc2 } +~~  
group by { device } +  
group by { loc1, loc2 } +  
group by { loc1, device } +  
~~group by { loc2, device } +~~  
group by { loc1, loc2, device }

# What to pre-compute? Data Cubes!

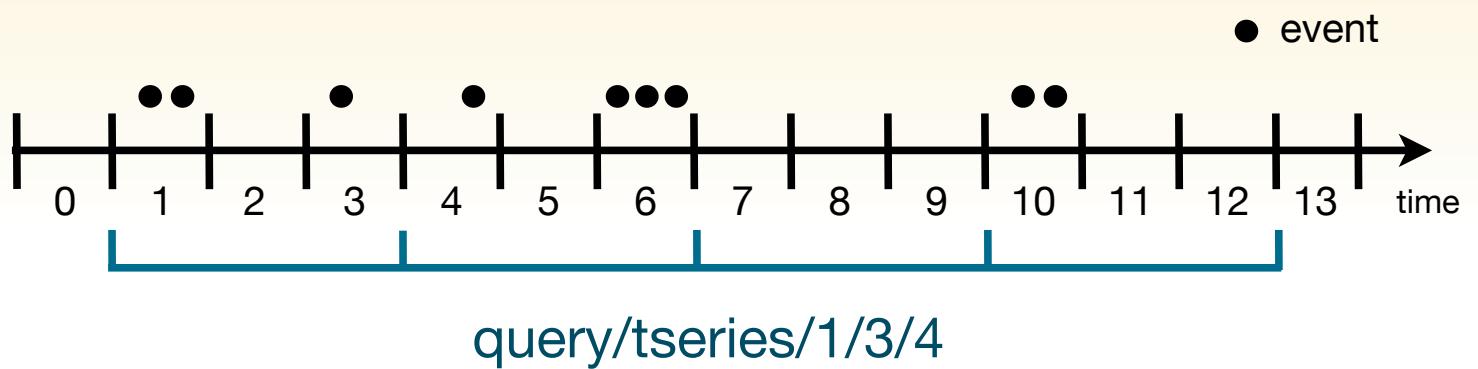
ROLL UP Cube { < loc1, loc2 >, <device> }



**relation**

Loc1	Loc2	device	Loc1	Loc2	device	orig. records
C	C2	Android	*	*	*	{r1,r2,r3,r4,r5}
C	C2	iPhone	B	*	*	{r3,r5}
B	B3	iPhone	C	*	*	{r1,r2}
D	D1	Android	D	*	*	{r4}
B	B4	iPhone	*	*	Android	{r1,r4}
			*	*	iPhone	{r2,r3,r5}
			C	C2	*	{r1,r2}
			B	B3	*	{r3}
			B	B4	*	{r5}
			D	D1	*	{r4}
			C	*	Android	{r1}
			C	*	iPhone	{r2}
			B	*	iPhone	{r3,r5}
			D	*	Android	{r4}
			C	C2	Android	{r1}
			C	C2	iPhone	{r2}
			B	B3	iPhone	{r3}
			B	B4	iPhone	{r5}
			D	D1	Android	{r4}

# Time Series Encoding



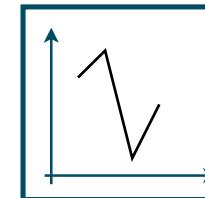
start at bin 1, use buckets of 3 bins each, and collect 4 of these buckets

solve using...

bin: 1 accum: 2	bin: 3 accum: 3	bin: 4 accum: 4	bin: 6 accum: 7	bin: 10 accum: 9
--------------------	--------------------	--------------------	--------------------	---------------------

A Summed Table Sparse Representation for Counts

result



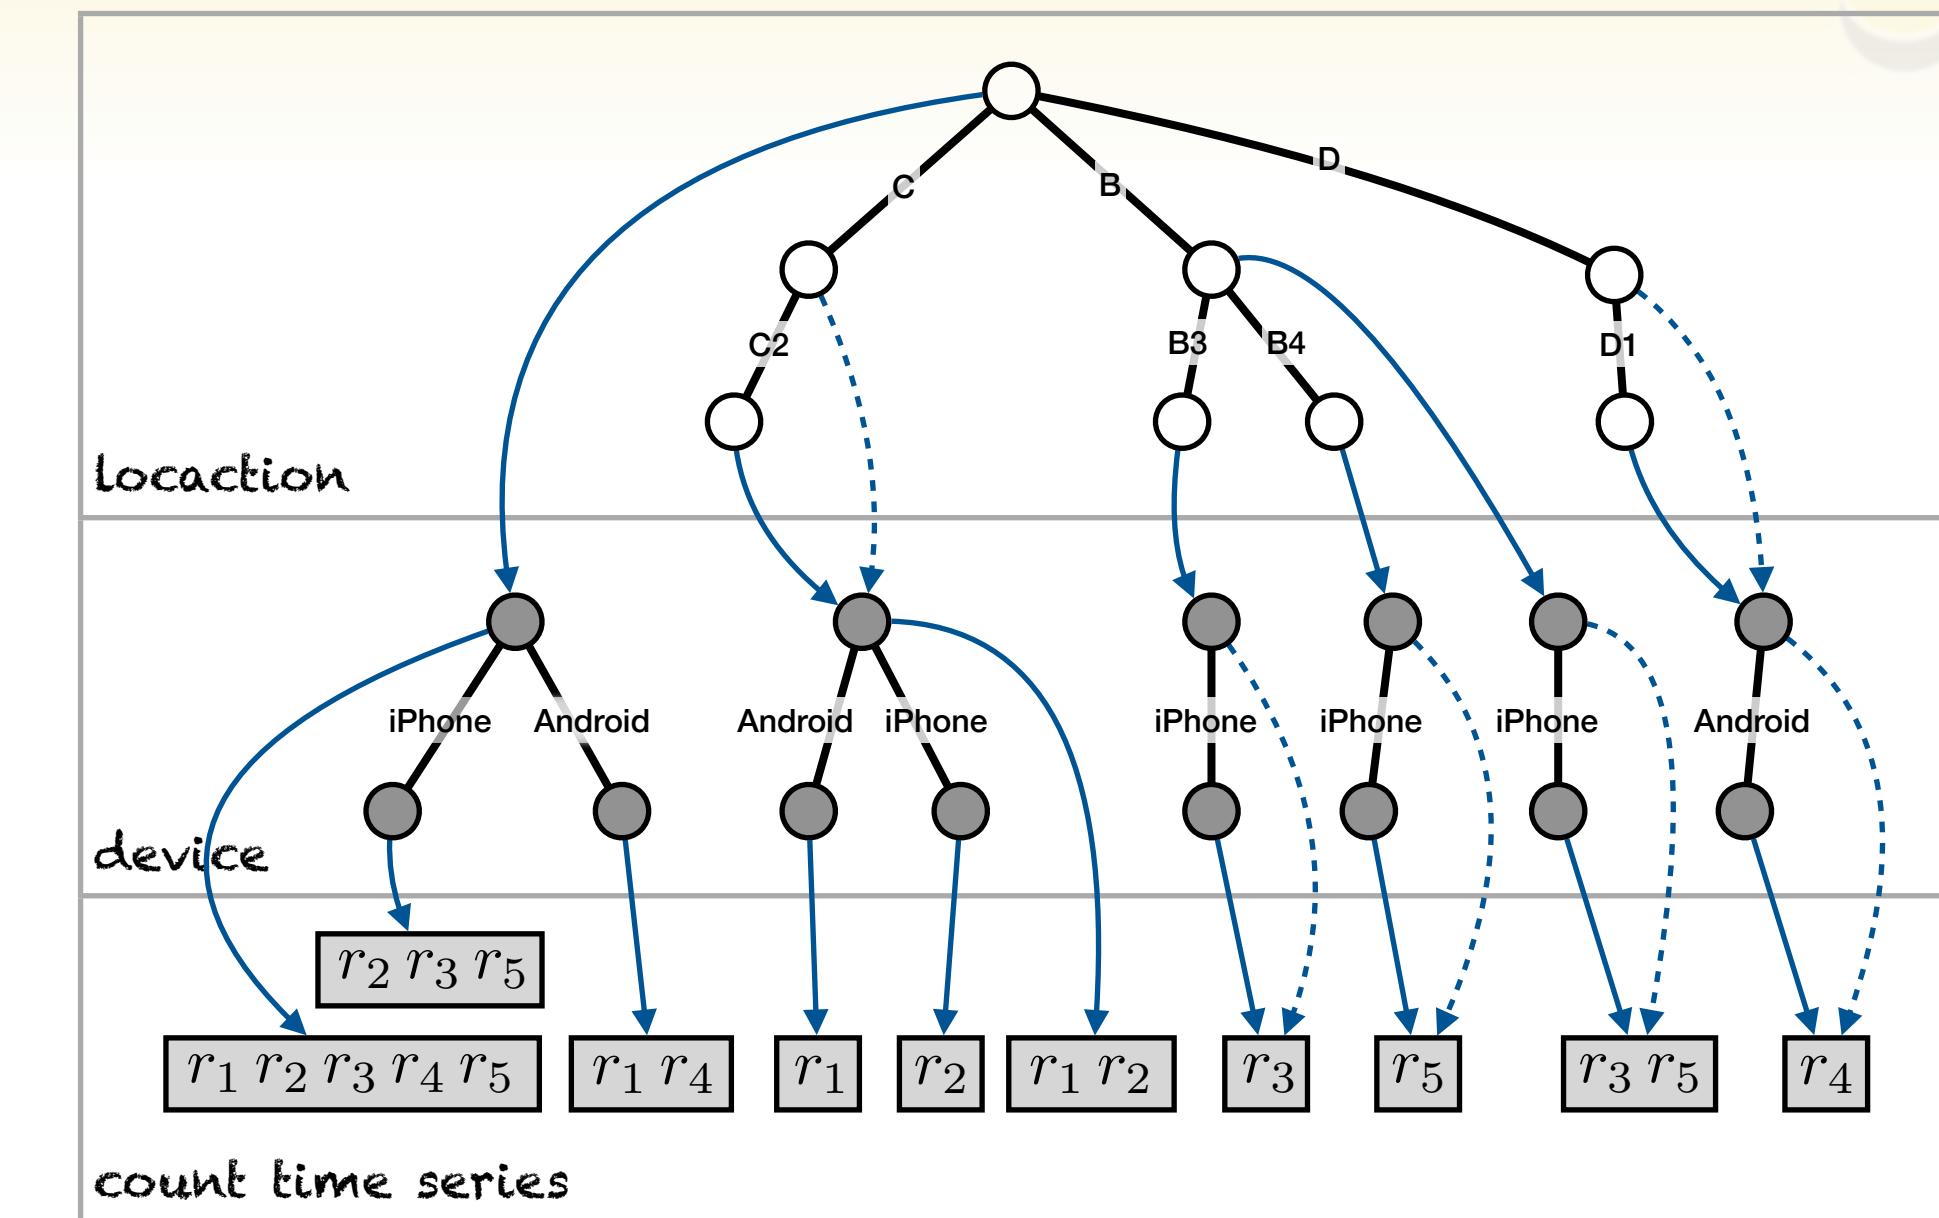
3	4	0	2
---	---	---	---

# Querying a Nanocube

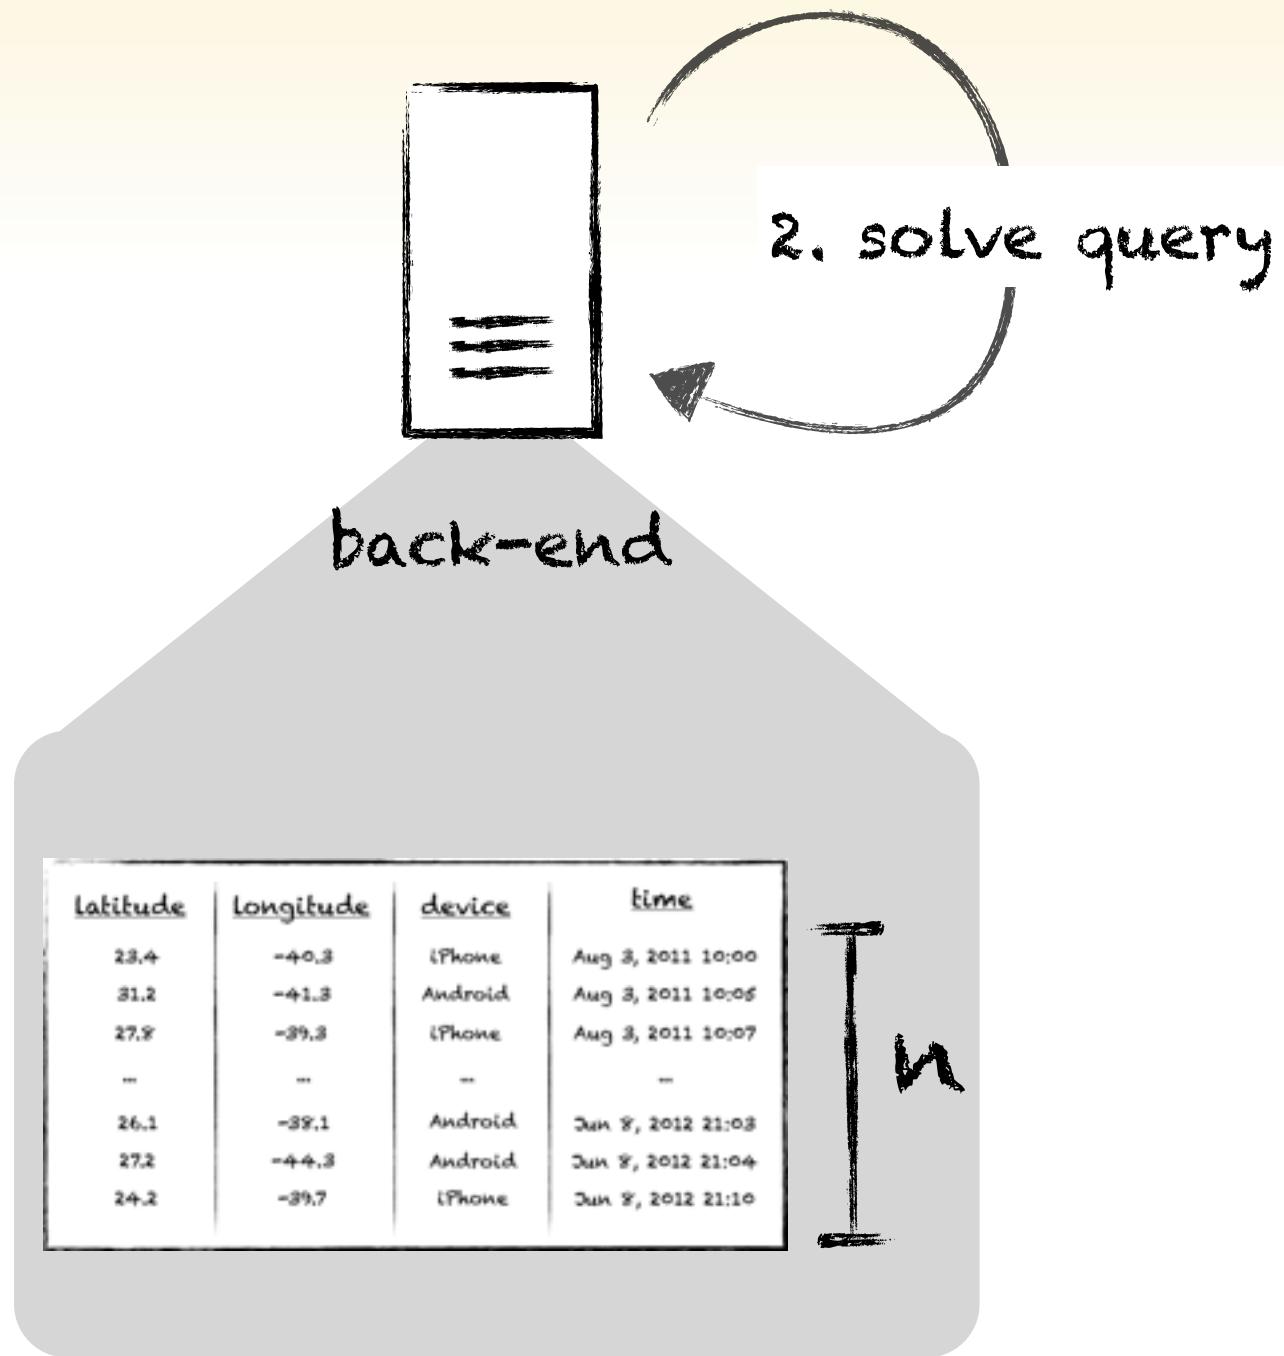
Q1: ALL Tweets

Q2: ALL iPhone Tweets

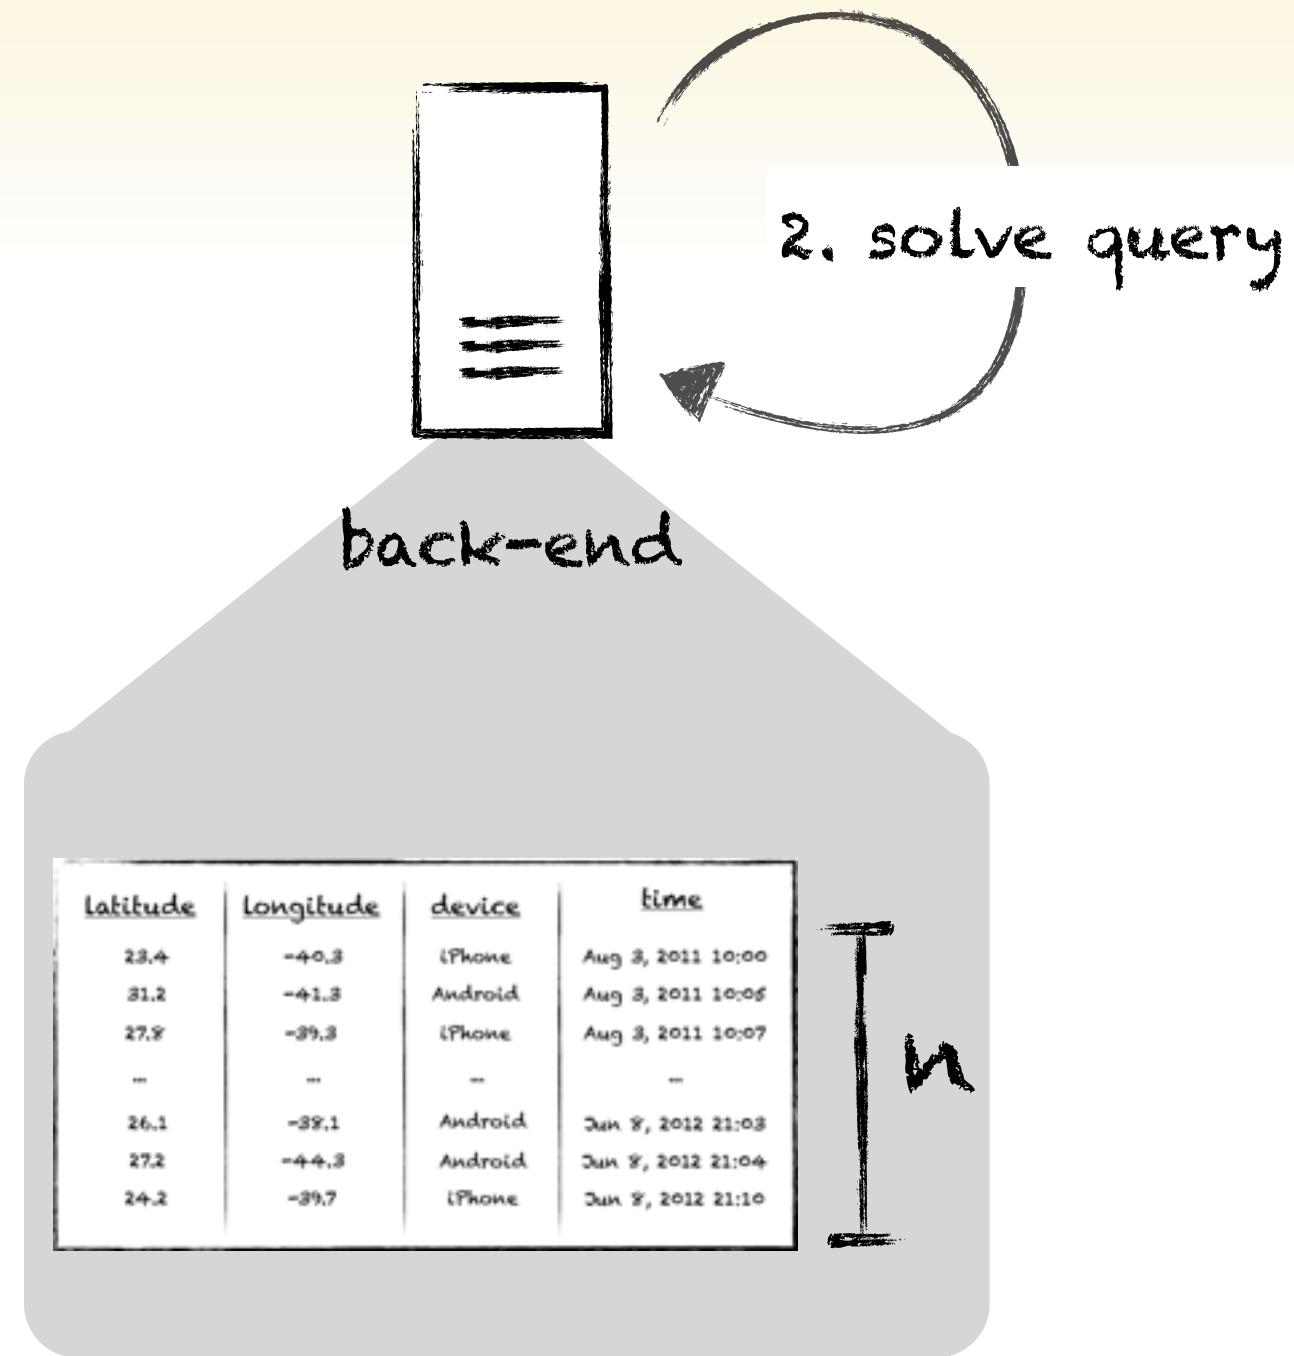
Q2: All Android Tweets  
in Space Cell C



# Solving Queries



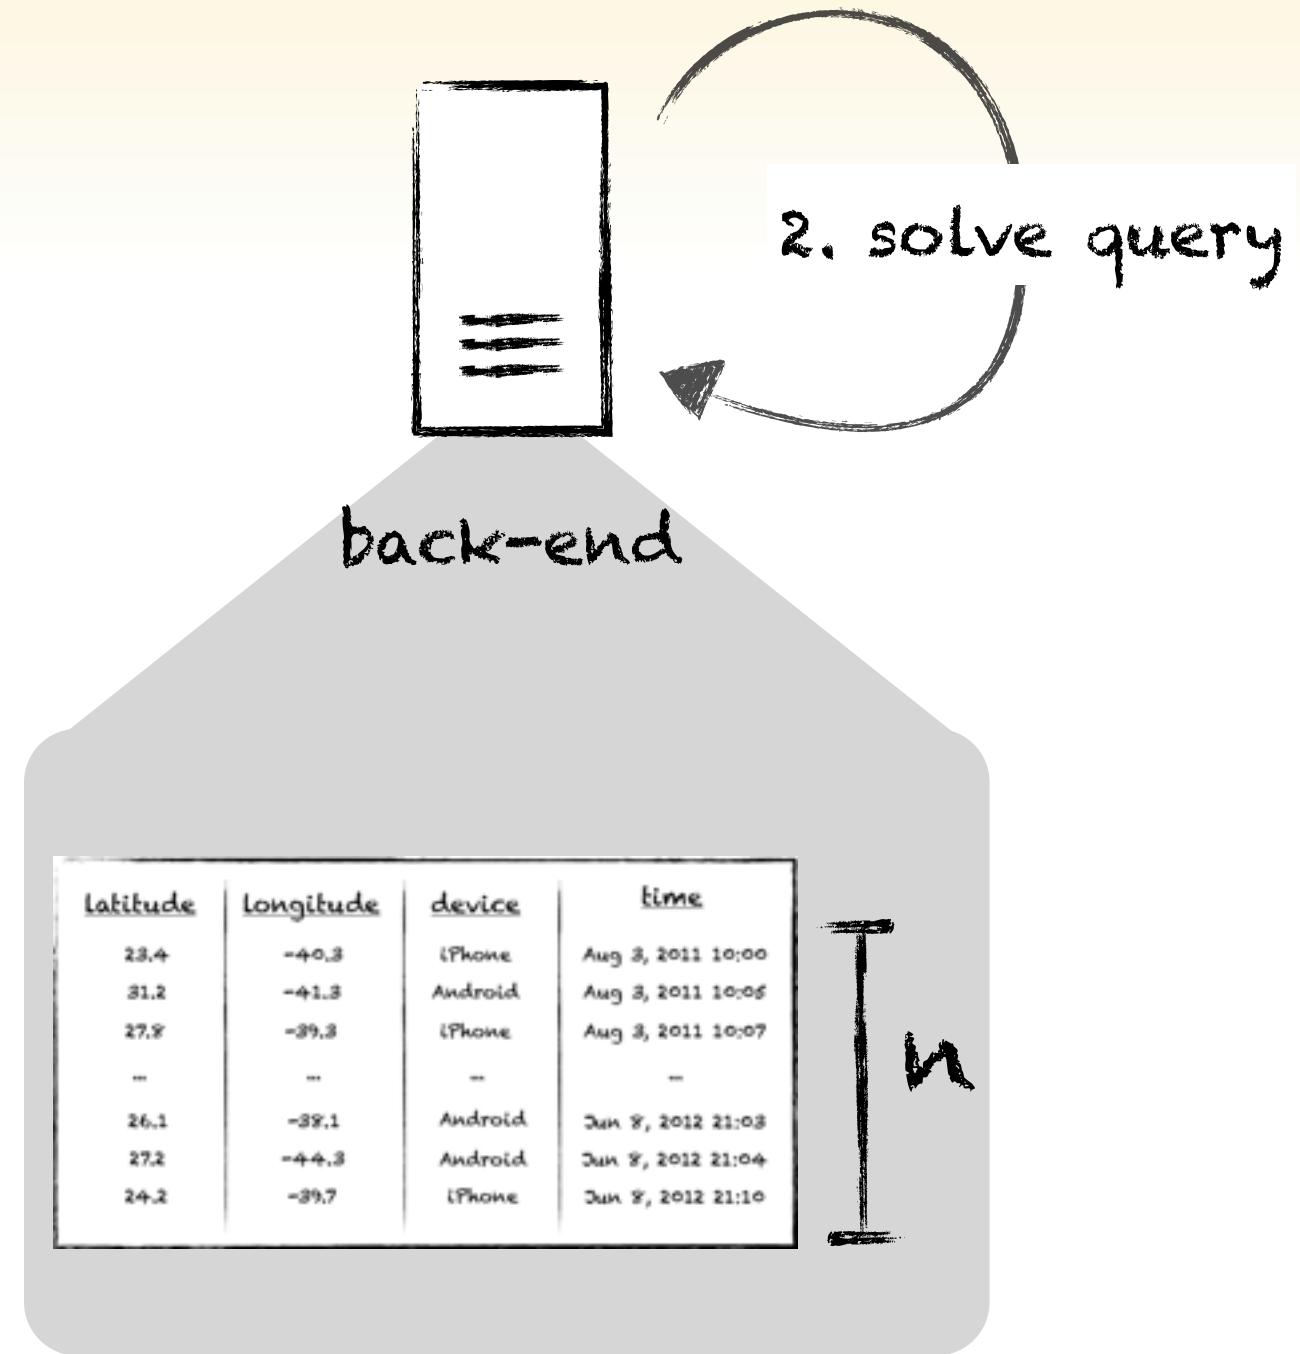
# Solving Queries



# Approaches

(small) If "n" is small we can scan the data and solve the query

# Solving Queries

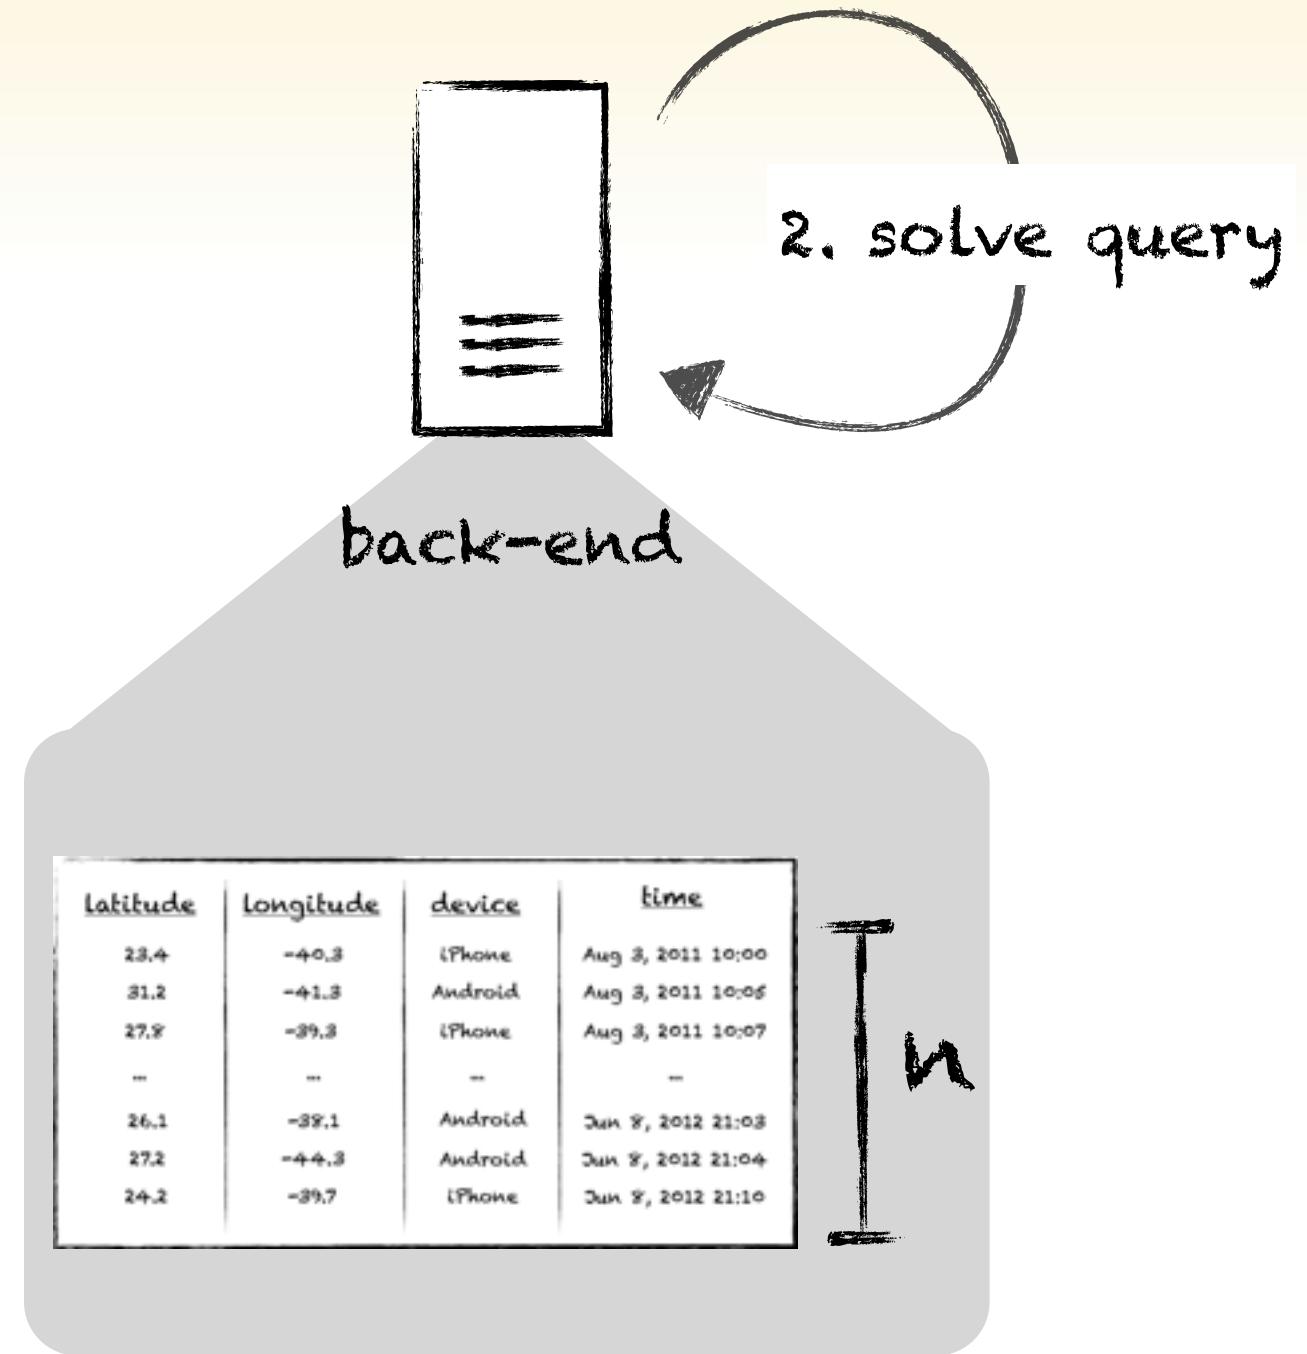


# Approaches

(small) If "n" is small we can scan the data and solve the query

(parallel) If "n" is large we can make "k" parallel processors deal with " $n/k$ " data points

# Solving Queries



# Approaches

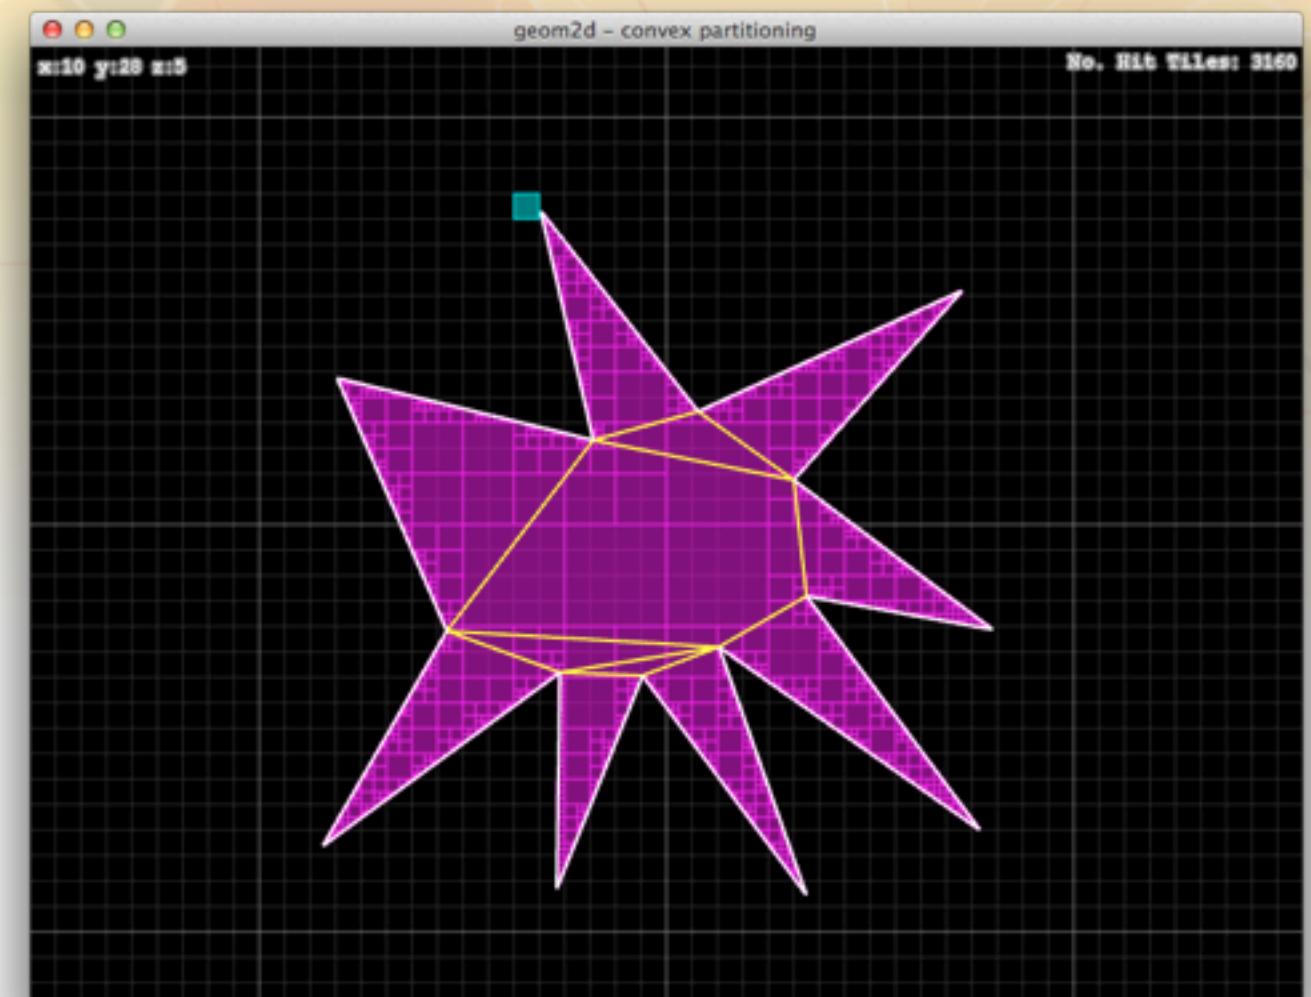
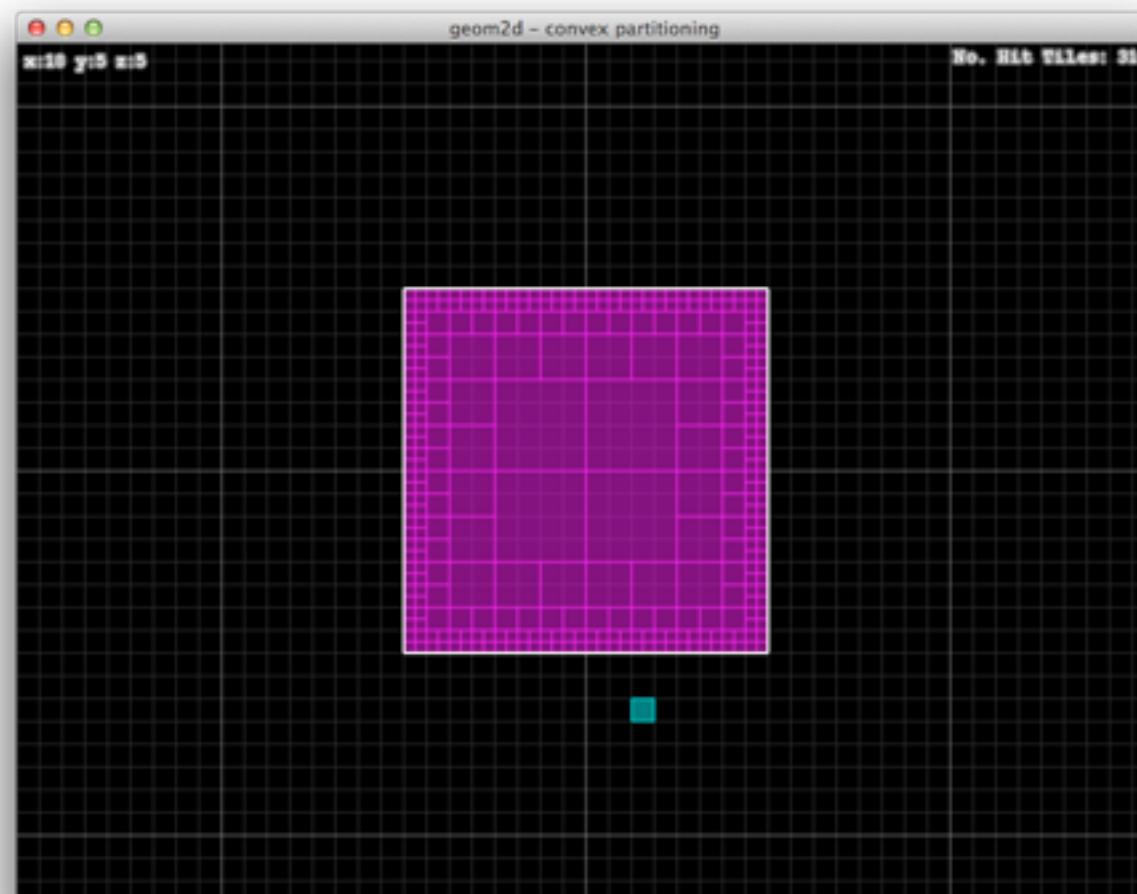
(small) If "n" is small we can scan the data and solve the query

(parallel) If "n" is large we can make "k" parallel processors deal with " $n/k$ " data points

(pre-compute) Pre-compute answers so that at query time computation is not proportional to "n" but to something smaller (e.g. output size)

DEMO

What if The Shape of  
the query doesn't match  
a predefined shape?



# Results

4 to 6 dimensions  
(space counting 2)

dataset	n	memory	time	keys	cardinality	schema
brightkite	4.5 M	1.6 GB	3.50 m	3.5 M	$2^{74}$	spatial(50), time(16), weekday(3), hour(5)
customer tix	7.8 M	2.5 GB	8.47 m	7.8 M	$2^{69}$	spatial(50), time(16), type(3)
flights	121 M	2.3 GB	31.13 m	43.3 M	$2^{75}$	spatial(50), time(16), carrier(5), delay(4)
twitter-small	210 M	10.2 GB	1.23 h	116.0 M	$2^{53}$	spatial(34), time(16), device(3)
twitter	210 M	46.4 GB	5.87 h	136.0 M	$2^{60}$	spatial(34), time(16), lang(5), device(3), app(2)
splom-10	1 B	4.3 MB	4.13 h	7.4 K	$2^{20}$	d1(4), d2(4), d3(4), d4(4), d5(4)
splom-50	1 B	166 MB	4.72 h	1.9 M	$2^{30}$	d1(6), d2(6), d3(6), d4(6), d5(6)
cdrs	1 B	3.6 GB	3.08 h	96.3 M	$2^{69}$	spatial(50), time(16), duration(3)



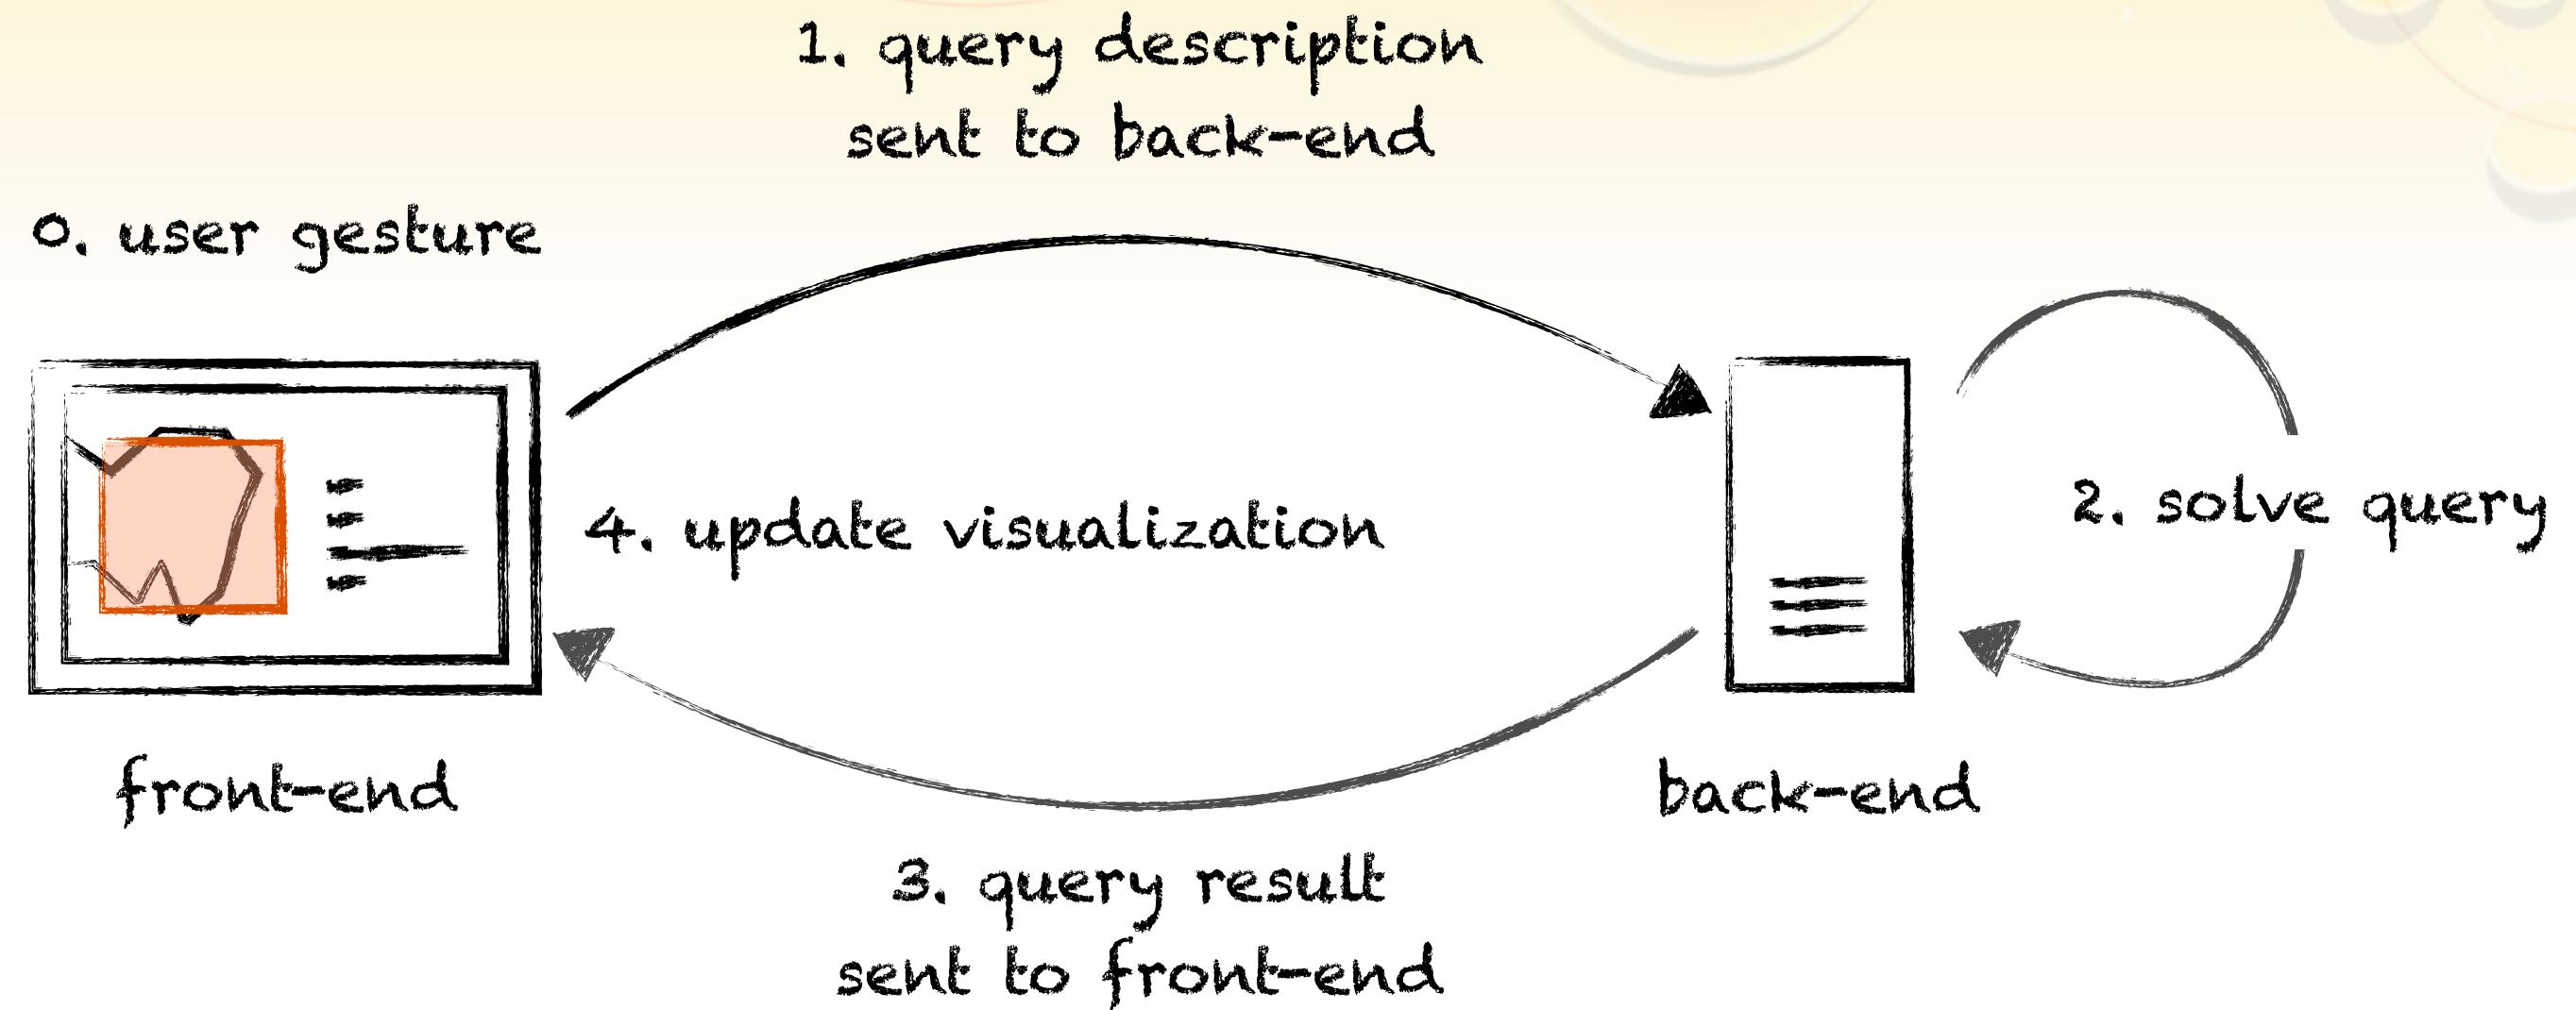
# Comparison with imMens

	imMens	Nanocubes
Technique	GPU & GPU Mem.	CPU & Main RAM
Hard Dimension Limit	$\leq 4$ dimensions	no hard limit
Max Bins. Across ALL Dim	$2^{130}$	no hard limit
Frame Rate	60fps	10 to 60fps
Memory Scheme	Dense	Sparse
Supports Streaming	No	Yes

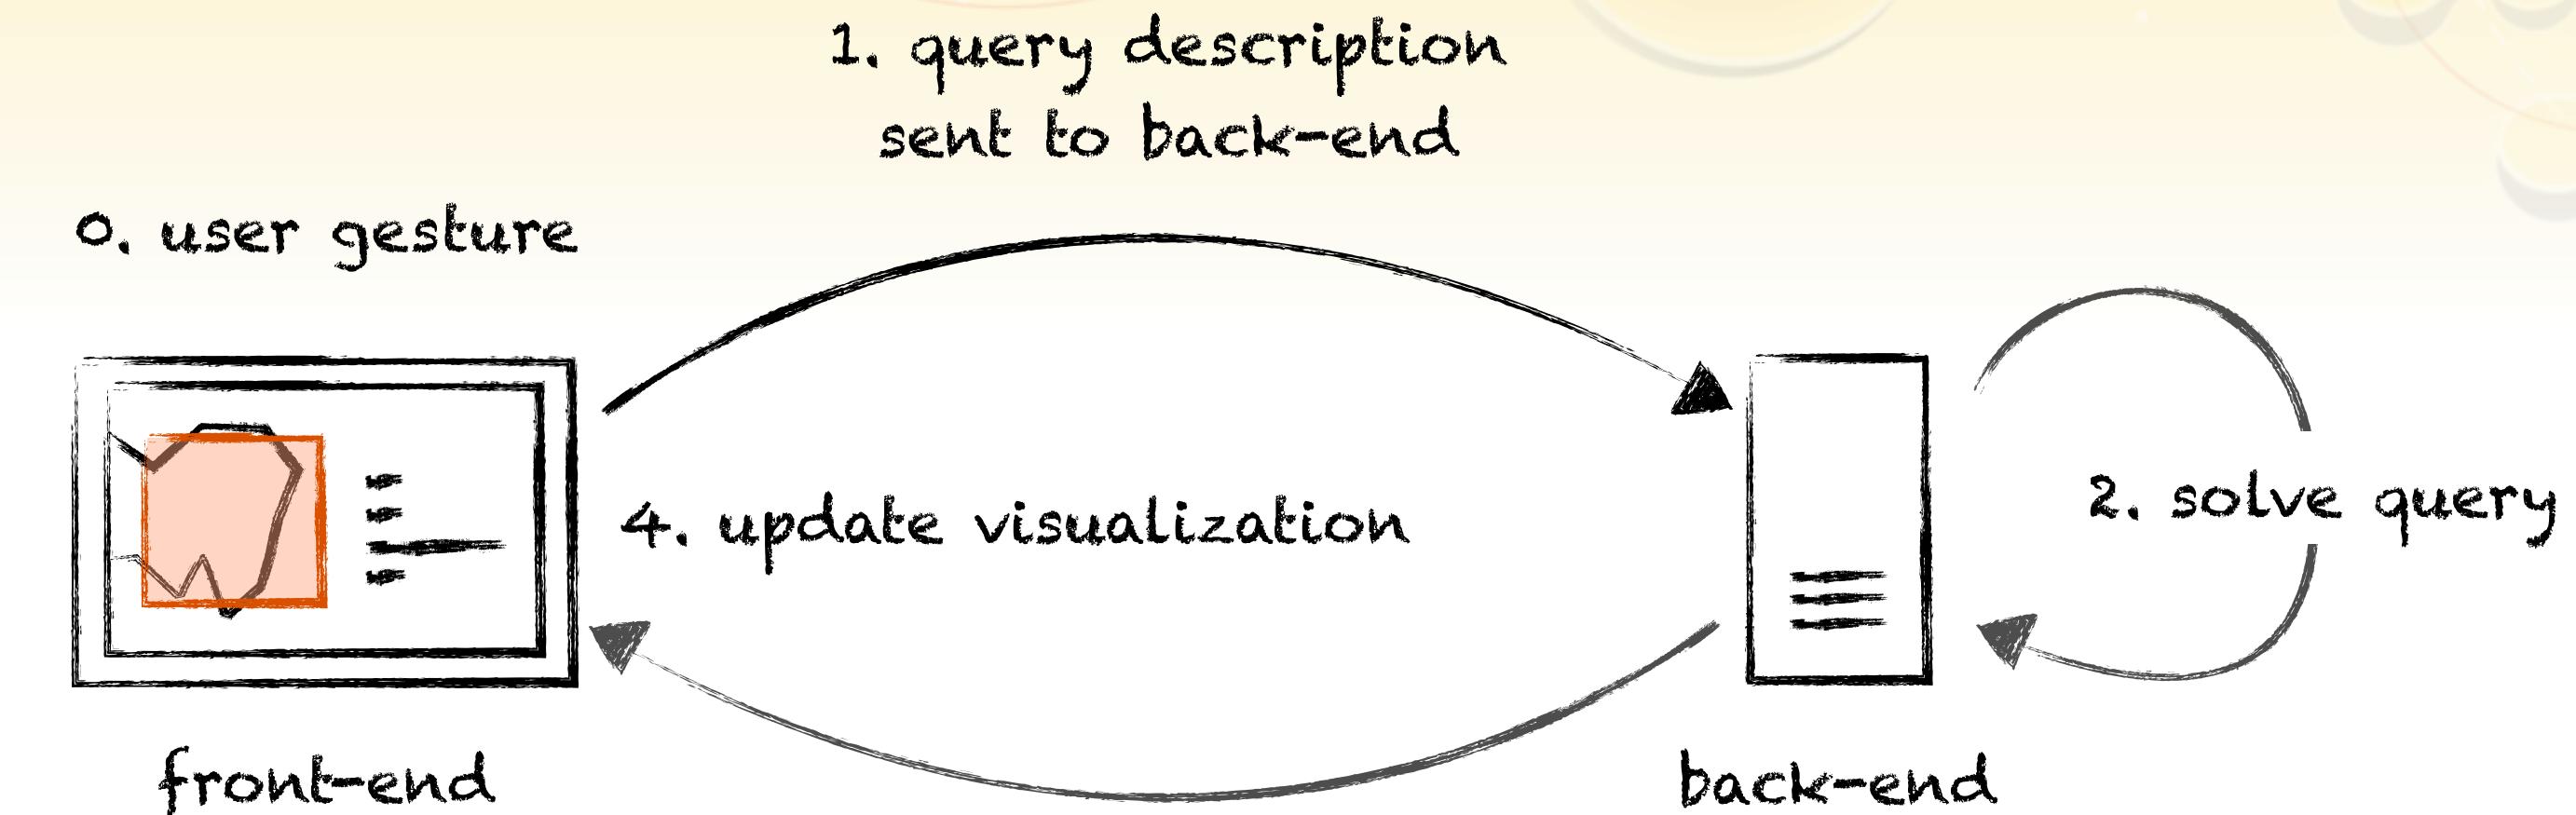
[EuroVis 2013, Z. L. Liu, B. Jiang, and J. Heer.

imMens: Real-time visual querying of big data.]

# Interaction Cycle



# Interaction Cycle



requirement

$$\text{cycle time} = t_1 + t_2 + t_3 + t_4 \leq \text{time budget}$$

# Interaction Cycle

